



# Top-Level View of a Computer

## Lecture Content

- **Instruction Cycle**
- Interrupts
- I/O Functionality
- Bus Systems
- QPI
- PCIe



# Top-Level View of a Computer

## Learning Objectives

- Understand the basic elements of an instruction cycle and the role of interrupts
- Understand the way different components are connected in a computer
- Understand the different types of bus systems
- Understand the need for multiple buses and the arrangement of them in a hierarchy


# What it Really Means

- We have learned so far how to program a computer and the basic principles of computer arithmetic and boolean algebra
- **Now we want to see how everything works together:**
  - What components are there?
  - How are they connected?
  - How does a computer execute an instruction?
- **Then we start digging deeper into the topic**
  - Memory hierarchies
  - CPU organization
  - CPU details

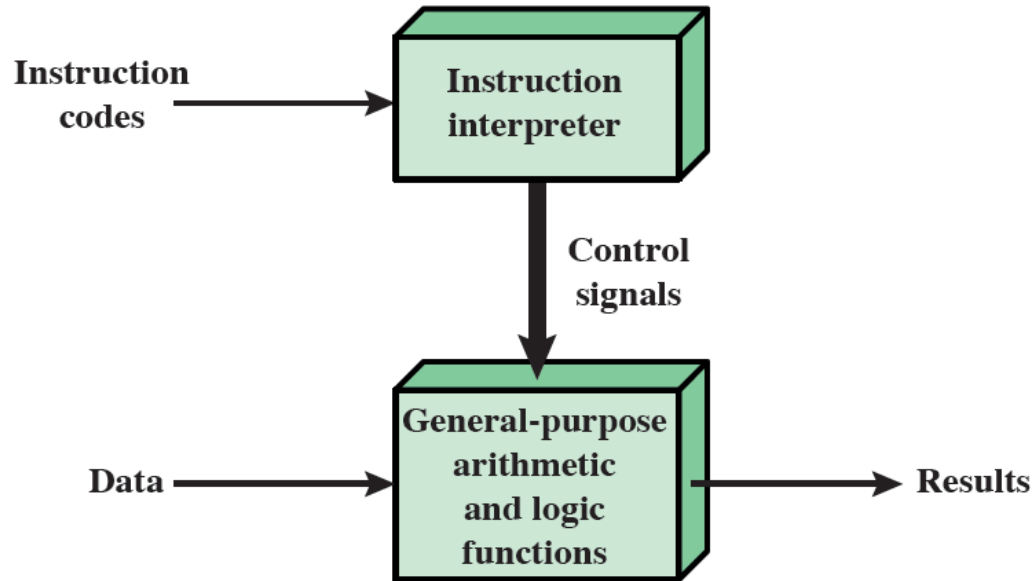
# Computer Components

- Contemporary computer designs are based on concepts developed by John von Neumann
- **Based on three key concepts**
  - Data and instructions are stored in a single read-write memory
  - The contents of this memory are addressable by location, without regard to the type of data contained there
  - Execution occurs in a sequential fashion (unless explicitly modified) from one instruction to the next
- **Hardwired program**
  - The result of the process of connecting the various components in the desired configuration

# General-Purpose vs. Hardwired



(a) Programming in hardware



(b) Programming in software

# The Top-Level View

## ■ Software

- A sequence of codes or instructions
- Is interpreted by hardware
- Much easier than rewiring the hardware ;-)

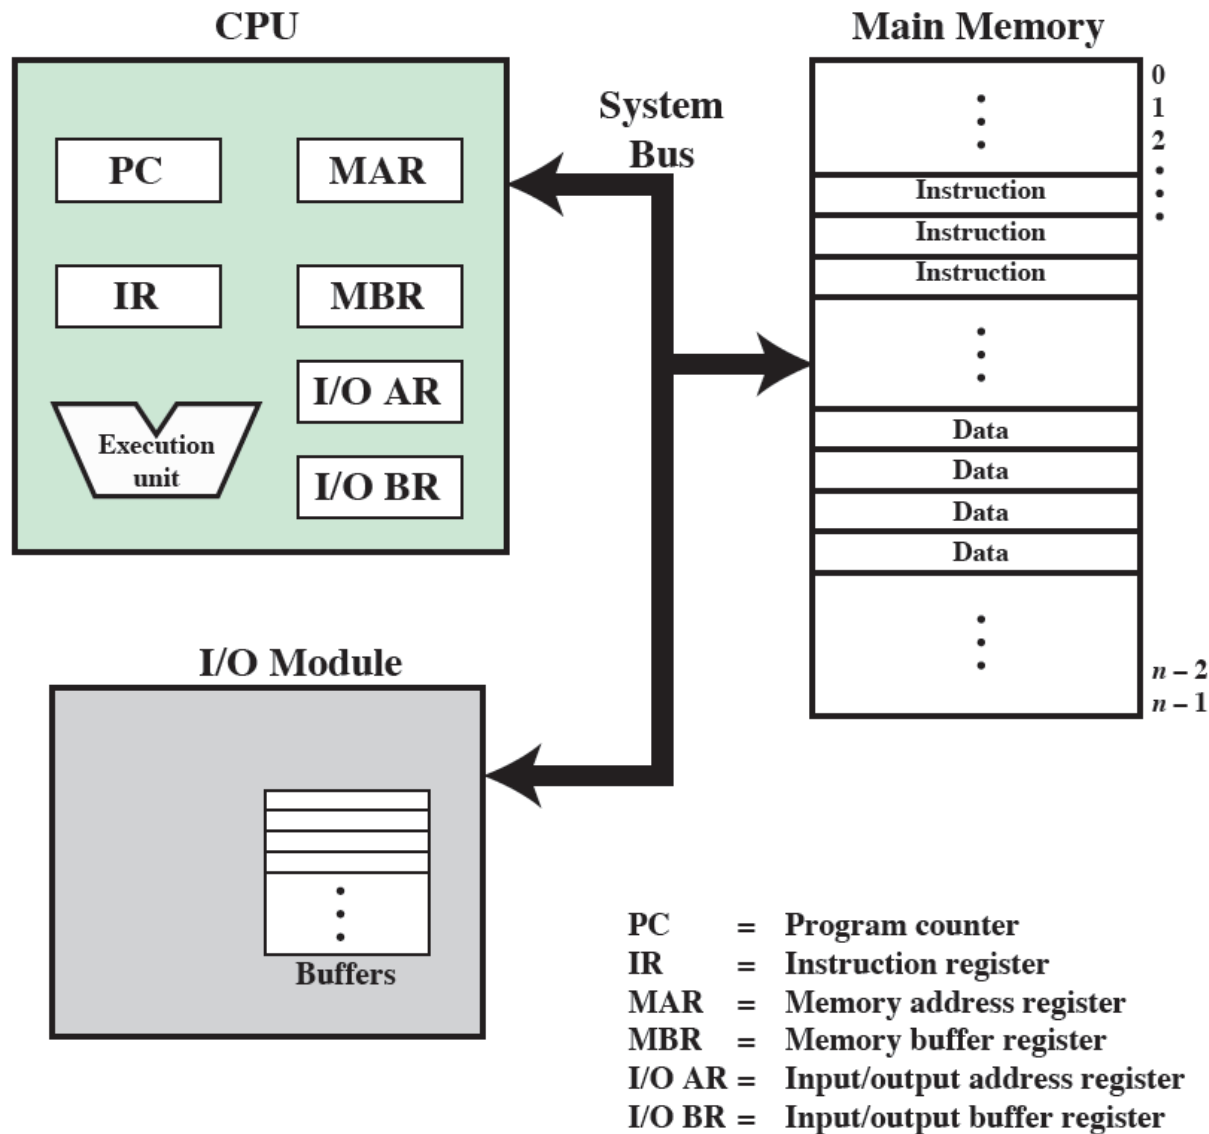
## ■ Major components

- CPU
  - Instruction interpreter
  - Module of general-purpose arithmetic and logic functions
  - I/O Components
- Input module
  - Contains basic components for accepting data and instructions and converting them into an internal form of signals usable by the system
- Output module
  - Means of reporting results

# Lets Build a Minimalistic Computer

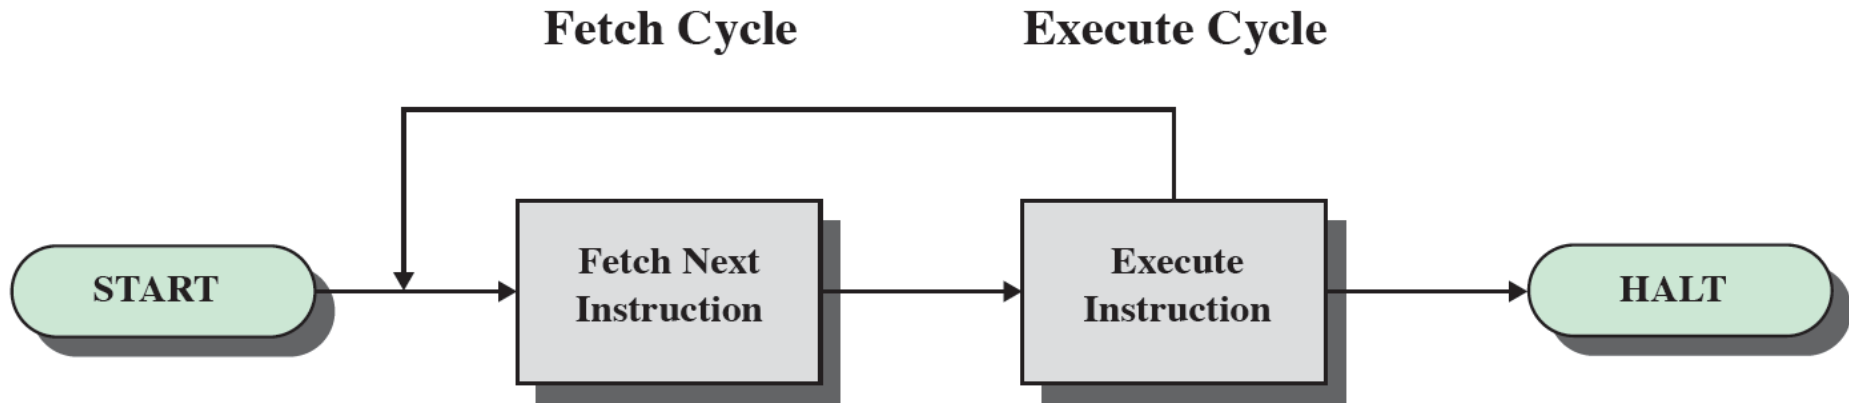
- First lets have a look at the CPU
- We definitely need
  - A **Program Counter (PC)** register pointing to the next instruction
  - A **Instruction Register (IR)** holding the current instruction
  - **Memory Address** and **Buffer** register (**MAR** and **MBR**) for fetching information from the memory
  - The same is required for I/O Operations (**I/O AR** and **I/O BR**)
- In the next step we will look how these components work together

# Our Minimalistic Computer





# How does Our Little Machine Work?



- The processing required for a single instruction is called an **instruction cycle**.
- In this simple form, a instruction cycle consists of:
  - Fetch Cycle
  - Execute Cycle

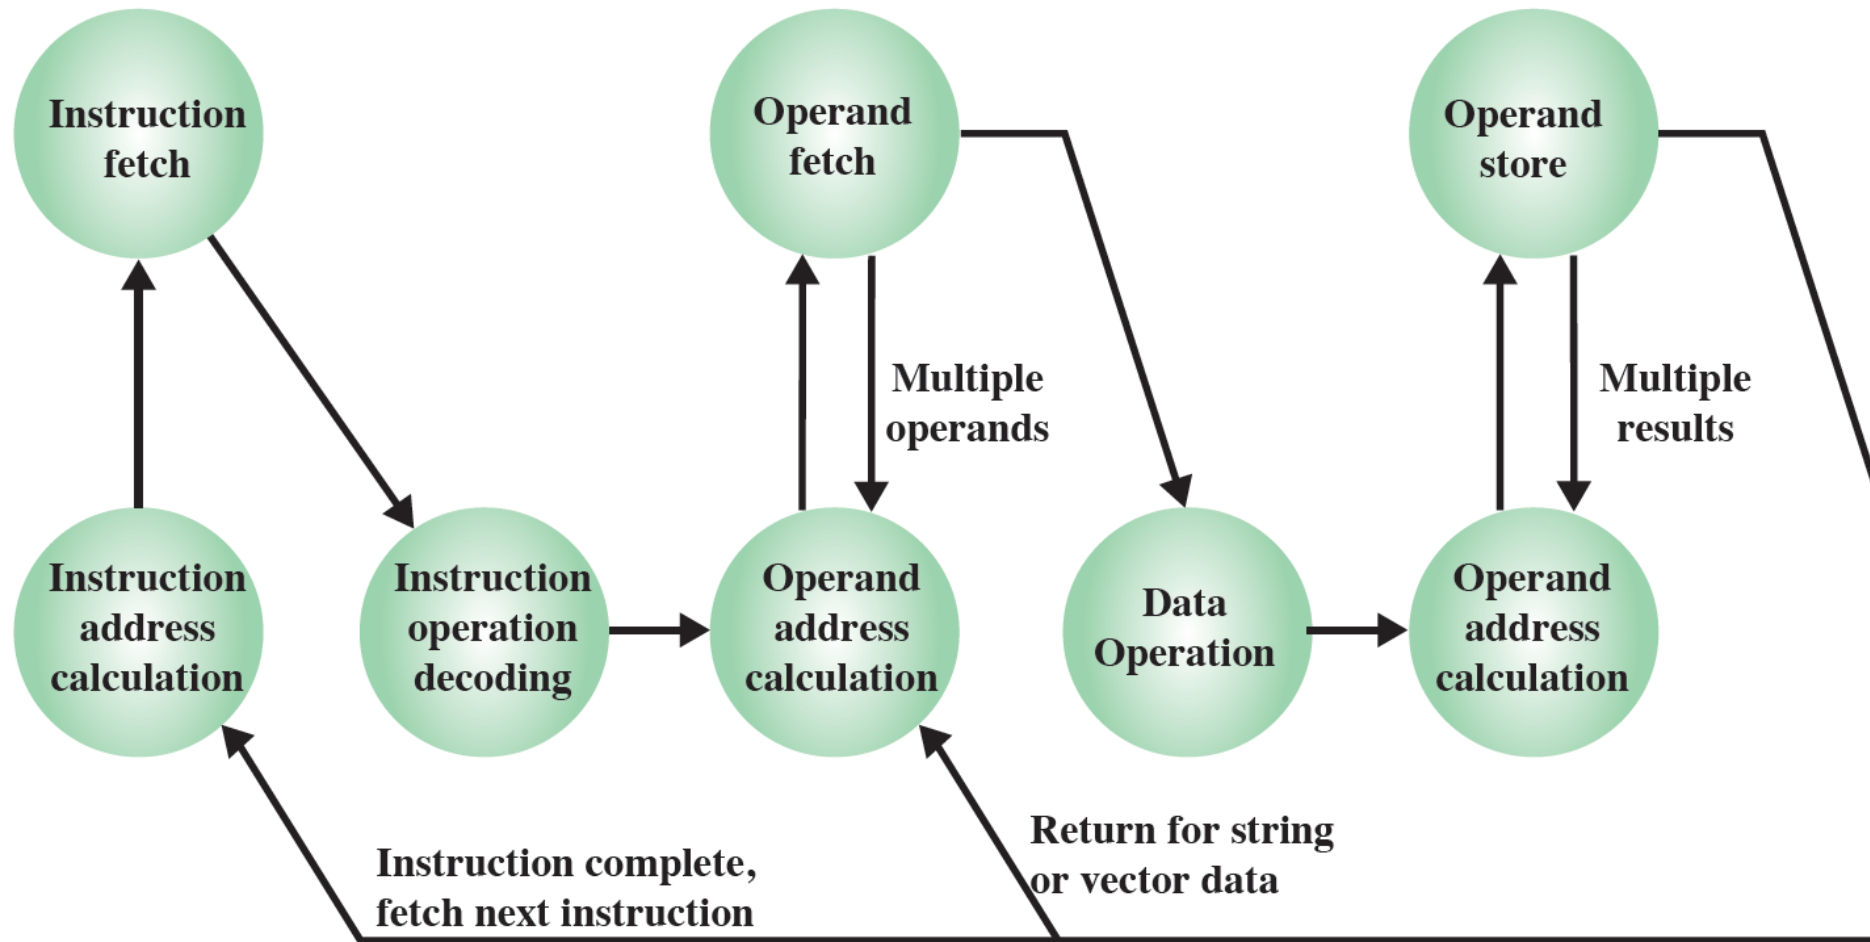
# The Fetch Cycle

- At the beginning of each instruction cycle the processor fetches an instruction from memory
- The program counter (PC) holds the address of the instruction to be fetched next
- The processor increments the PC after each instruction fetch
- The fetched instruction is loaded into the instruction register (IR)
- The processor interprets the instruction and performs the required action

# Are We Happy With That?

- Is this structure sufficient?
- Some thoughts:
  - Often, we are having several operands
  - What about indirect addressing
  - What about control flow changes
- These considerations require changes in our Instruction Cycle

# Advanced Instruction Cycle





# Top-Level View of a Computer

## Lecture Content

- Instruction Cycle
- **Interrupts**
- I/O Functionality
- Bus Systems
- QPI
- PCIe

# Further Mechanisms: Interrupts

- Virtually all computers provide a mechanism by which other modules may interrupt the normal processing
- We will hear more about interrupts later, for now we just need to know about the concepts
- We distinguish between various types of interrupts

# Types of Interrupts

## ■ Program

- Result of an instruction execution
  - arithmetic overflow, division by zero
  - attempt to execute an illegal machine instruction
  - reference outside a user's allowed memory space

## ■ Timer

- Generated by a timer within the processor. Used by the OS to perform certain functions on a regular basis, e.g., scheduling

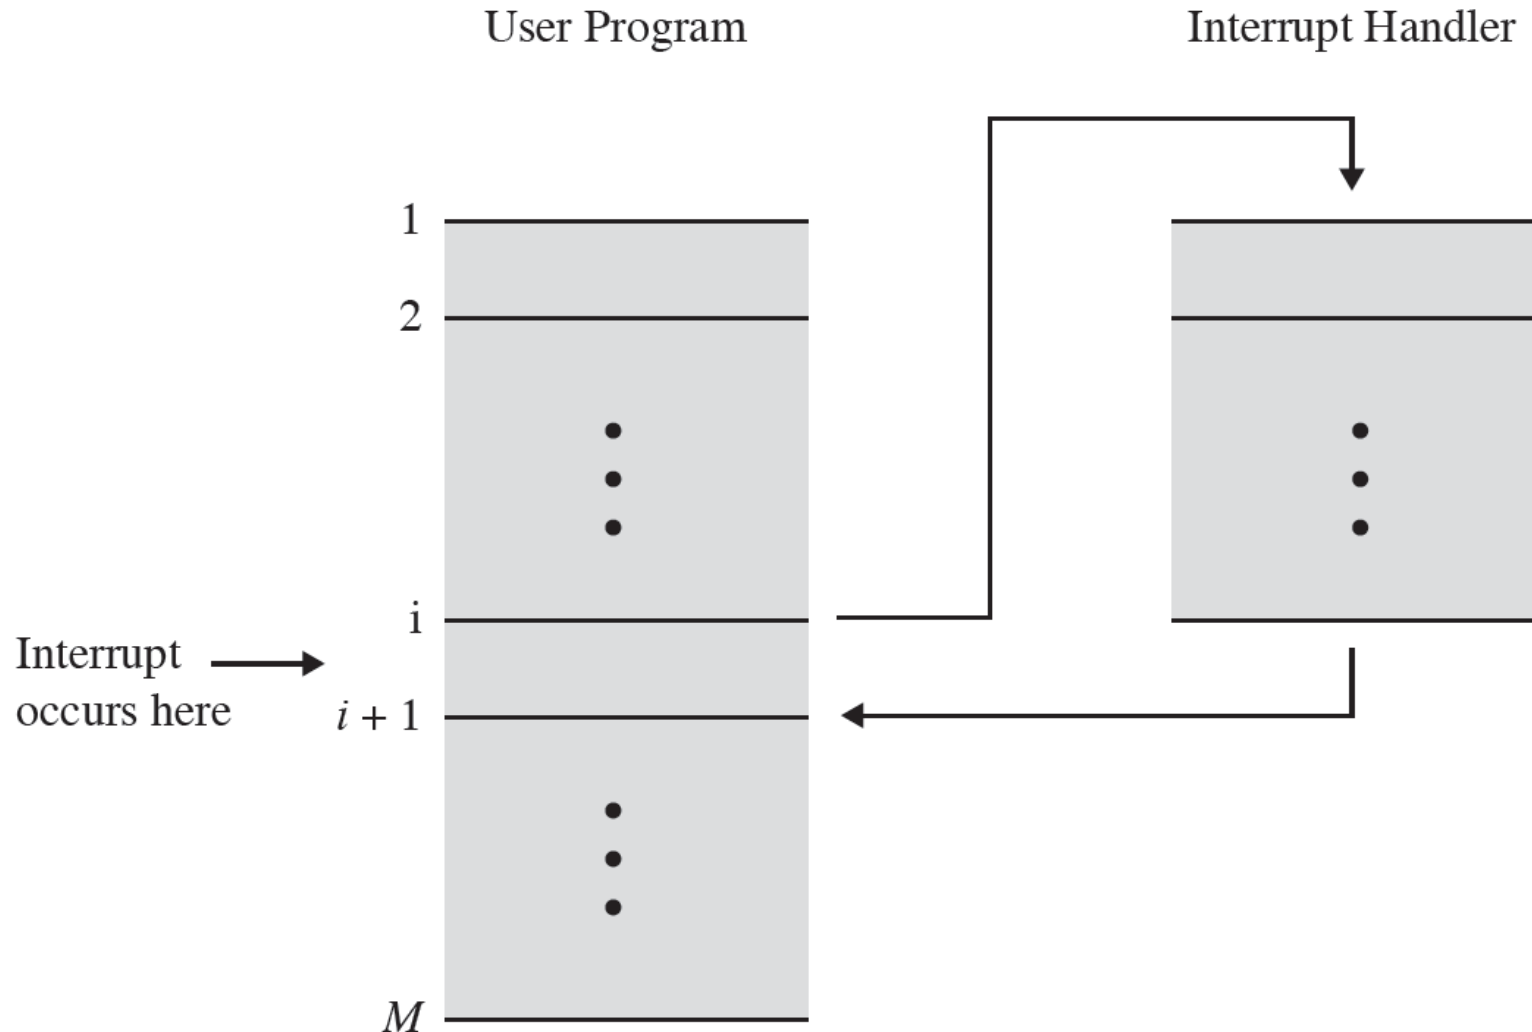
## ■ I/O

- Generated by an I/O controller
  - signal normal completion of an operation
  - variety of error conditions

## ■ Hardware failure

- Generated by a failure such as power failure or memory parity error.

# How Does an Interrupt Work?

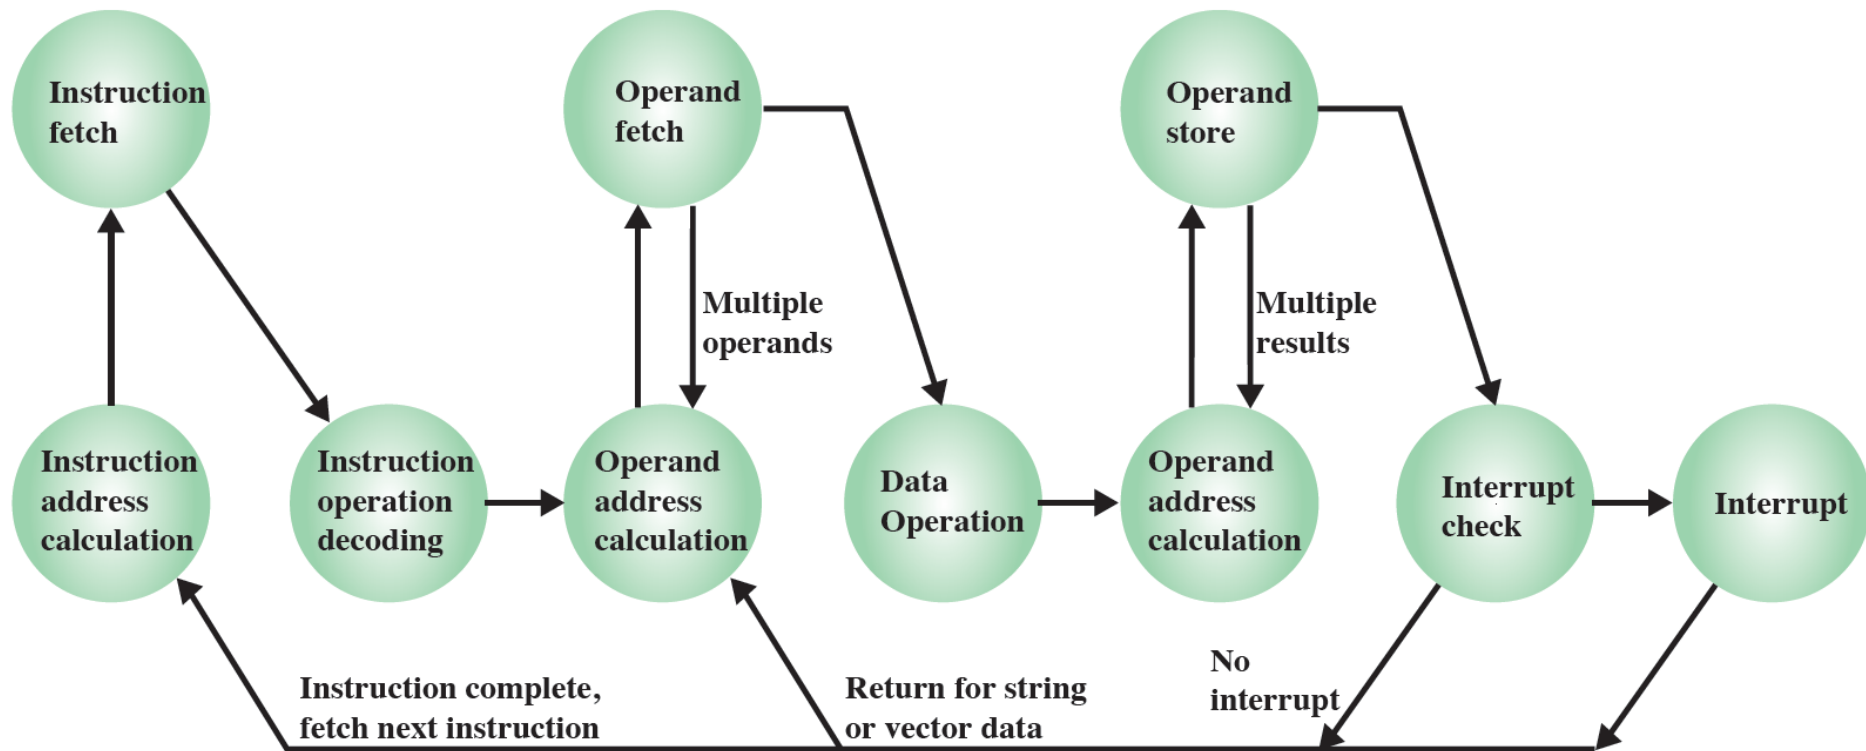





# Why Interrupts?

- Interrupts are provided primarily as a way to improve processing efficiency
- For example, most external devices are much slower than the processor
- Suppose that the processor always has to wait until it can finish the I/O operation and go to the next instruction
- Clearly, this is a very wasteful use of the processor.


# Instruction Cycle State Diagram With Interrupts




# Time Savings With Interrupts



# Two Ways of Handling Multiple Interrupts



(a) Sequential interrupt processing



(b) Nested interrupt processing


# Handling Multiple Interrupts: Sequential

- Interrupts have to wait until the first Interrupt was finished; thus, during one interrupts, others are **disabled**
- Doesn't consider time criticality of certain interrupts
- With a printer we don't care
- With communication lines indicating new data, we do
  - Buffers may overflow, no more data can be received
- Generally, we want the system to be responsive

# Prioritize Interrupts

- Normally, Interrupts are prioritized
- Interrupts with a higher priority may interrupt ongoing interrupts of lower priority
- Higher maintenance for the machine
- Small Example:
  - We have printer with priority 0 ( $t=10$ )
  - A hard disk with priority 4 ( $t=20$ )
  - A network with priority 5 ( $t=15$ )

# Handling Multiple Interrupts





# Top-Level View of a Computer

## Lecture Content

- Instruction Cycle
- Interrupts
- **I/O Functionality**
- Bus Systems
- QPI
- PCIe






# I/O Functionality

- I/O module can exchange data directly with the processor
- Processor can read data from or write data to an I/O module
  - This happens using special I/O instructions rather than memory referencing instructions
- In some cases it is desirable to allow I/O exchanges to occur directly with memory
  - The processor grants to an I/O module the authority to manipulate part of the memory
  - The I/O module issues read or write commands to memory relieving the processor of responsibility for the exchange
  - This operation is known as direct memory access (DMA)

# What Needs to be Interchanged?

- **Memory to processor:**
  - The processor reads an instruction or a unit of data from memory.
- **Processor to memory:**
  - The processor writes a unit of data to memory.
- **I/O to processor:**
  - The processor reads data from an I/O device via an I/O module.
- **Processor to I/O:**
  - The processor sends data to the I/O device.
- **I/O to or from memory:**
  - For these two cases, an I/O module is allowed to exchange data directly with memory (DMA)

# The Computer as a Set of Components





# Top-Level View of a Computer

## Lecture Content

- Instruction Cycle
- Interrupts
- I/O Functionality
- **Bus Systems**
- QPI
- PCIe

# Connecting these Modules: BUS

- A communication pathway connecting two or more devices
  - Key characteristic is that it is a shared transmission medium
- Signals from on device are visible to all devices attached to the bus
  - Two Senders: Garbage
- Typically consists of multiple communication lines
- Computers has several busses for various levels of the computer system hierarchy
- **System bus**
  - A bus that connects major computer components (processor, memory, I/O)


# Data Bus

- Data lines that provide a path for moving data among system modules
- May consist of 32, 64, 128, or more separate lines
- The number of lines is referred to as the width of the data bus
- This basically limits the amount of data which can be transferred at a time
  - The width of the data bus is a key factor in determining overall system performance
  - But there are more points to performance than the bus width

# Control Bus

- Because the data and address lines are shared by all components there must be a means of controlling their use
- Control signals transmit both command and timing information among system modules
- Timing signals indicate the validity of data and address information
- Command signals specify operations to be performed
- Typical control lines include:
  - Memory or I/O read/write
  - Transfer ACK
  - Bus request/grant
  - Interrupt request/ACK
  - Clock
  - Reset

# A Typical Bus Layout





# Central Elements of a Bus Design

## ■ Bus Line Types

### ■ **Dedicated:**

- Such a bus line is always assigned to one function or component
- E.g.: Permanent control line or address lines

### ■ **Multiplexed:**

- Fewer lines
- More complex, less parallelism

## ■ Method of Arbitration

### ■ **Centralized Arbitration:**

- A single hardware device (arbiter) is responsible for allocating time on the bus

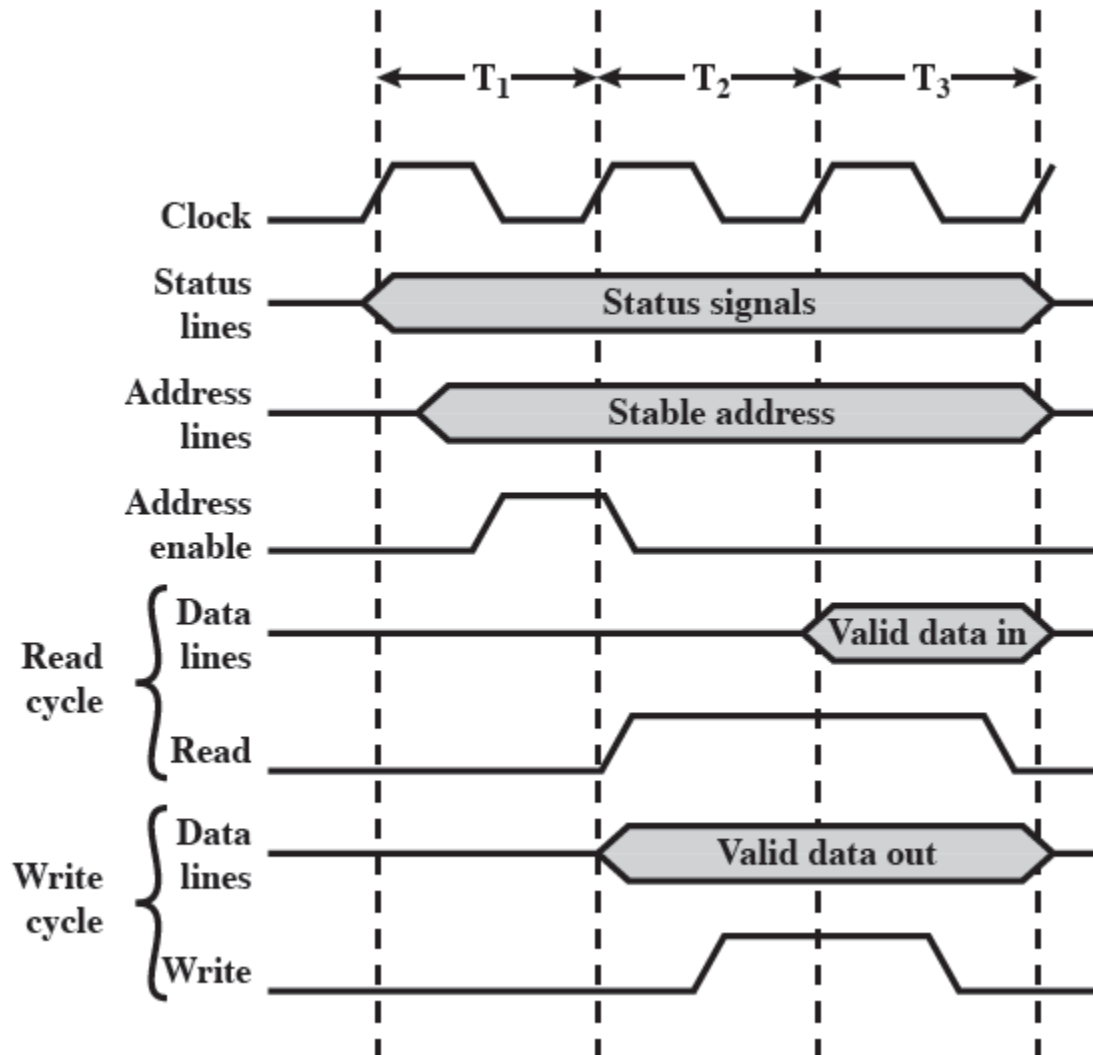
### ■ **Distributed Arbitration:**

- The use of the bus is negotiated between all modules

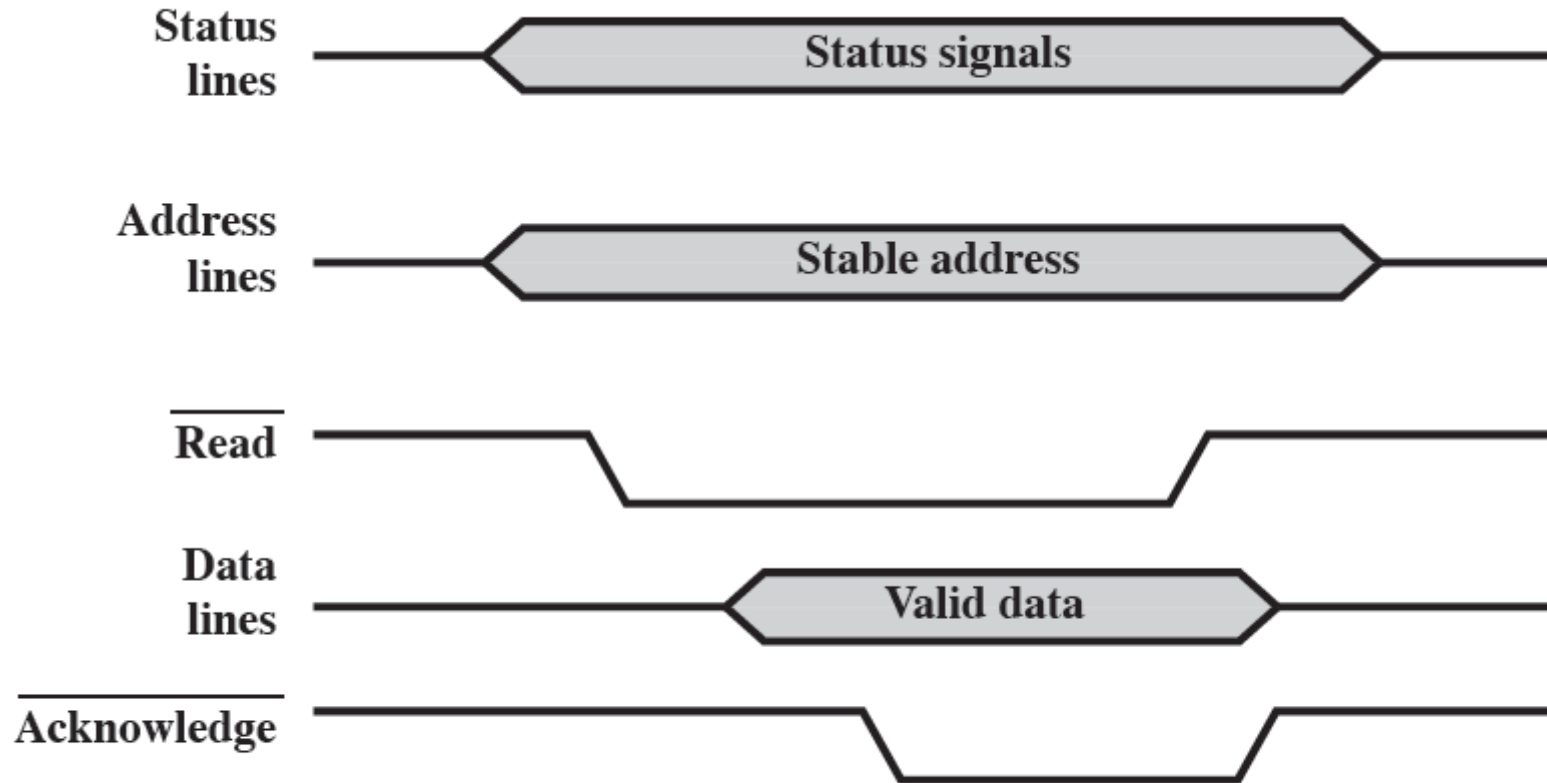
# Central Elements of a Bus Design

- **Bus width**
  - Address
  - Data
- **Data Transfer Types**
  - Read
  - Write
  - Read-modify-write
  - Read-after-write
  - Block

# Synchronous Bus

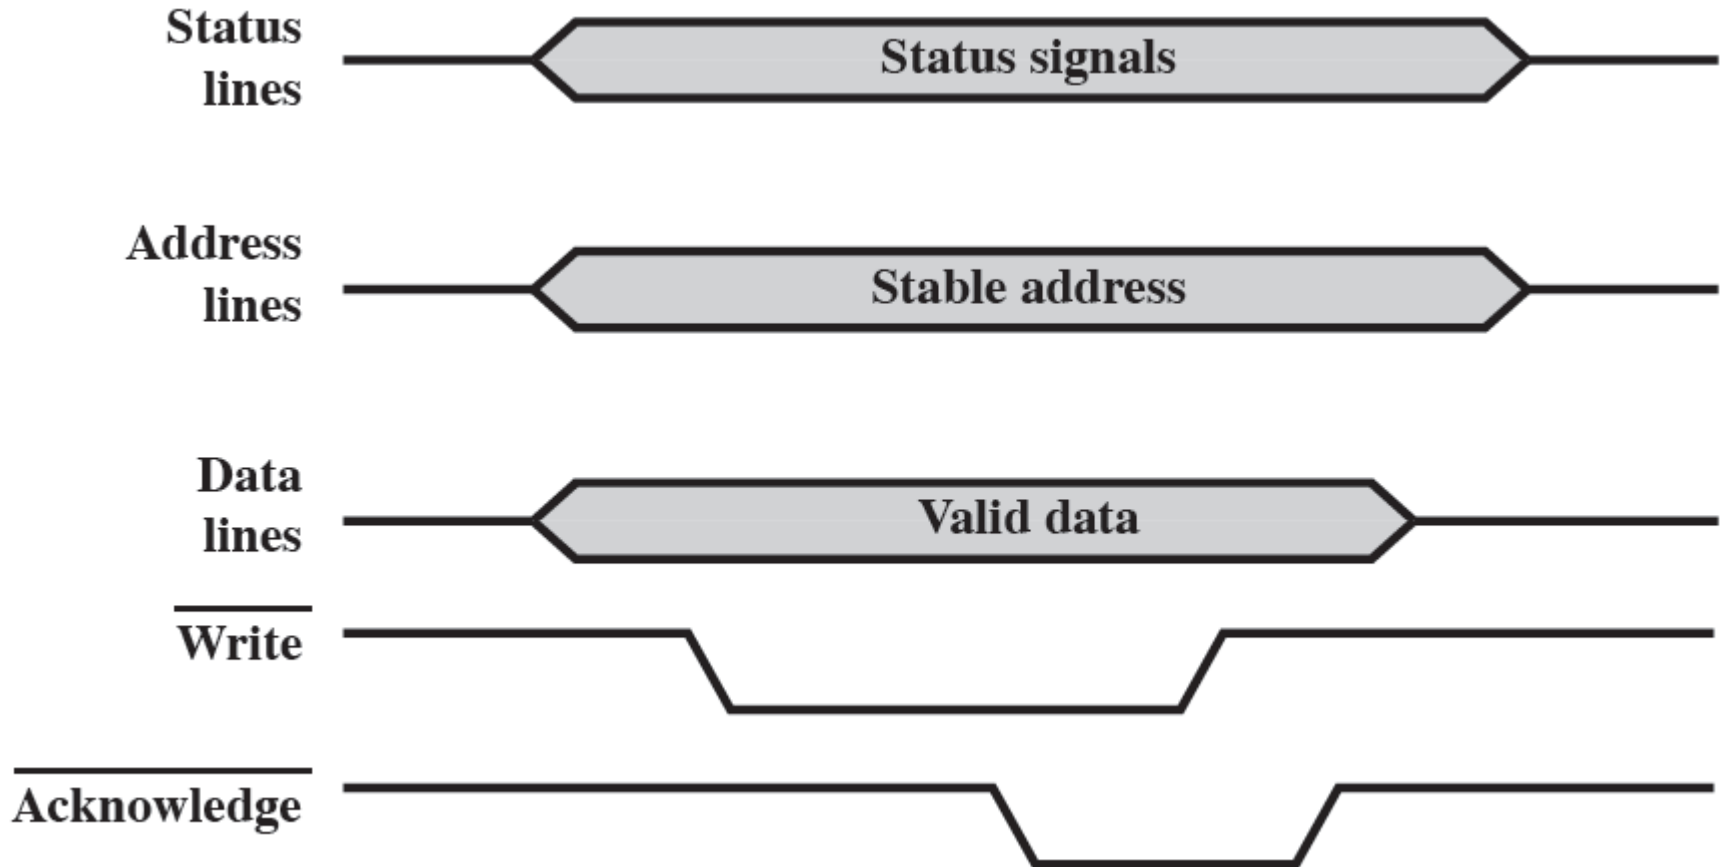


# Asynchronous Read



(a) System bus read cycle

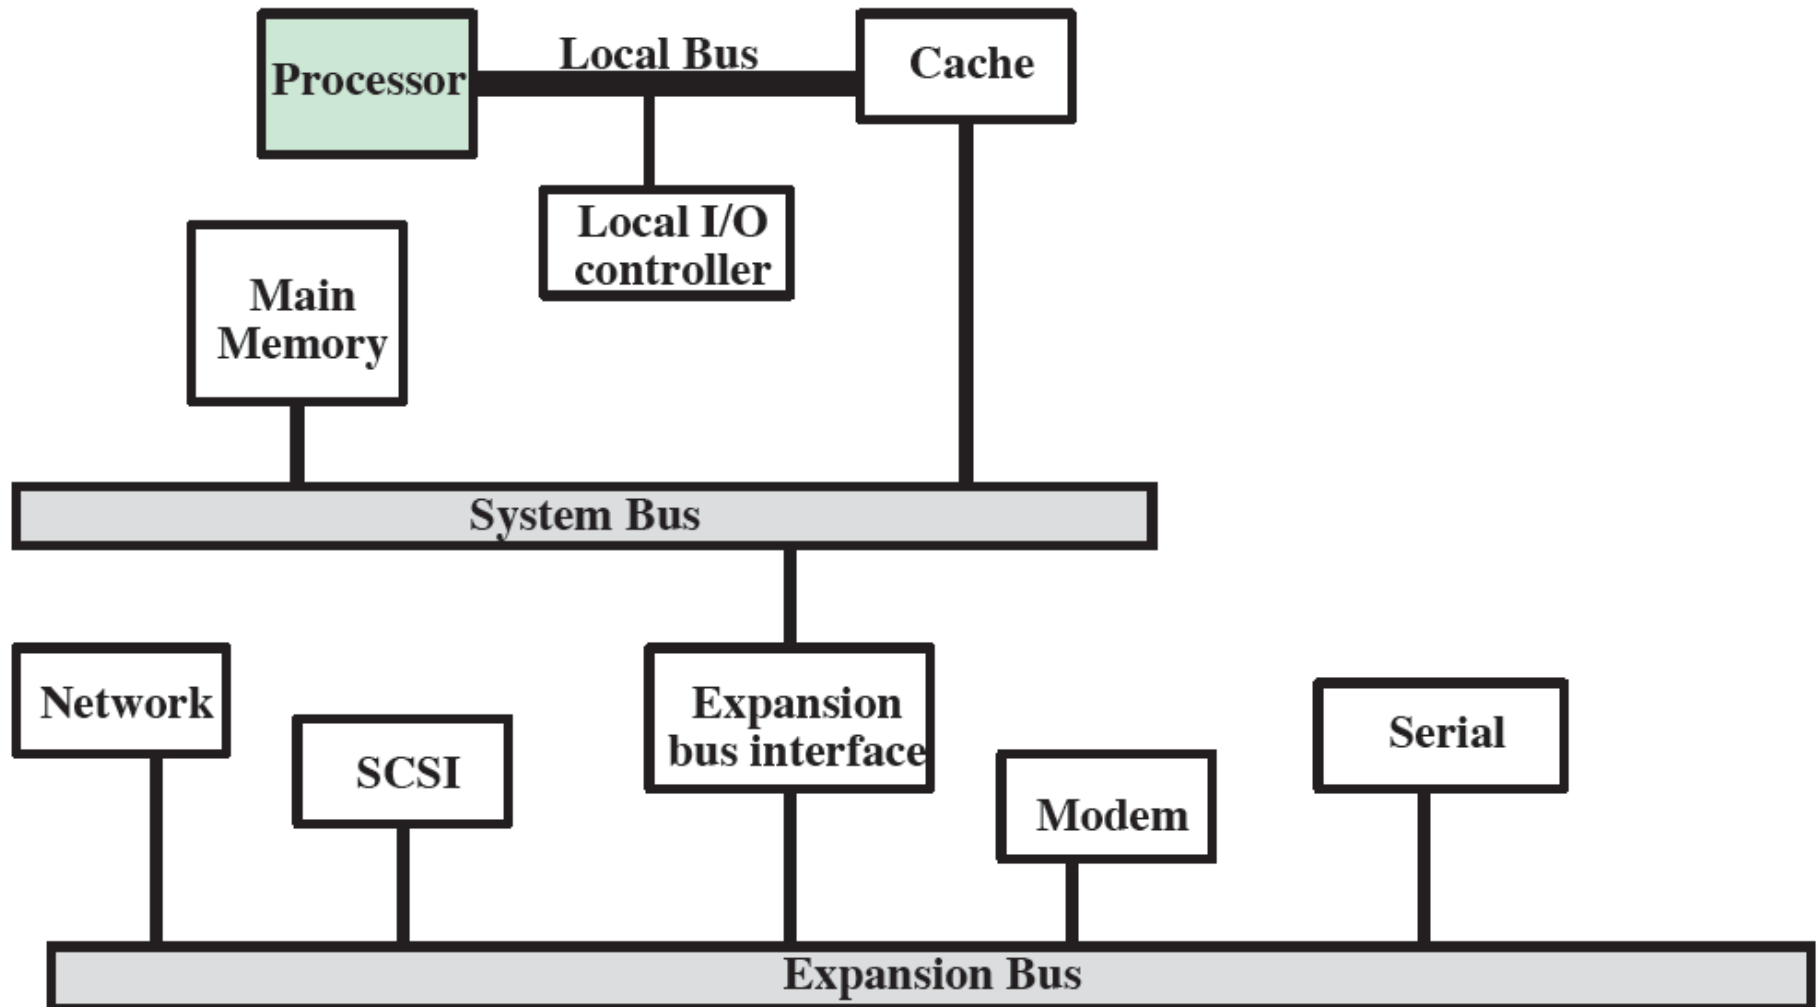
# Asynchronous Write



# Factors Limiting the Speed of a Bus

- If a great number of devices are connected to the bus, performance will suffer:
  - the more devices, the longer the bus (physically) -> greater latency
  - higher coordination overhead
  - The bus may become a bottleneck as the aggregate data transfer demand approaches the capacity of the bus.
- Counter measures:
  - Increase the data width
  - Increase clock speed of bus
- These measures won't ultimately solve that problem

# A Classical Hierarchical Bus Layout



# Using a Bus Hierarchy

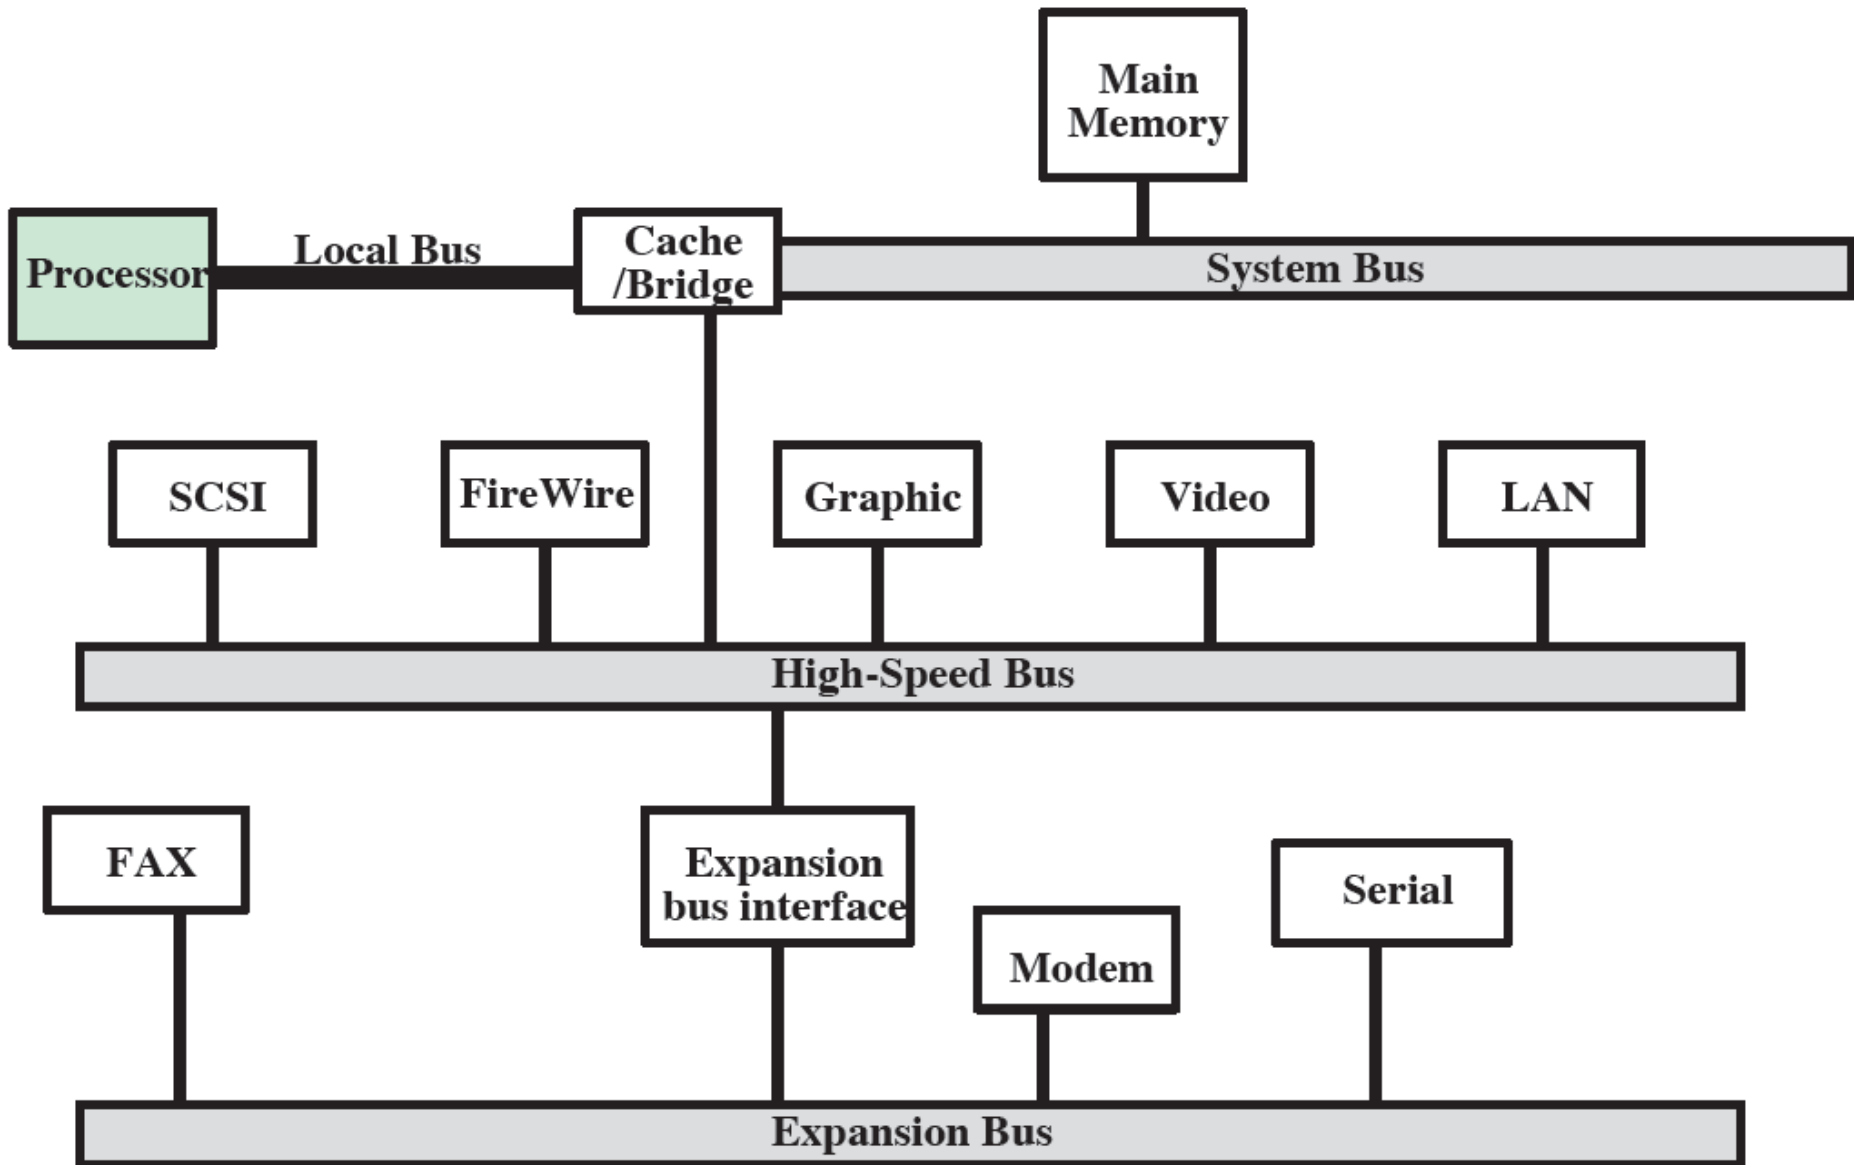
- In our Example:
  - Local Bus: Connecting processor to cache
  - Cache Memory Controller connects the cache to a system bus
  - System Bus: Connection to main memory and I/O
- The cache reduces the memory access which enabled the move of the memory to the System Bus
- This in turn enables the direct communication of I/O devices over the system bus without interfering with the processor
- Normally, I/O devices are not directly connected to system bus but to a secondary bus, the Expansion bus interface



# Using a Bus Hierarchy

- This set-up is also reaching its boundaries
- Due to the ever increasing demand in I/O performance, a newer approach is taken:
  - So called mezzanine architecture
  - A high speed bus tightly integrated into the system
  - The cache is integrated into a bridge or buffering device which connects the processor's bus and the high-speed bus is required

# Nowadays Design



# Advantages of the Mezzanine Architecture

- High-speed bus brings high demand devices into closer integration with the processor
- At the same time it is independent of the processor
- Thus, differences in processor and high-speed bus speeds and signal line definitions are tolerated
- Changes in processor architecture do not affect the high-speed bus, and vice versa



# Top-Level View of a Computer

## Lecture Content

- Instruction Cycle
- Interrupts
- I/O Functionality
- Bus Systems
- **QPI**
- PCIe

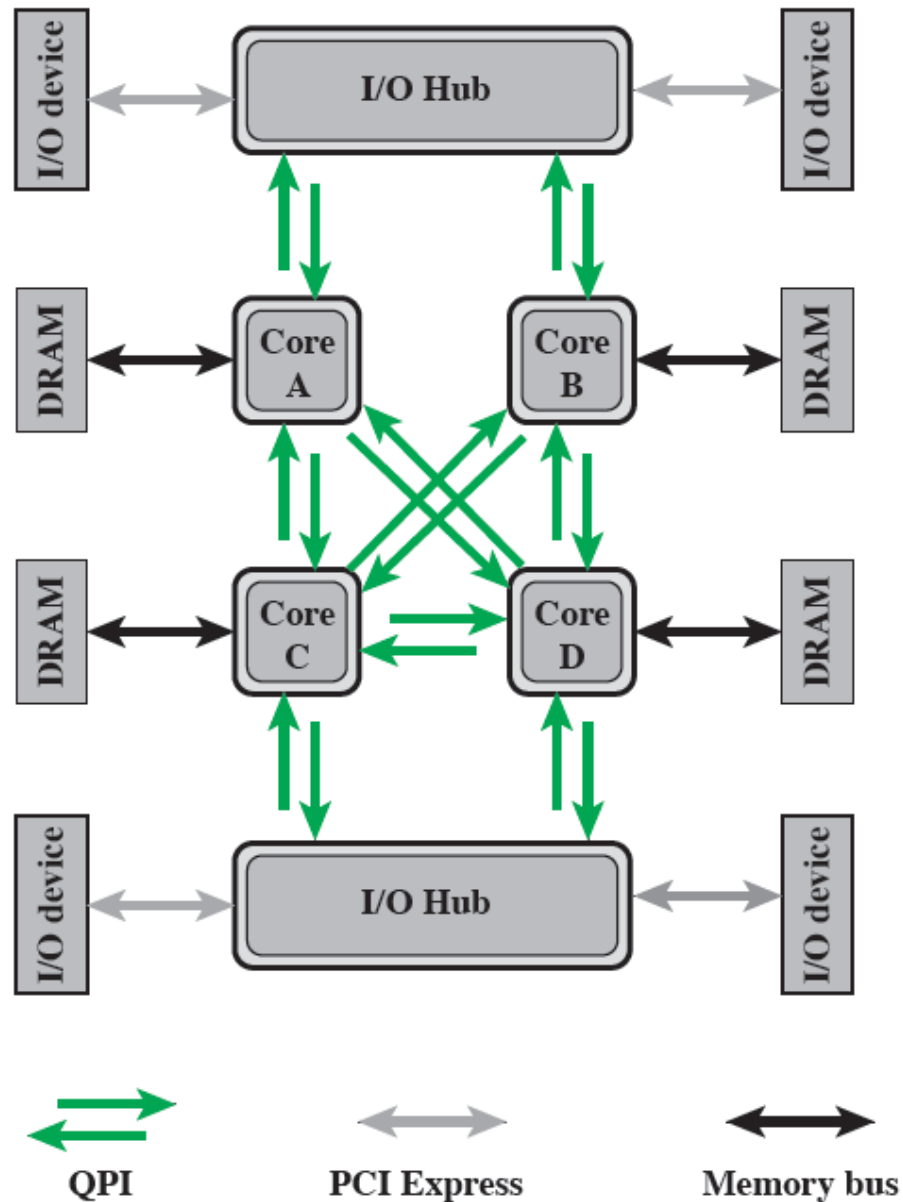
# Point to Point Interconnect

- Wide synchronous buses came to a limit with the ever increasing clock speed
- At higher and higher data rates it becomes increasingly difficult to perform the synchronization and arbitration functions in a timely fashion
- Especially a conventional shared bus on the same chip magnified the difficulties
- Solution: Point to Point Interconnect. Has lower latency, higher data rate, and better scalability

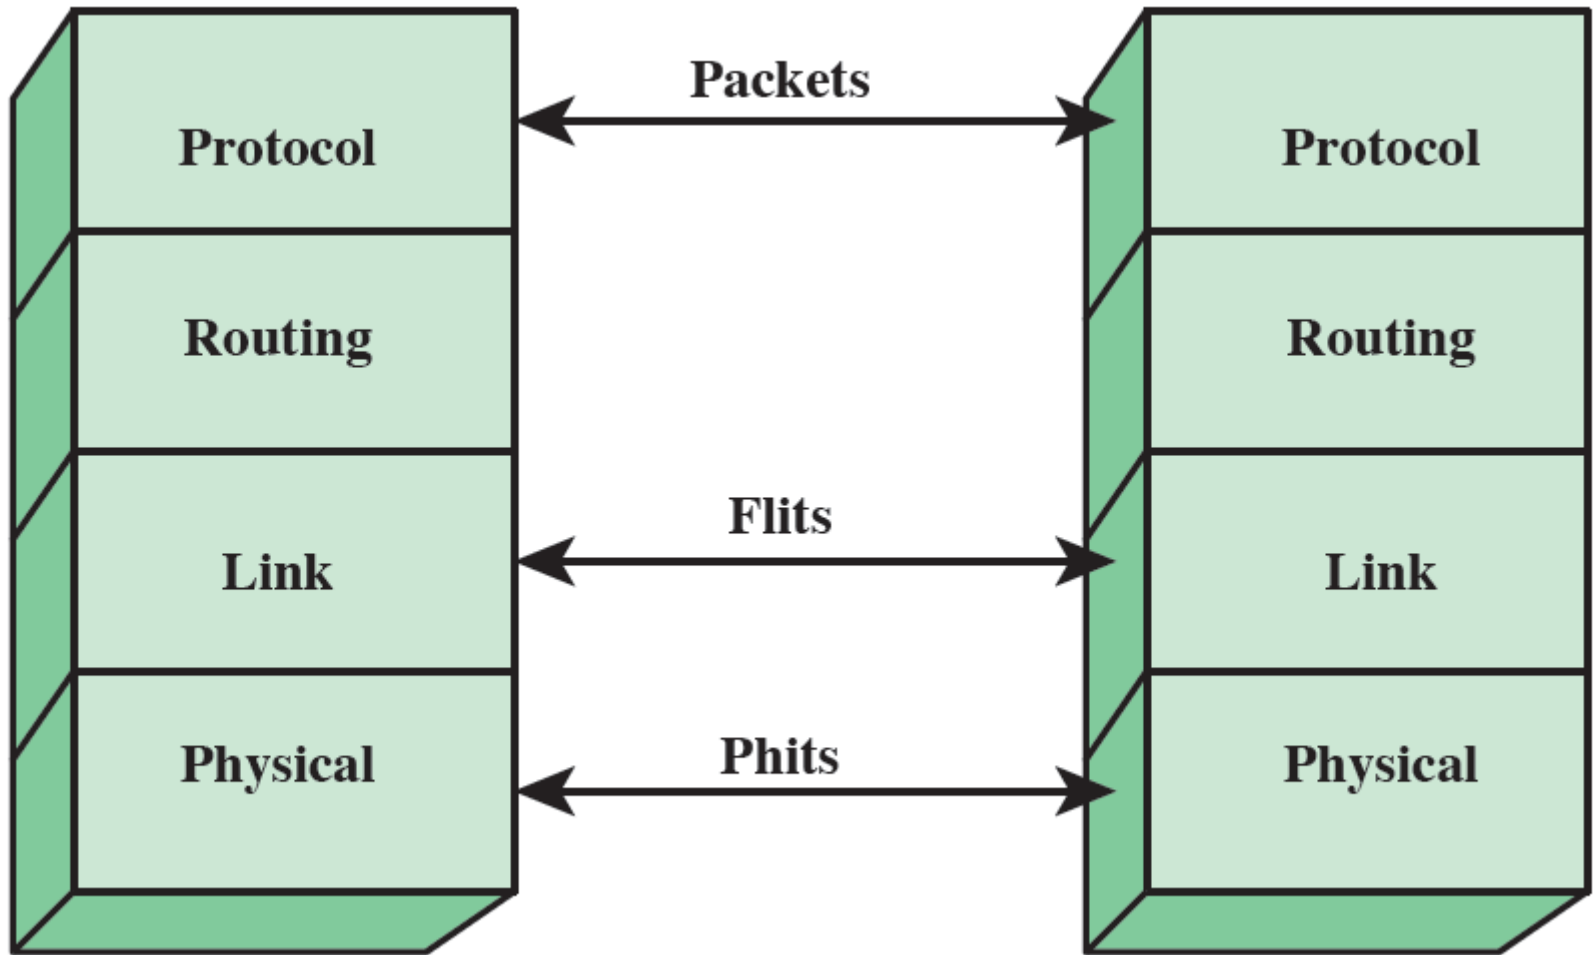
# Important Example: Quick Path from Intel

- Introduced in 2008, replaced the Front Side Bus (FSB)
- Direct pairwise connections to other components
  - No need for arbitration
- Layered protocol architecture
  - In contrast to simple use of control signals found in shared bus
  - 20 Bit physical (**phits**), 80 Bit logical (**flits**)
- Packetized data transfer
  - Data are sent as a sequence of packets each of which includes control headers and error control codes

# Example with 4 Cores

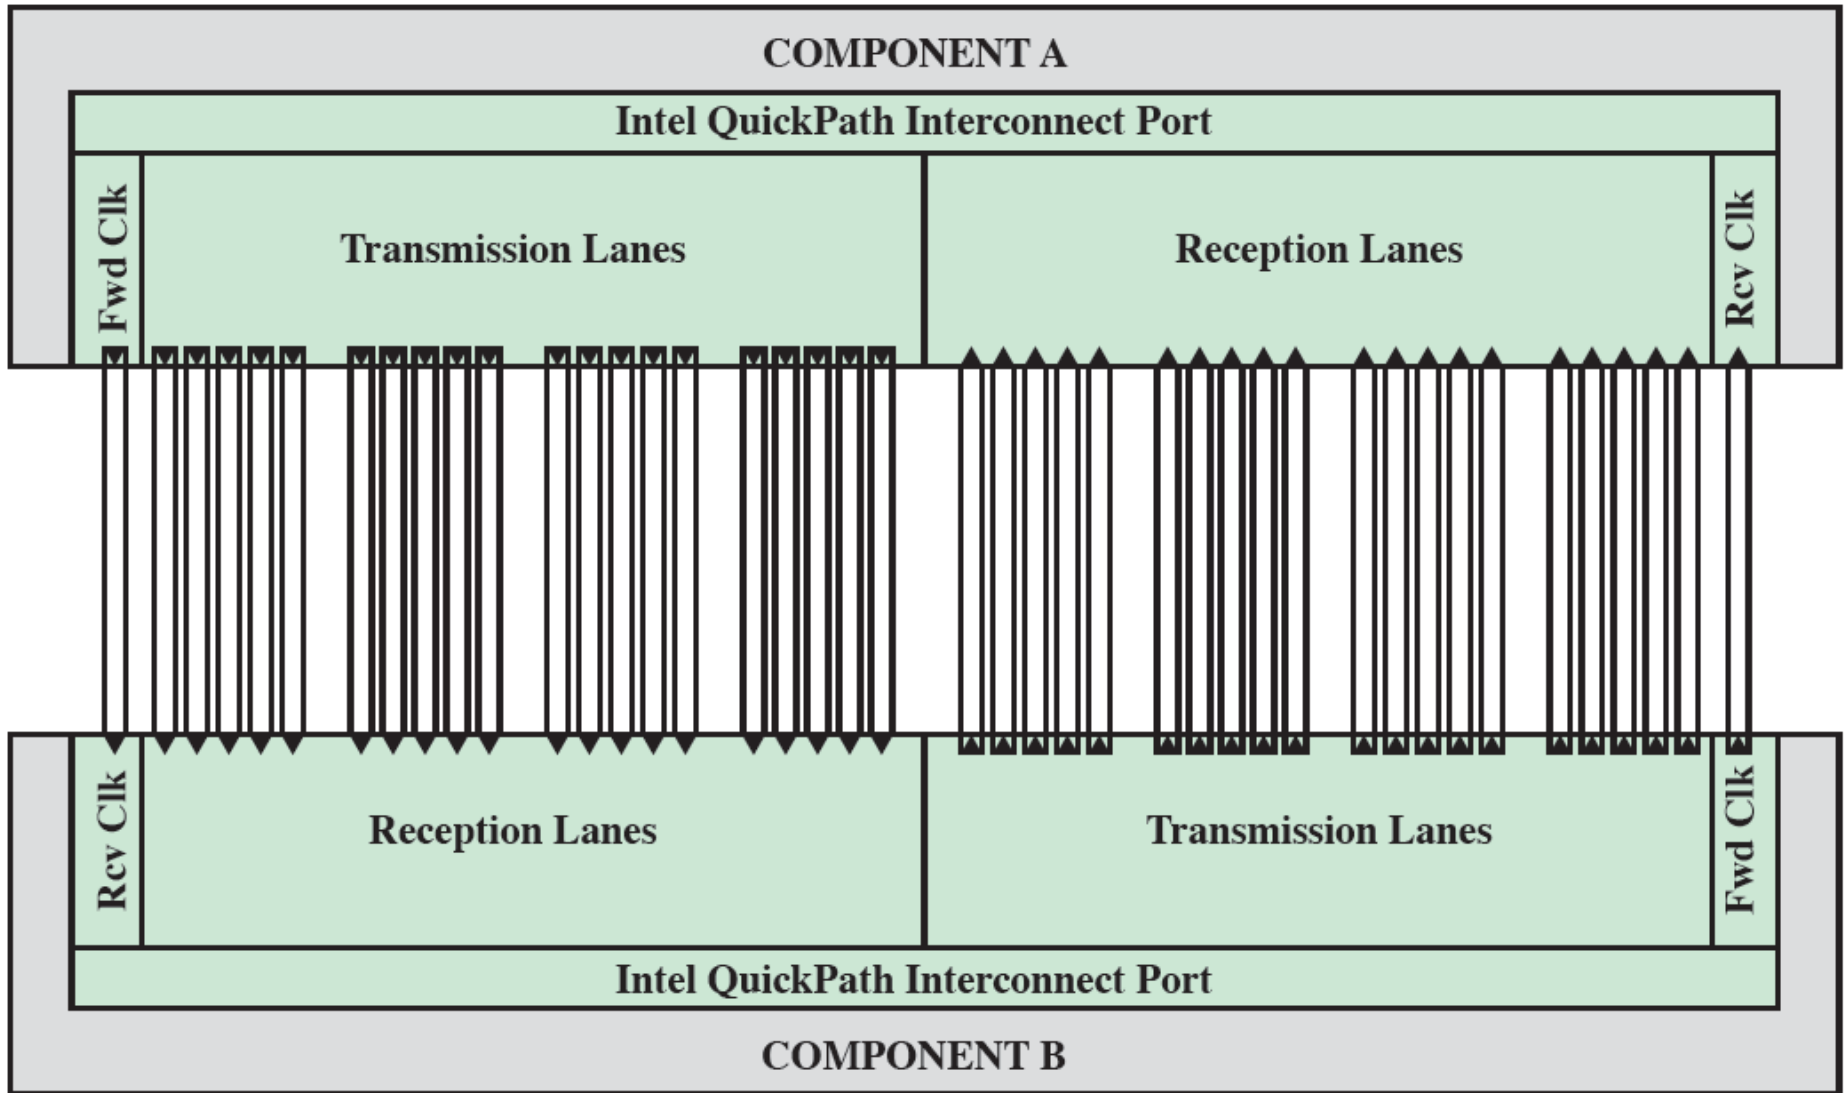


# The QPI Layers





# QPI Physical Connections



# QPI Facts

- 84 individual links; each data path consists of a pair of wires, called a lane.
- Operates at 6.4 GT/s. At 20 bits per transfer, that adds up to 16 GB/s in one direction, the total capacity is 32 GB/s.
- The form of transmission on each lane is known as low-voltage differential signaling, or balanced transmission.
  - The signal is transmitted by small voltage differences
  - The receiver consists of a resistor
  - The polarity of the current going through the resistor defines 0 or 1
  - Very low electro magnetic noise

# The QPI Link Layer

- Performs two key functions: **flow control** and **error control**
  - Operate on the level of the flit (flow control unit)
  - Each flit consists of a 72-bit message payload and an 8-bit error control code called a *cyclic redundancy check* (CRC)
- Flow control function
  - In order to not overwhelm the receiver
  - Credit based system
- Error control function
  - Detects and recovers from bit errors, and so isolates higher layers from experiencing bit errors
  - Requires buffering of the sender (in case something needs to be resent)

# QPI Routing Layer

- Used to determine the course that a packet will traverse across the available system interconnects
- Defined by firmware and describe the possible paths that a packet can follow
- The routing table options may be quite complex, depending on how
  - devices are populated
  - system resources are partitioned
  - reliability events result in mapping around a failing resource

# Protocol Layer

- Packet is defined as the unit of transfer
- One key function performed at this level is a cache coherency protocol which deals with making sure that main memory values held in multiple caches are consistent
- A typical data packet payload is a block of data being sent to or from a cache
- Contains 8 Bit for routing information
  - That means one 80 Bit flit contains 64 bit payload, 8 bit routing and 8 bit CRC



# Top-Level View of a Computer

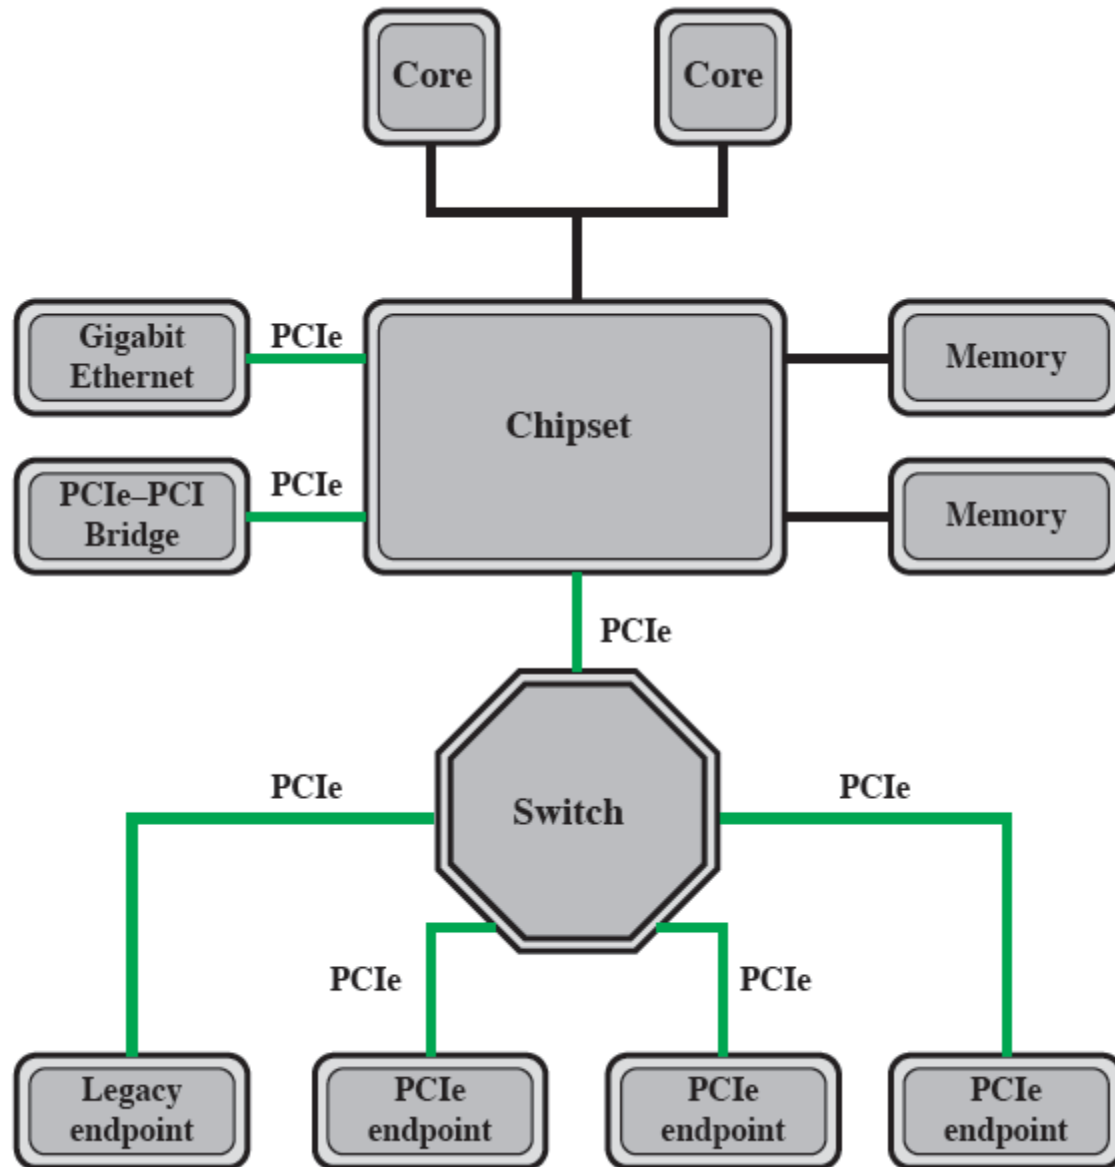
## Lecture Content

- Instruction Cycle
- Interrupts
- I/O Functionality
- Bus Systems
- QPI
- **PCIe**

# Peripheral Component Interconnect (PCI)

- A popular high bandwidth, processor independent bus
- Delivers better system performance for high speed I/O subsystems
- PCI Special Interest Group (SIG)
  - First developed by Intel
  - Widely adopted
- PCI Express (PCIe)
  - Point-to-point interconnect scheme intended to replace bus-based schemes such as PCI
  - Key requirement is high capacity to support the needs of higher data rate I/O devices, such as Gigabit Ethernet
  - Another requirement deals with the need to support time dependent data streams

# PCIe Logical Architecture






# PCIe Configuration

- **Root Complex device**
  - also referred to as a chipset or a host bridge
  - connects the processor and memory subsystem to the PCI Express switch
  - Acts as buffering device to deal with the difference in the data rate
- The chipset will typically support multiple PCIe ports
  - some of which attached directly to a PCIe device
  - some are attached to a switch that manages multiple PCIe streams.
- Legacy endpoints
  - Allows for existing designs to be ported to PCIe
- PCIe/PCI bridge

# PCIe also has a Protocol



# PCIe Protocol Layers

## ■ Physical Layer

- The actual wires carrying the signals
- Circuitry and logic required in the transmission and receipt of data

## ■ Data Link Layer

- Responsible for reliable transmission and flow control
- Data packets generated and consumed by the DLL are called Data Link Layer Packets (DLLPs).


## ■ Transaction Layer

- Generates and consumes data packets
- Manages the flow control of those
- Data packets generated and consumed by the TL are called Transaction Layer Packets (TLPs).

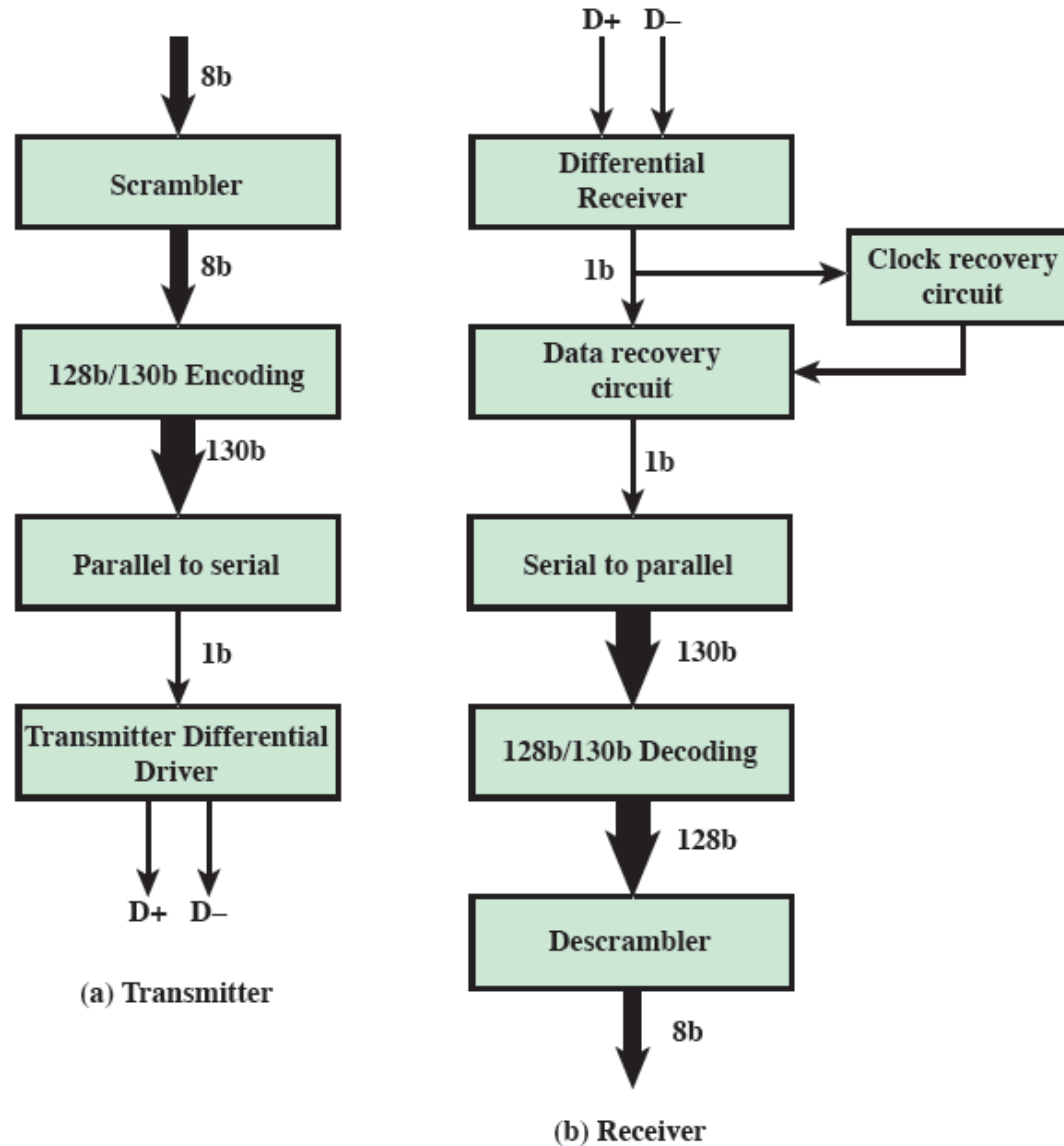
# Details on the Physical Layer

- 1, 4, 6, 16, or 32 bidirectional lanes per port, again differential signaling
- As with QPI, PCIe uses a multilane distribution technique.
- At each physical lane, data are buffered and processed
  - Blocks of 16 bytes (128 bits) at a time
  - Each block is encoded into a unique 130-bit codeword for transmission
  - 128b/130b encoding.
- 128b/130b
  - PCIe has no clock. The two additional bytes are used to synchronize

# PCIe Multilane Distribution



# PCIe Transmit and Receive Block Diagram



(a) Transmitter

(b) Receiver

# PCIe Transaction Layer

- Receives read and write requests from the software above the TL and creates request packets for transmission to a destination via the link layer
- Most transactions use a split transaction technique
  - A request packet is sent out by a source PCIe device which then waits for a response called a completion packet
- TL messages and some write transactions are posted transactions (meaning that no response is expected)
- TL packet format supports 32-bit memory addressing and extended 64-bit memory addressing

# PCIe Transaction Layer: Address Spaces

- **Memory**

- The memory space includes system main memory. It also includes PCIe I/O devices. Certain ranges of memory addresses map into I/O devices.

- **I/O**

- This address space is used for legacy PCI devices, with reserved memory address ranges used to address legacy I/O devices.

- **Configuration**


- This address space enables the TL to read/write configuration registers associated with I/O devices.

- **Message**


- This address space is for control signals related to interrupts, error handling, and power management.



# PCIe Protocol Data Unit Format



(a) Transaction Layer Packet



(b) Data Link Layer Packet

# TLP Memory Request Format

