



Operating System Support

Lecture Content

- **OS in General**
- **Memory Management**
- **Virtual Memory**
- **Segmentation**
- **Memory Management of the Pentium II**

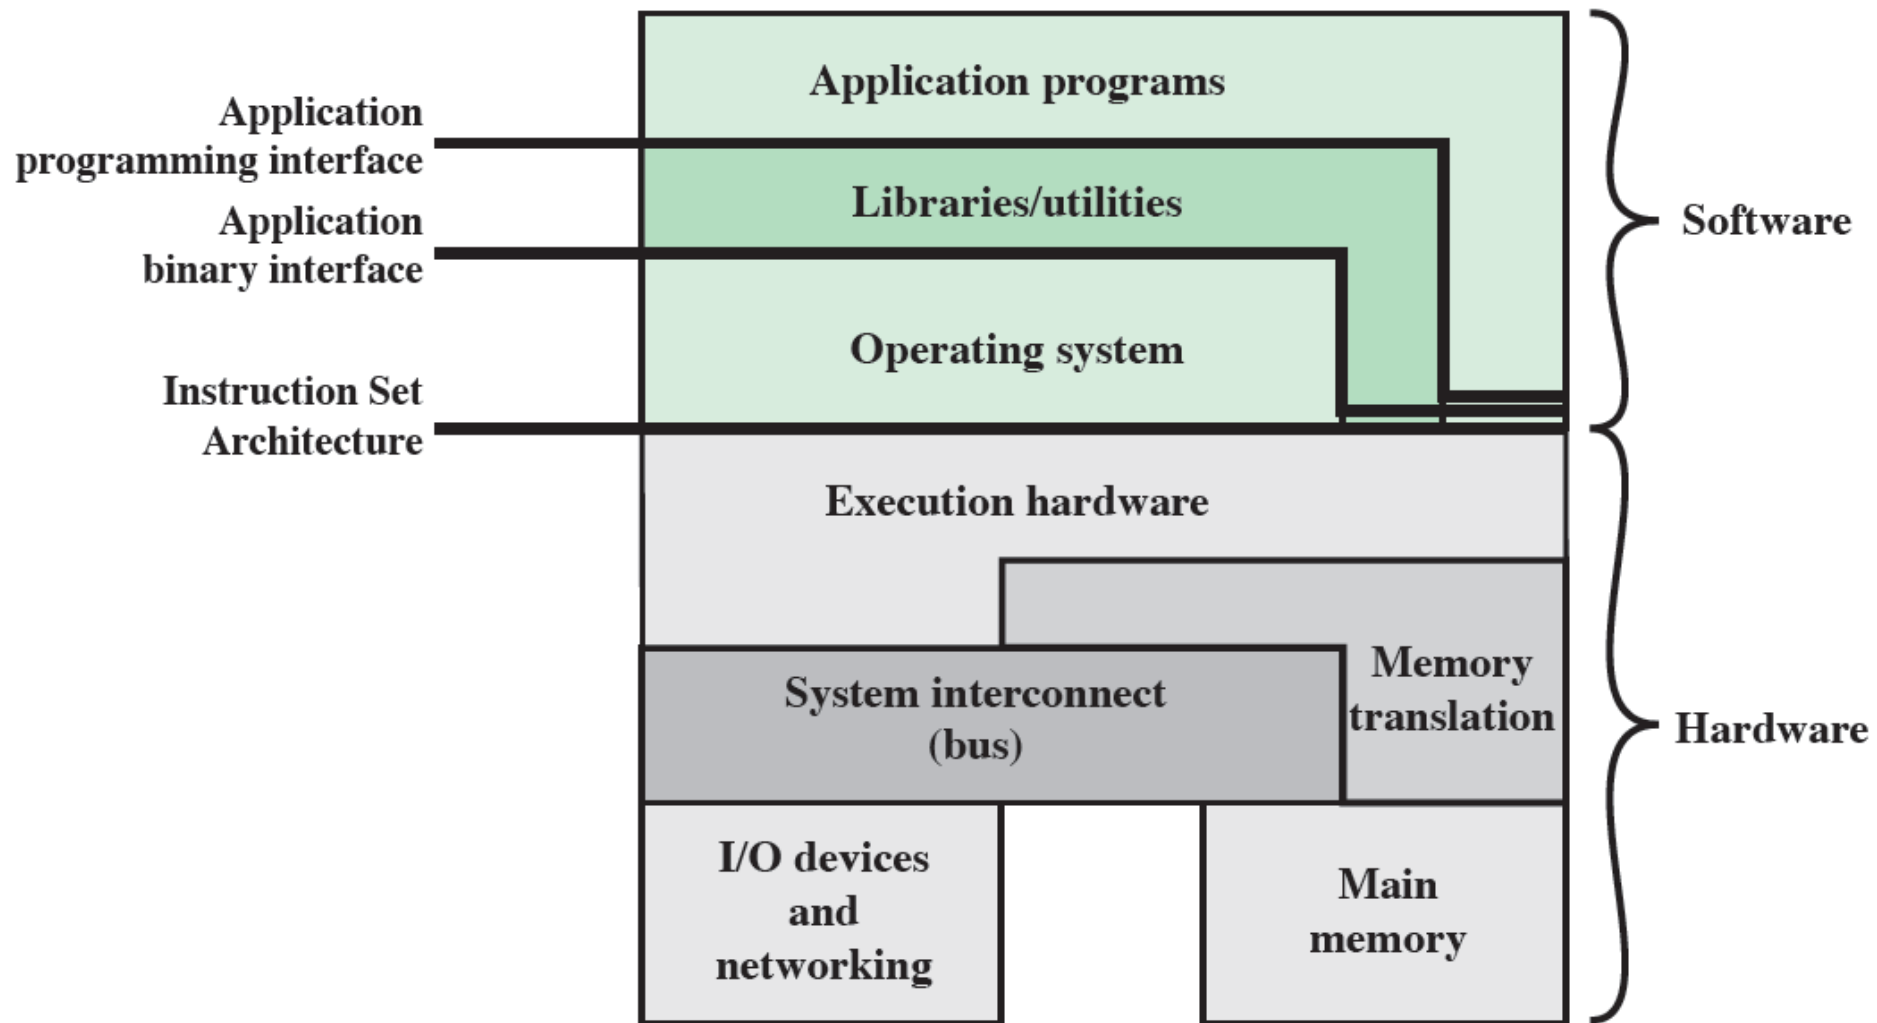


Operating System Support

Learning Objectives

- Understand the OS from the point of view of the CPU
- Understand the concepts and necessity of scheduling
- Understand paging and segmentation
- Understand the concepts of virtual memory

Computer Hard and Software Structure



Key Interfaces For Computers

- **Instruction set architecture (ISA)**
 - Defines the machine language instructions of a computer
 - Boundary between hardware and software
- **Application binary interface (ABI)**
 - Defines a standard for binary portability across programs
 - Defines the system call interface to the operating system
- **Application programming interface (API)**
 - Gives a program access to the hardware resources and services available in a system through the user ISA supplemented with high-level language (HLL) library calls
 - Using an API enables application software to be ported easily to other systems that support the same API

Small Overview over the x86_64 Linux ABI

- Functions Calls:
 - Integer Parameters: %rdi, %rsi, %rdx, %rcx, %r8, %r9
 - Float Parameters: %xmm0-%xmm7
 - Further Parameters on the stack
 - Return Values: %rax, %rdx, %xmm0, %xmm1
 - Callee registers: %rsp, %rbx, %r12-%r15
- Memory and Stack Design
- Recommended reading:
<http://www.x86-64.org/documentation/abi.pdf>

Operating System as Resource Manager

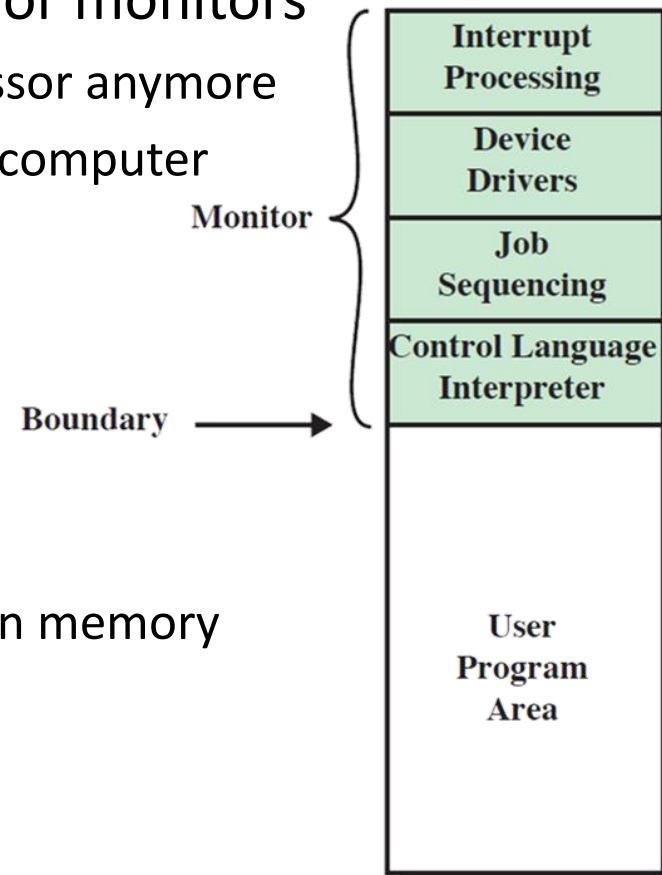
- A computer is a set of resources for the movement, storage, and processing of data and for the control of these functions
- The OS is responsible for managing these resources
- The OS as a control mechanism is unusual in two respects:
 - The OS functions in the same way as ordinary computer software – it is a program executed by the processor
 - The OS frequently relinquishes control and must depend on the processor to allow it to regain control

Back then ...

- From the late 1940s to the mid-1950s the programmer interacted directly with the computer hardware – there was no OS
- Problems:
 - **Scheduling**
 - Sign-up sheets were used to reserve processor time
 - This could result in wasted computer idle time if the user finished early
 - If problems occurred the user could be forced to stop before resolving the problem
 - **Setup Time**
 - A single program could involve:
 - Loading the compiler plus the source program into memory
 - Saving the compiled program
 - Loading and linking together the object program and common functions

Batch Operating Systems

- The wasted time due to scheduling and setup time was unacceptable
- Advent of simple batch operating systems or monitors
 - Programmer has no direct access to the processor anymore
 - The user submits the job on cards or tape to a computer operator
 - Put on an input device, for use by the monitor
- Point of View of the Monitor
 - Controls the sequence of events
 - Thus, most of the monitor has to be in the main memory
 - This is called **resident monitor**



Point of View of the Processor

- Processor executes instructions from the monitor
 - These instructions cause the next job to be read into main memory
 - The processor executes the instruction in the user's program until it encounters an ending or error condition
 - Either event causes the processor to fetch its next instruction from the monitor program
- The monitor handles setup and scheduling
 - A batch of jobs is queued up and executed with no idle time
 - Job control language (JCL)
 - Special type of programming language controlling the monitor
- **Monitor, or batch OS, is simply a computer program and relies on the processors ability to fetch instructions from various portions of main memory and to seize and relinquish control alternately**

Desirable CPU Features

- **Memory Protection**
- **Timer**
- **Privileged Instructions**
- **Interrupts**

Desirable CPU Features

- **Memory Protection**

- User program must not alter the memory area containing the monitor
- The processor hardware should detect an error and transfer control to the monitor
- The monitor aborts the job, prints an error message, and loads the next job

- **Timer**

- **Privileged Instructions**

- **Interrupts**

Desirable CPU Features

- **Memory Protection**
- **Timer**
 - Used to prevent a job from monopolizing the system
 - If the timer expires an interrupt occurs and control returns to monitor
- **Privileged Instructions**
- **Interrupts**

Desirable CPU Features

- **Memory Protection**
- **Timer**
- **Privileged Instructions**
 - Can only be executed by the monitor
 - If the processor encounters such an instruction while executing a user program an error interrupt occurs
 - I/O instructions are privileged so the monitor retains control of all I/O devices
- **Interrupts**

Desirable CPU Features

- **Memory Protection**
- **Timer**
- **Privileged Instructions**
- **Interrupts**
 - Gives the OS more flexibility in relinquishing control to and regaining control from user programs

Multiprogramming and Time Sharing Systems

■ Multiprogramming Batch Processing

- The monitor executes several programs at the same time
- More efficient, as wait times can be used by the processor
- Nevertheless, there is no interaction possible

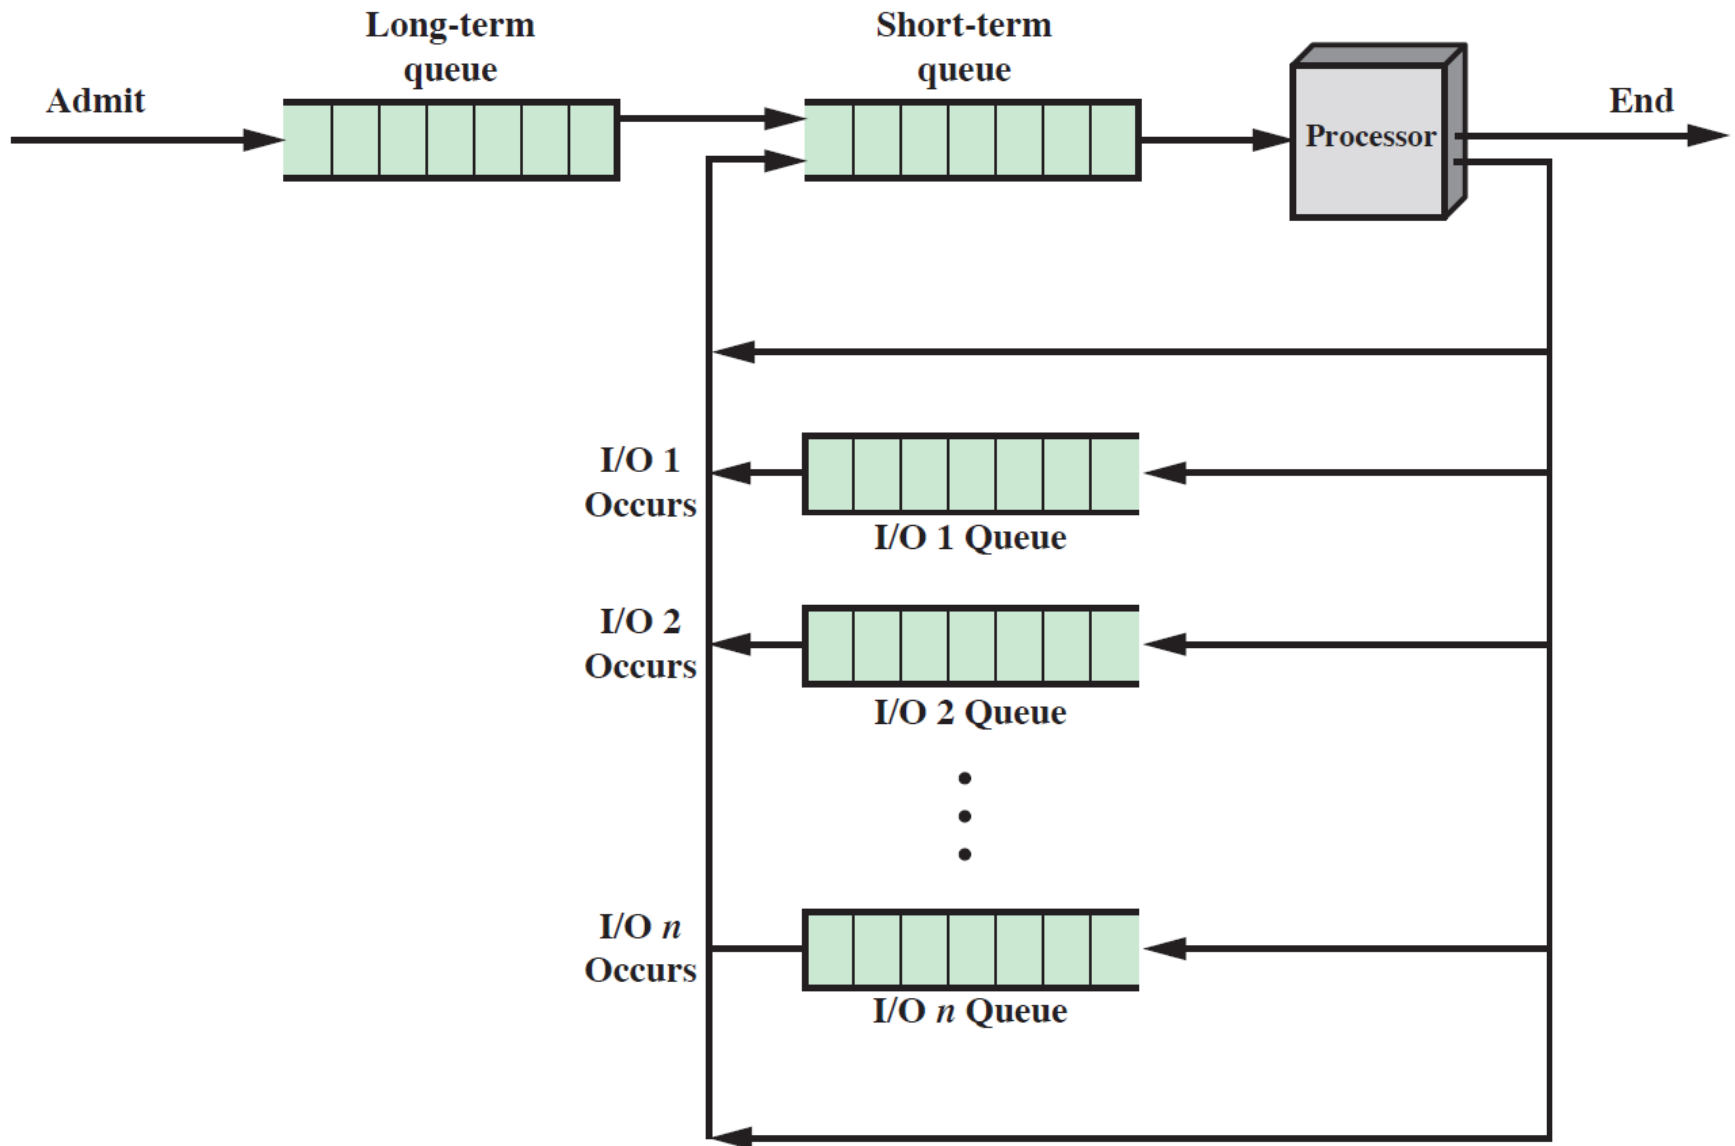
■ Time Sharing Systems

- Used when the user interacts directly with the computer
- Processor's time is shared among multiple users
- Multiple users simultaneously access the system through terminals, with the OS interleaving the execution of each user program in a short burst or quantum of computation

Scheduling

- The key to multiprogramming
- **Long-Term Scheduling**
 - The decision to add to the pool of processes to be executed
- **Medium-Term Scheduling**
 - The decision to add to the number of processes that are partially or fully in main memory
- **Short-Term Scheduling**
 - The decision as to which available process will be executed by the processor
- **I/O Scheduling**
 - The decision as to which process's pending I/O request shall be handled by an available I/O device

Scheduling





Operating System Support


Lecture Content

- OS in General
- **Memory Management**
- Virtual Memory
- Segmentation
- Memory Management of the Pentium II

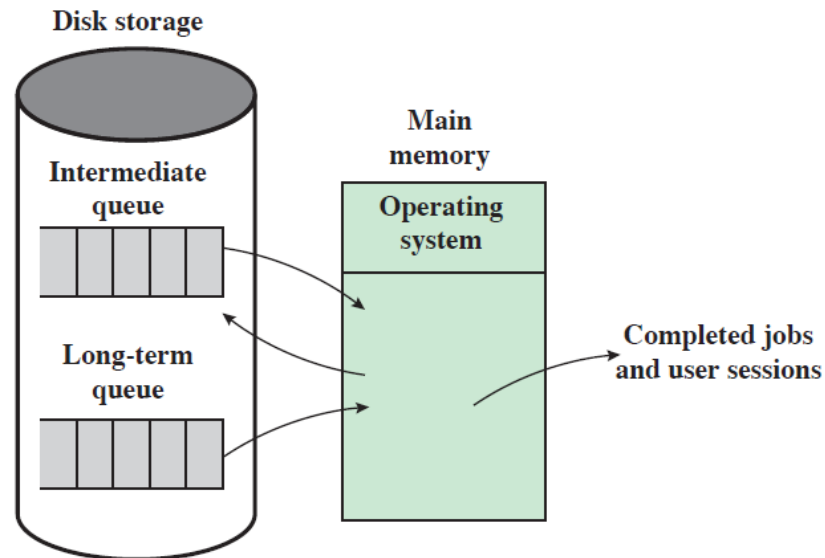
Memory Management

- Programs being executed have to be in main memory
- In order to accommodate more and more programs, we have to increase the amount of memory available
- But there are two flaws in this approach:
 - First, main memory is expensive, even today
 - Second, the appetite of programs for memory has grown as fast as the cost of memory has dropped
 - So larger memory results in larger processes, not more processes

Swapping




(a) Simple job scheduling




(b) Swapping

Fixed-Size Partitions

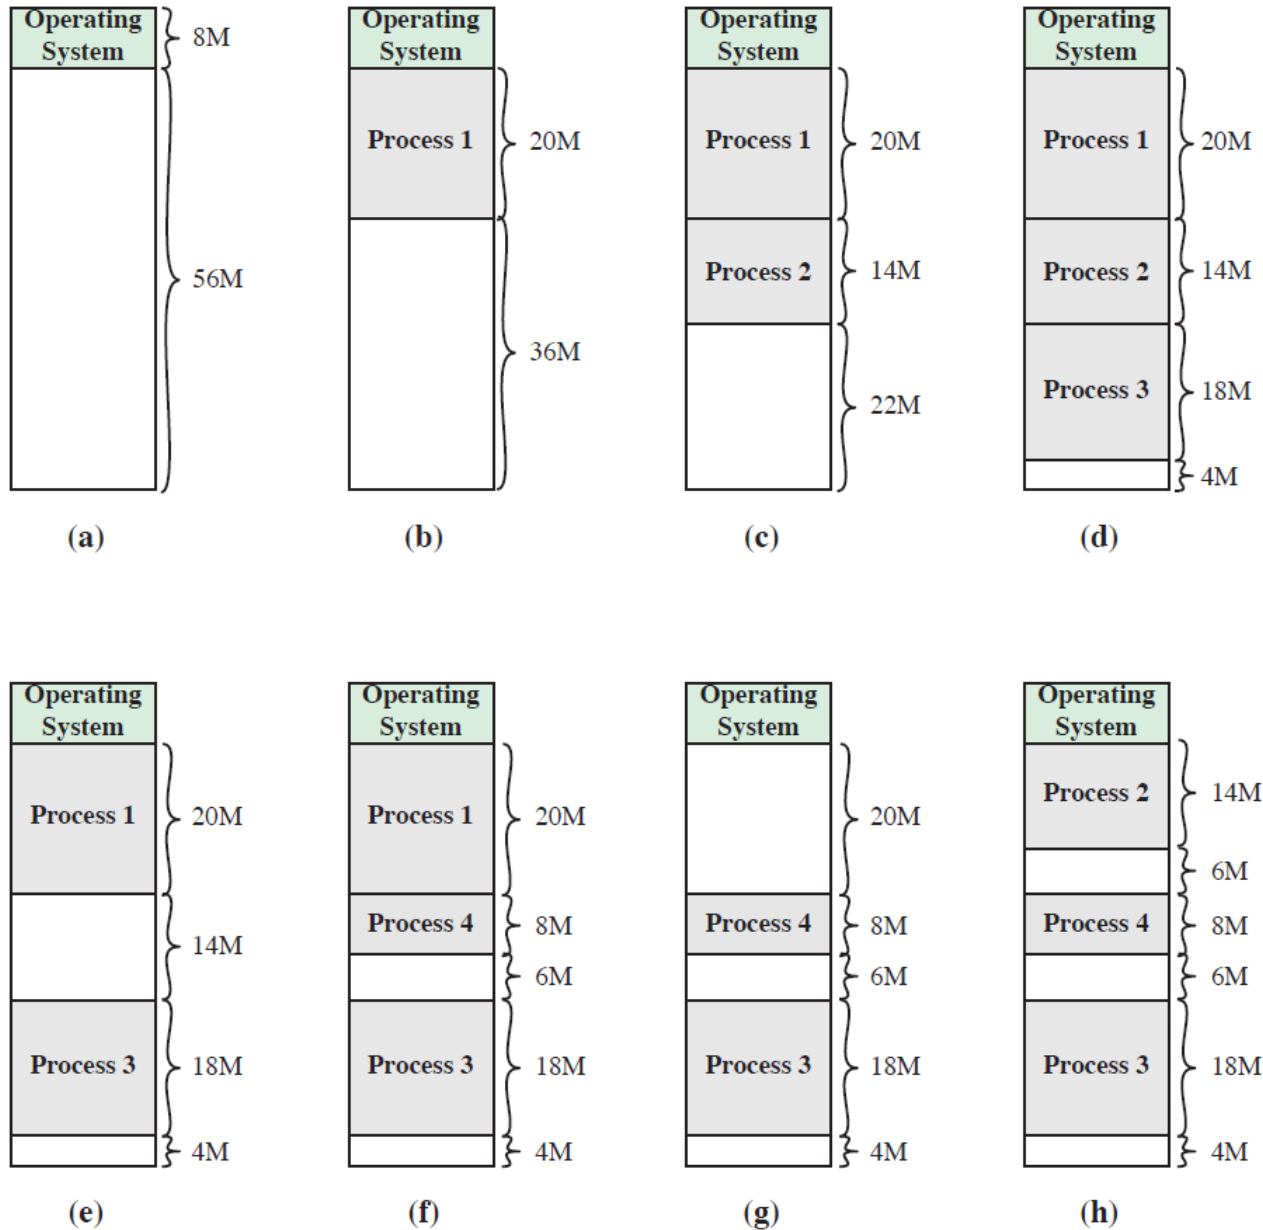


(a) Equal-size partitions



(b) Unequal-size partitions

Dynamic Partitioning



Observations

- It is not likely that a program is in the same place in main memory each time it is swapped in
- A process in memory consists of instructions and data
 - Addresses of data items
 - Addresses of instructions, used for branching instructions
- To solve this problem, programs only use logical addresses
- **Logical Address**
 - expressed as a location relative to the beginning of the program
- **Physical Address**
 - an actual location in main memory
- **Base Address**
 - current starting location of the process

Limitations

- With a simple virtual address, we can fit a program and its data anywhere in the memory
- **But:**
 - Do we really know how much memory do we need?
 - What, if the program requires more space during runtime, but the space is already occupied?
 - Where to place the stack?
 - ...



Operating System Support

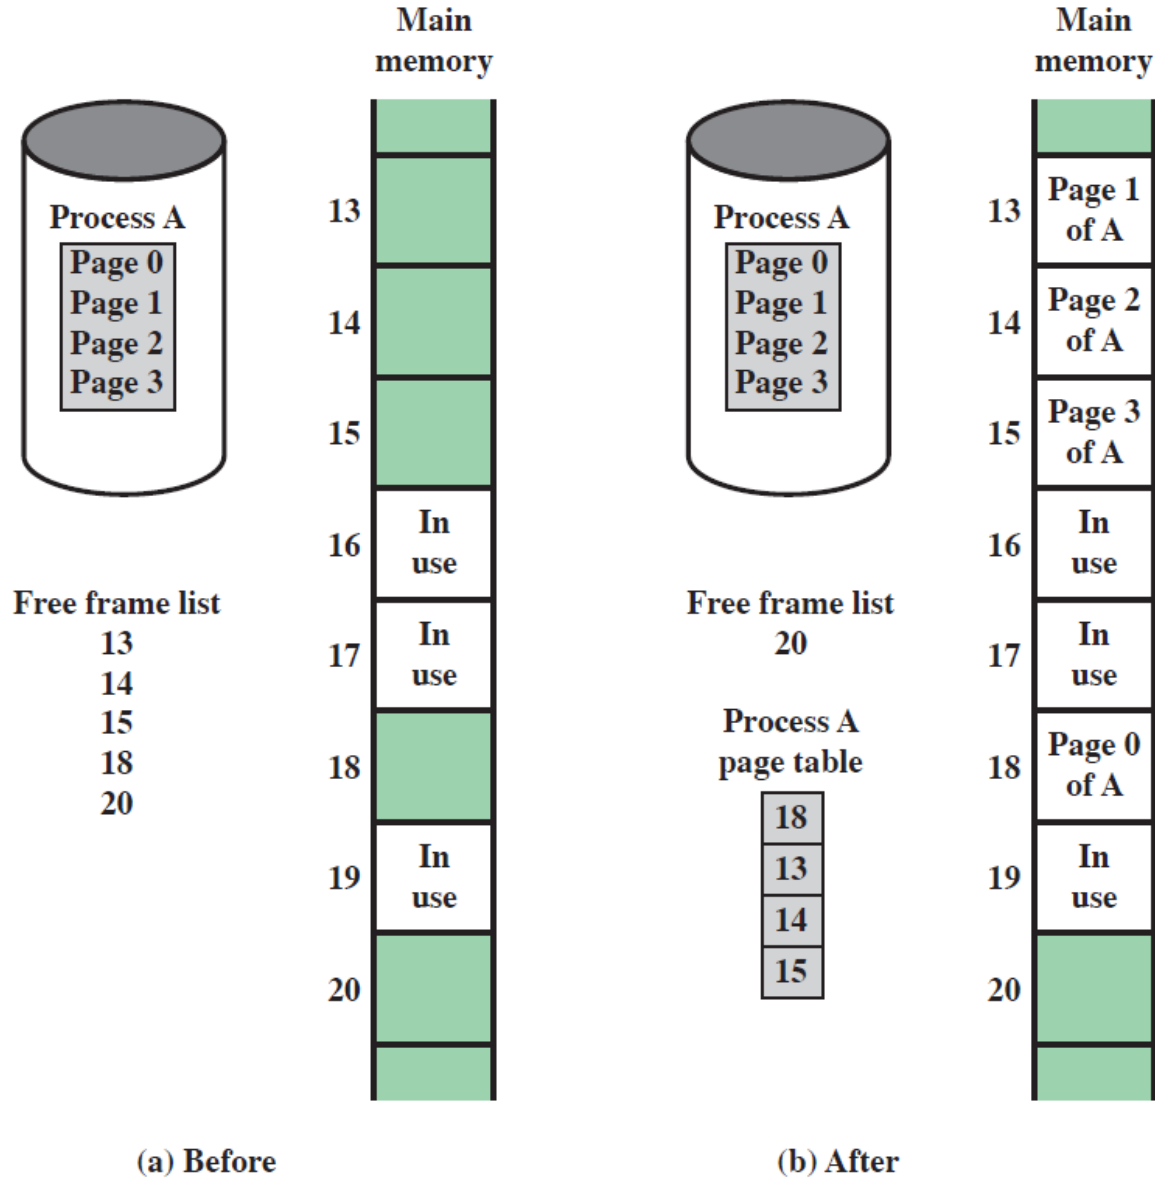
Lecture Content

- OS in General
- Memory Management
- **Virtual Memory**
- Segmentation
- Memory Management of the Pentium II

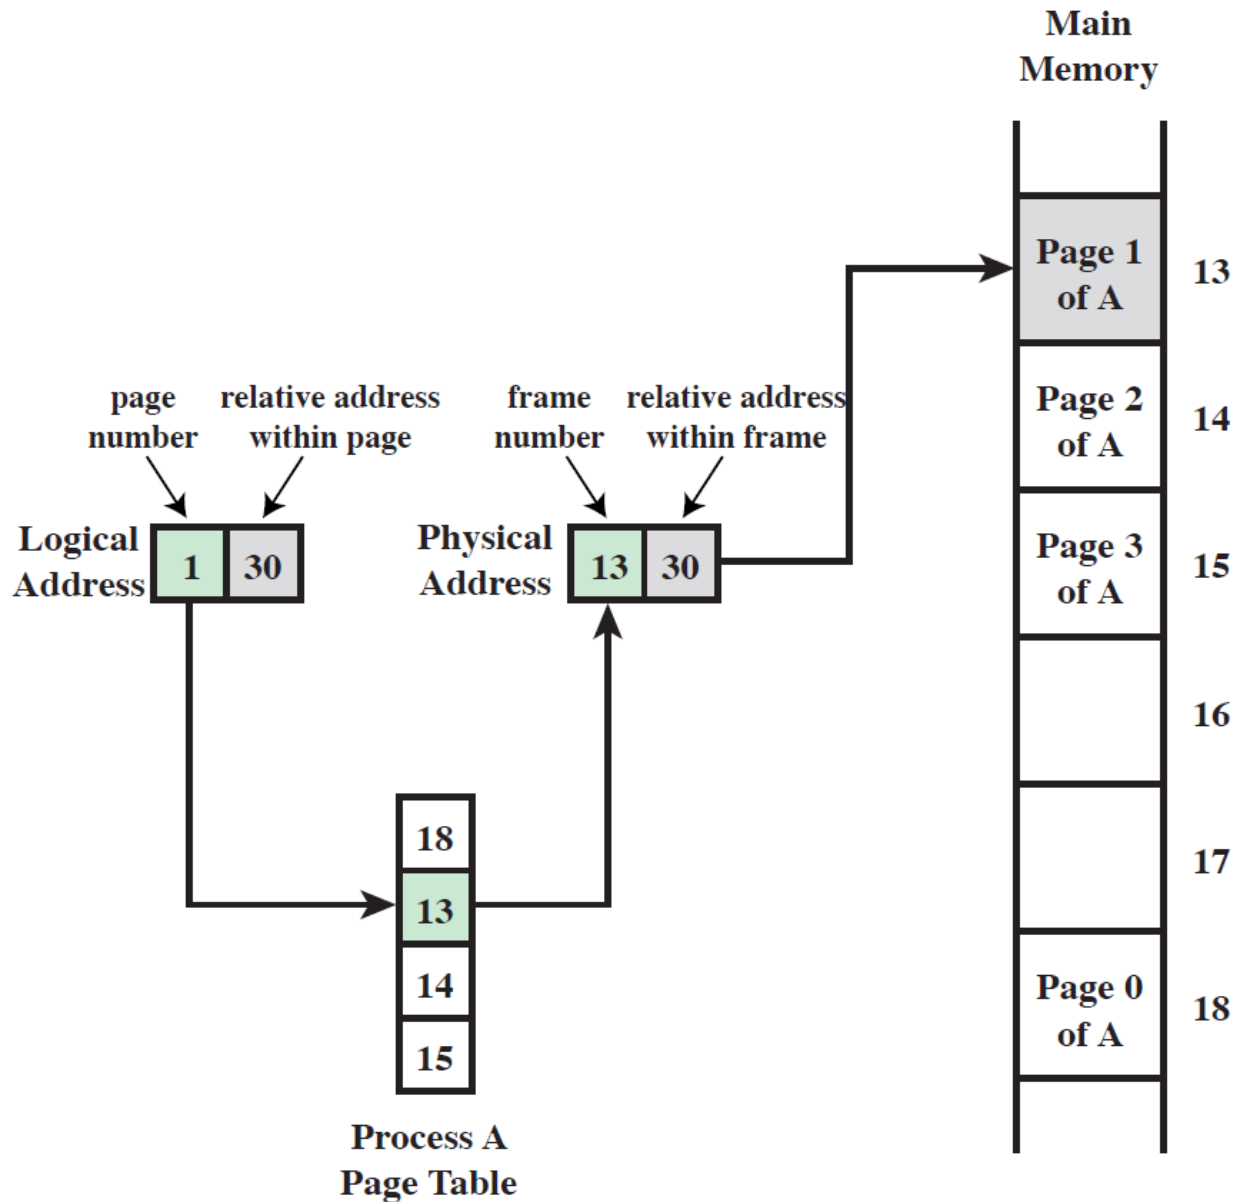
Drawbacks of Fixed-Size or Dynamic Portioning

- Both unequal fixed-size and variable-size partitions are inefficient in the use of memory
- Solution: Paging
 - Memory is partitioned into equal fixed-size small chunks
 - Each process is also divided into small fixed-size chunks
 - Then the chunks of a program, known as pages, are assigned to available chunks of memory, known as frames, or page frames
 - Wasted space in memory for a process is only a fraction of the last page
 - The list of free frames is maintained by the OS

Paging



Paging



Paging

- Paging extends the concept of logical addresses
- A simple base address will no longer suffice
- Rather, the OS maintains a page table for each process
- The page table shows the frame location for each page of the process
- The processor translates the logical address into a physical address
- The processor must know how to access the page table of the current process
- Presented with a logical address (page number, relative address), the processor uses the page table to produce a physical address (frame number, relative address)

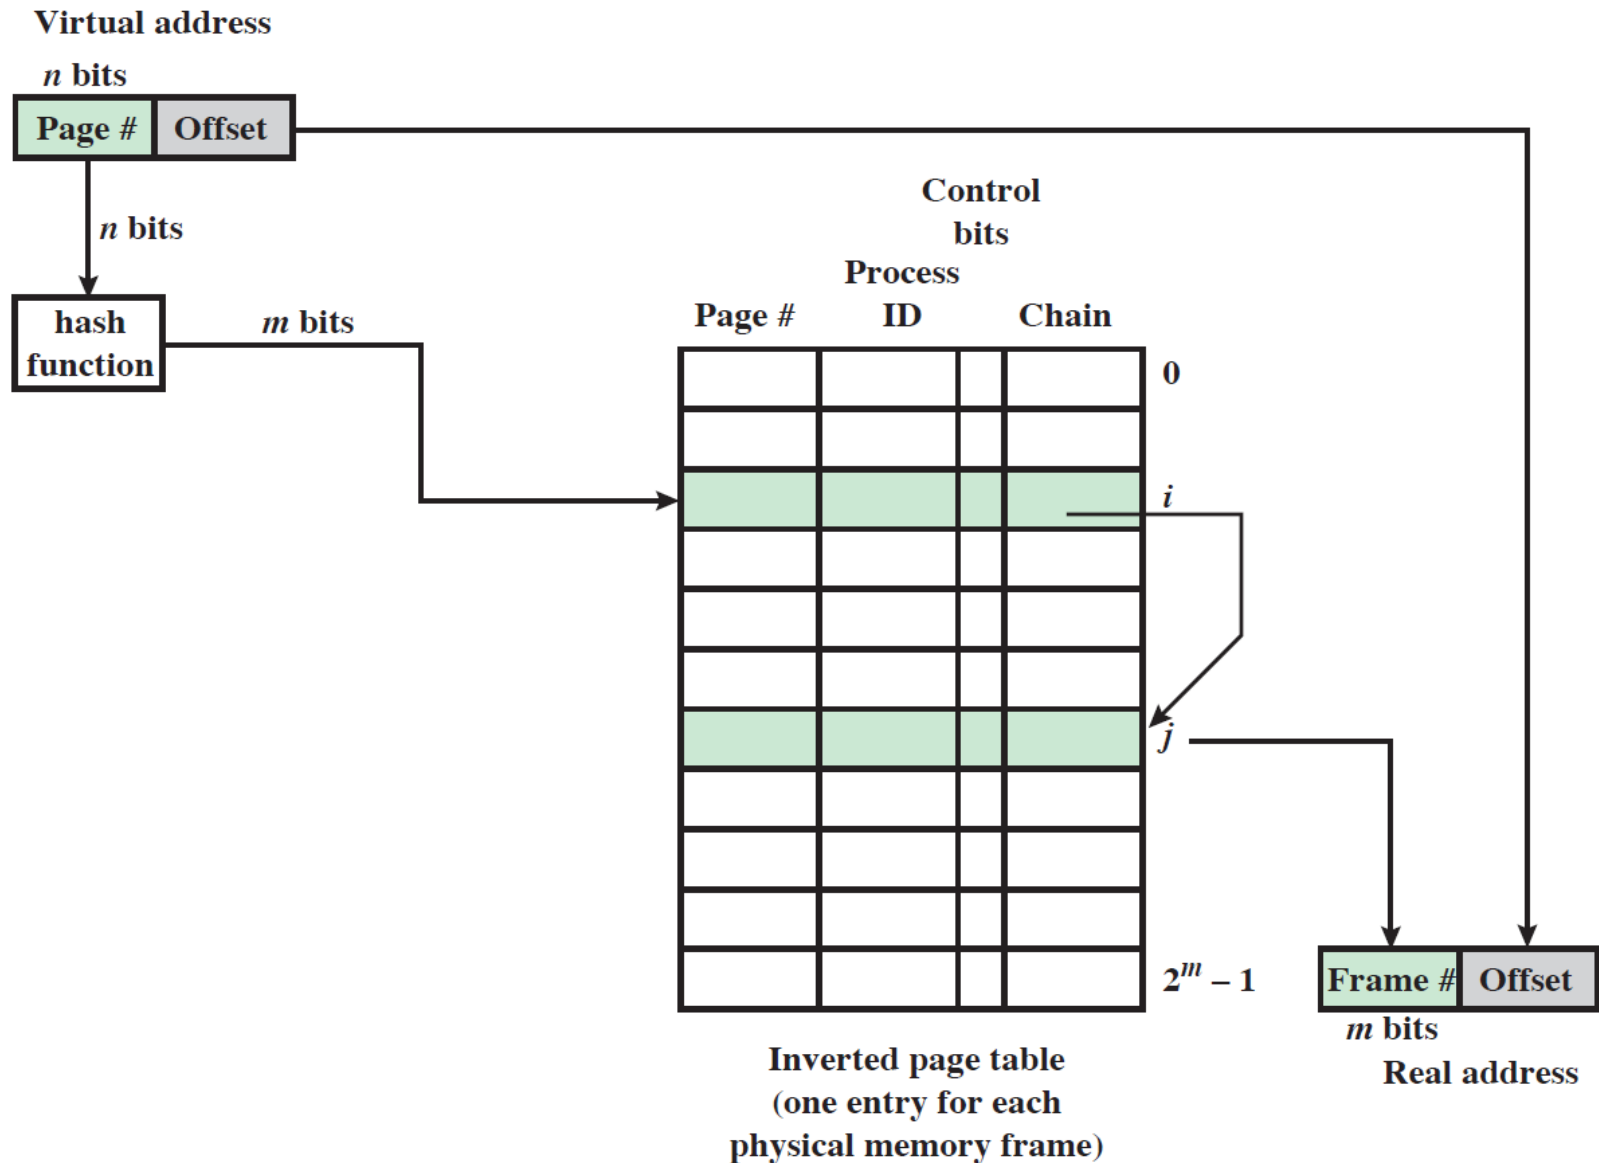
Virtual Memory

- Each page of a process is brought in only when it is needed
 - Principle of locality
 - If the program references to a page not in main memory, a page fault is triggered which tells the OS to bring in the desired page
- **Advantages**
 - More processes can be maintained in memory
 - Time is saved because unused pages are not swapped in and out of memory
- **Disadvantages**
 - Page replacement: A new page throws out an existing page
 - This replacements can get inefficient
 - Thrashing
 - When the processor spends most of its time swapping pages rather than executing instructions

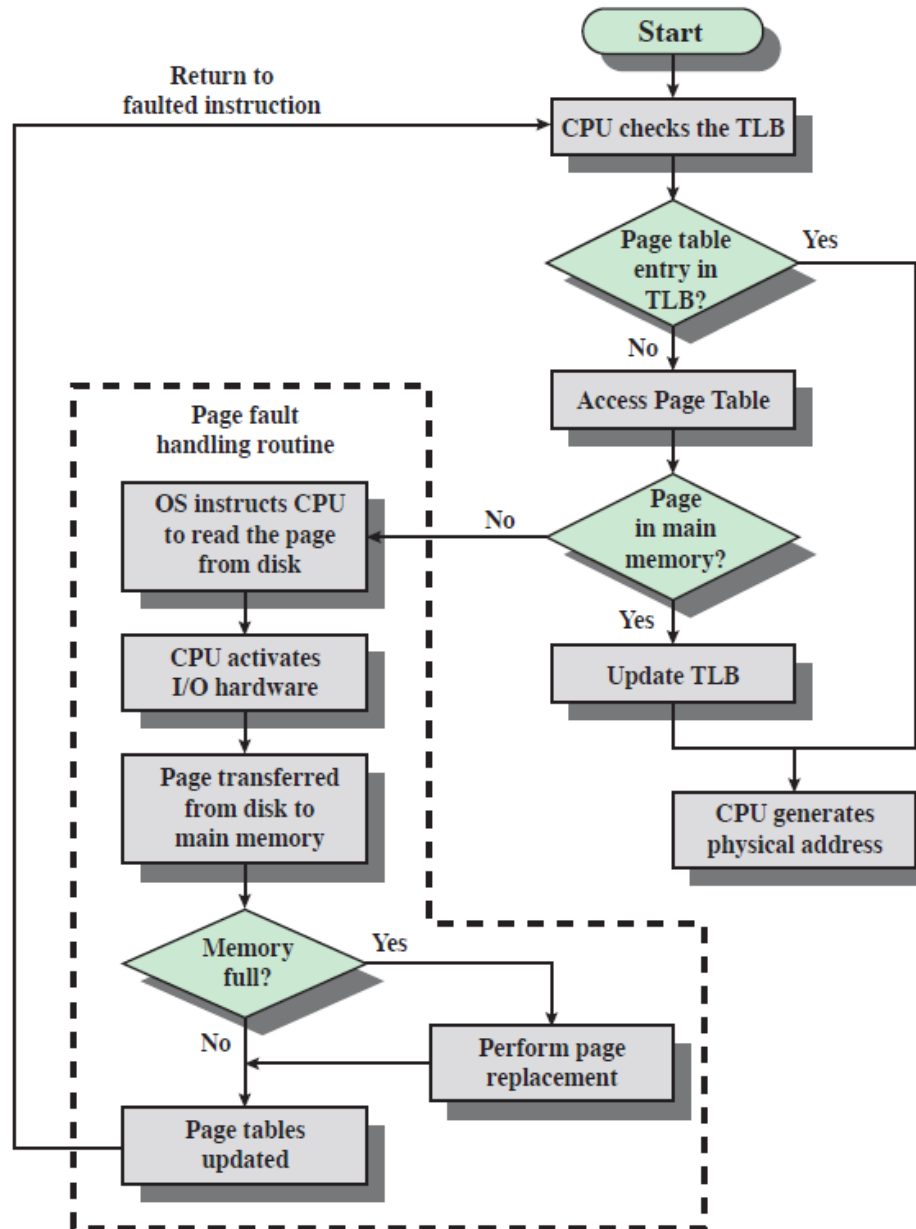
Handling the Page Tables

- In most systems, there is one page table per process
- Allowing each process to have up to $2^{31} = 2$ Gbytes of virtual memory using $2^9 = 512$ -byte pages results in 2^{22} page table entries
- To overcome this problem
 - store page tables in virtual memory rather than real memory
 - This means that page tables are subject to paging just as other pages are
 - Two-level scheme
 - There is a page directory, in which each entry points to a page table
 - Typically, the maximum length of a page table is restricted to be equal to one page

Inverted Page Table Structure



Paging and Translation Lookaside Buffer (TLB)


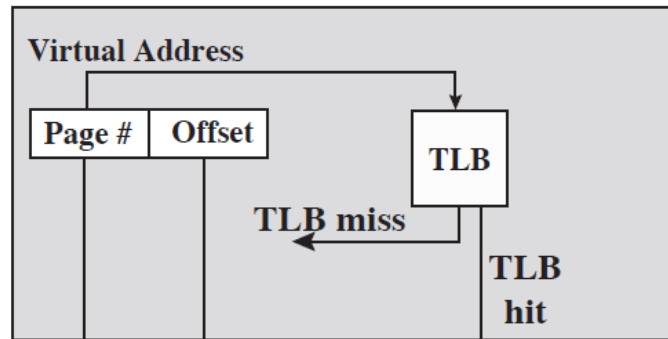


Where to put the cache?

- Where would you put the cache?
 - Before TLB using virtual addresses and PID?
 - After TLB using real addresses?
- What are the implications?

TLB & Cache



TLB Operation



Page Table

\oplus

Cache Operation





Operating System Support

Lecture Content

- OS in General
- Memory Management
- Virtual Memory
- **Segmentation**
- Memory Management of the Pentium II

Segmentation

- Usually visible to the programmer
- Provided as a convenience for organizing programs and data
- Means for associating privilege and protection attributes with instructions and data
- Allows the programmer to view memory as consisting of multiple address spaces or segments

Segmentation

- Usually visible to the programmer
- Provided as a convenience for organizing programs and data
- Means for associating privilege and protection attributes with instructions and data
- Allows the programmer to view memory as consisting of multiple address spaces or segments
- **Advantages**
 - Simplifies the handling of growing data structures
 - Allows programs to be altered and recompiled independently without requiring that an entire set of programs be re-linked and re-loaded
 - Lends itself to sharing among processes
 - Lends itself to protection

Different Operation Modes

- **Unsegmented Unpaged Memory**
- **Unsegmented Paged Memory**
- **Segmented Unpaged Memory**
- **Segmented Paged Memory**

Different Operation Modes

- **Unsegmented Unpaged Memory**
 - Virtual address is the same as the physical address
 - Useful in low-complexity, high performance controller applications
- **Unsegmented Paged Memory**
- **Segmented Unpaged Memory**
- **Segmented Paged Memory**

Different Operation Modes

- **Unsegmented Unpaged Memory**
- **Unsegmented Paged Memory**
 - Memory is viewed as a paged linear address space
 - Protection and management of memory is done via paging
 - Favored by some operating systems
- **Segmented Unpaged Memory**
- **Segmented Paged Memory**

Different Operation Modes

- **Unsegmented Unpaged Memory**
- **Unsegmented Paged Memory**
- **Segmented Unpaged Memory**
 - Memory is viewed as a collection of logical address spaces
 - Affords protection down to the level of a single byte
 - Guarantees that the translation table needed is on-chip when the segment is in memory
 - Results in predictable access times
- **Segmented Paged Memory**

Different Operation Modes

- **Unsegmented Unpaged Memory**
- **Unsegmented Paged Memory**
- **Segmented Unpaged Memory**
- **Segmented Paged Memory**
 - Segmentation is used to define logical memory partitions subject to access control, and paging is used to manage the allocation of memory within the partitions
 - Operating systems such as UNIX System V favor this view

Segmentation of the Pentium II

- Each virtual address consists of a 16-bit segment reference and a 32-bit offset
 - Two bits of segment reference deal with the protection mechanism
 - 14 bits specify segment
- Unsegmented virtual memory is $2^{32} = 4\text{Gbytes}$
- Segmented virtual memory is $2^{46} = 64\text{ terabytes (Tbytes)}$
- Physical address space employs a 32-bit address for a maximum of 4 Gbytes
- Virtual address space is divided into two parts
 - One-half is global, shared by all processes
 - The remainder is local and is distinct for each process

Segment Protection

- Associated with each segment are two forms of protection:
 - Privilege level
 - Access attribute
- There are four privilege levels
 - Most protected (level 0)
 - Least protected (level 3)
- Privilege level of data segment is called “**classification**”
- Privilege level of a program is called “**clearance**”
- A program can only access data segments with $\text{classification} \leq \text{clearance}$
- The privilege mechanism also limits the use of certain instructions

Paging on Pentium II

- Segmentation may be disabled
 - In which case linear address space is used
- Two level page table lookup
 - First, page directory
 - 1024 entries max
 - Splits 4 Gbyte linear memory into 1024 page groups of 4 Mbyte
 - Each page table has 1024 entries corresponding to 4 Kbyte pages
 - Can use one page directory for all processes, one per process or mixture
 - Page directory for current process always in memory
 - Use TLB holding 32 page table entries
 - Two page sizes available, 4k or 4M

Pentium II Address Translation Mechanism

