



Processor Structure and Function

Lecture Content

- **Processor Organization**
- **Pipelining & Branch Prediction**
- **Intel Pipelining & Processor Structure**



Processor Structure and Function

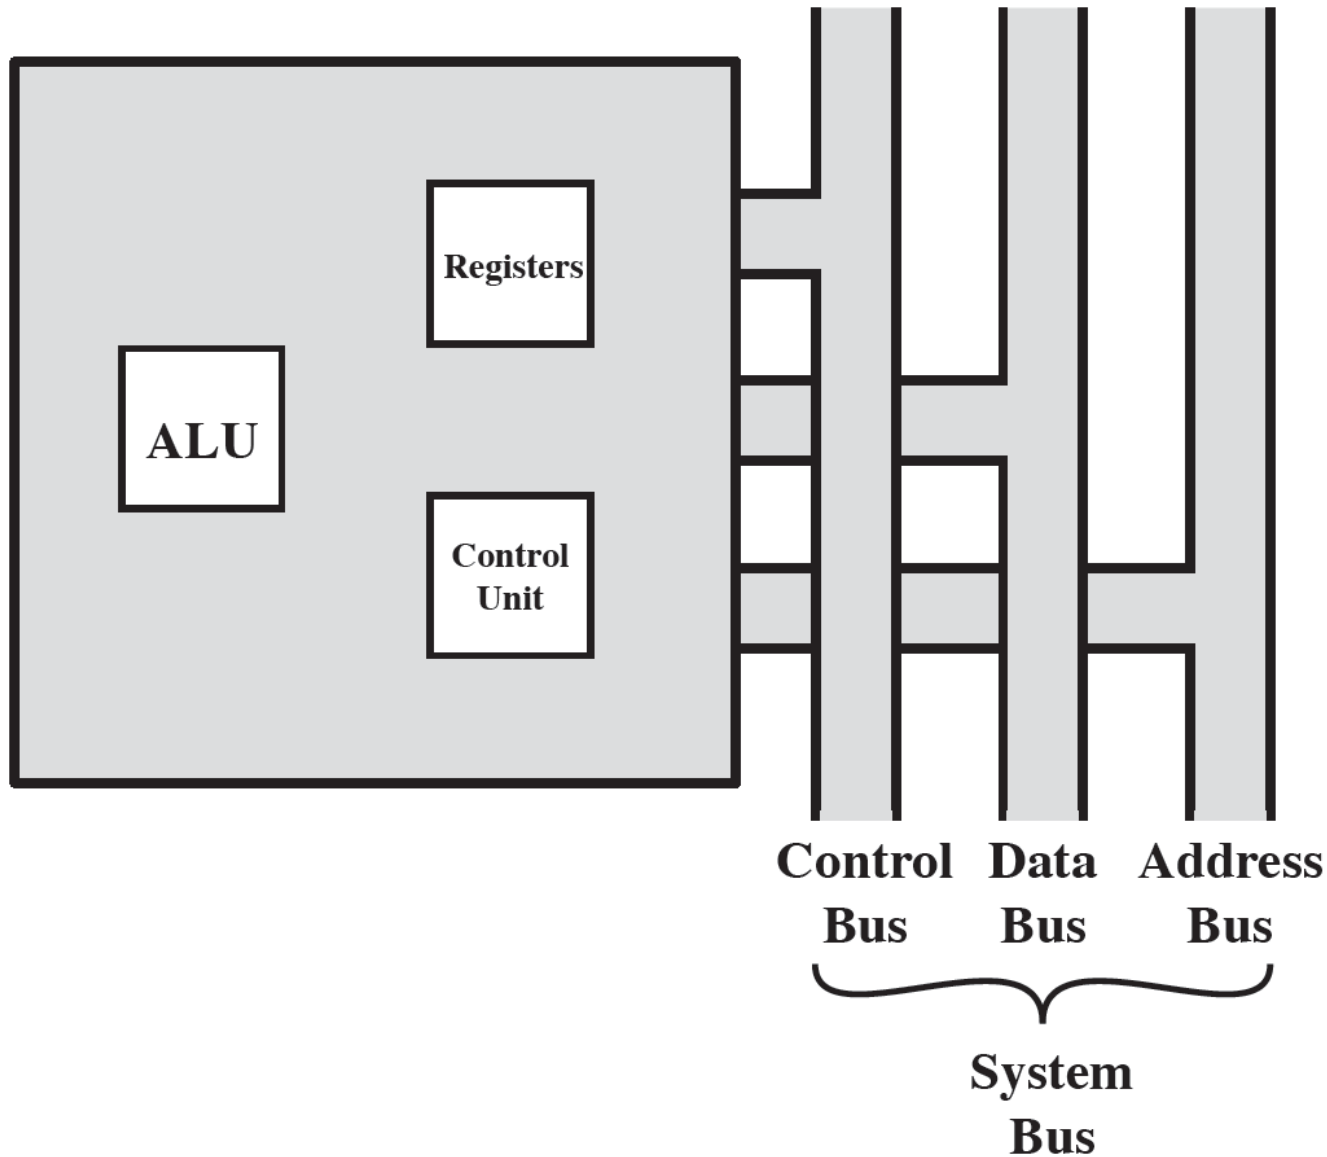
Learning Objectives

- **Recap of the instruction cycle**
- **Distinguish between user-visible and control/status register**
- **Understand the principles of pipelining**
- **Understand pipeline hazards**
- **Understand branch prediction**

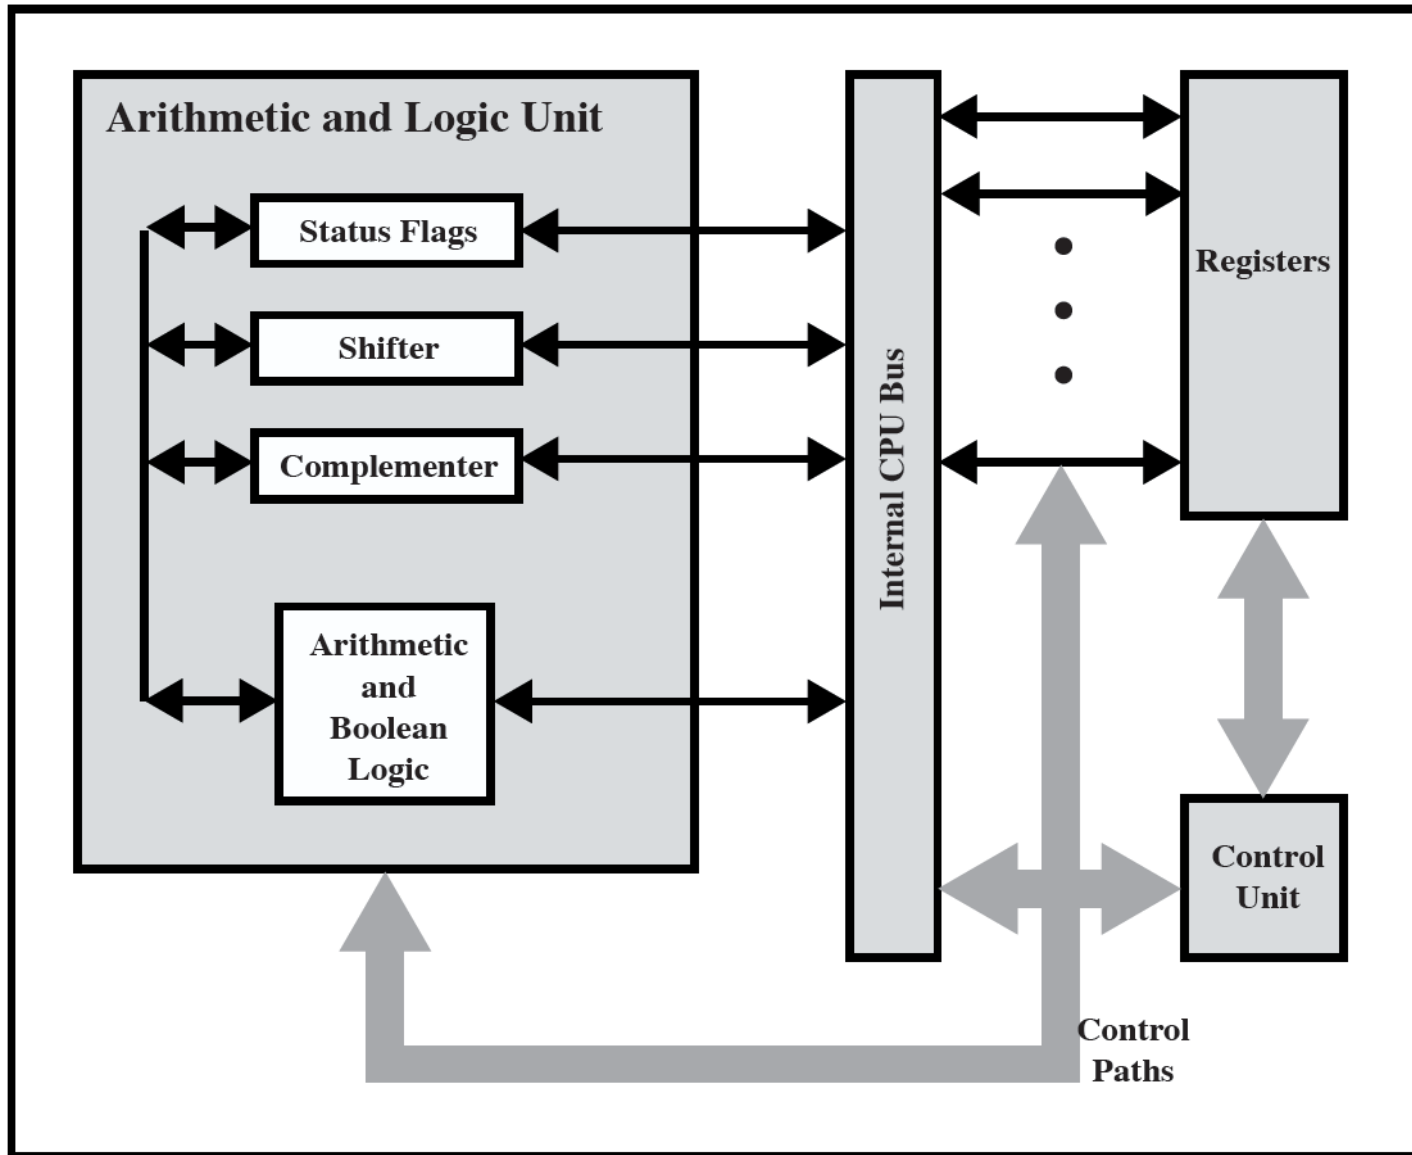
Requirements for a Processor

- **Fetch instruction**
 - The processor reads an instruction from register, cache, main memory
- **Interpret instruction**
 - The instruction is decoded to determine what action is required
- **Fetch data**
 - The instruction may require reading data from memory or an I/O module
- **Process data**
 - The instruction may be performing some arithmetic or logical operation
- **Write data**
 - The results may require writing data to memory or an I/O module
- In order to do these things the processor needs to store some data temporarily and therefore needs a small internal memory

Schematic View of a Processor



Schematic Internal Structure



Register Organization

- Within the processor there is a set of registers that function as a level of memory above main memory and cache in the hierarchy
- The registers in the processor perform two roles:
 - **User Visible Registers**
 - Enable the machine or assembly language programmer to minimize main memory references by optimizing use of registers
 - **Control and Status Registers**
 - Used by the control unit to control the operation of the processor and by privileged operating system programs to control the execution of programs

User-Visible Registers

- **General Purpose**
 - Can be assigned to a variety of functions by the programmer
- **Data**
 - May be used only to hold data and cannot be employed in the calculation of an operand address
- **Address**
 - May be somewhat general purpose or may be devoted to a particular addressing mode
 - Examples: segment pointers, index registers, stack pointer
- **Condition Codes**
 - Also referred to as flags
 - Bits set by the processor hardware as the result of operations

Design Issues

■ General-Purpose vs. Specialized

- Affects instruction set design
- Specialized registers are less complicated and slim the ISA
- Specialization limits the programmer's flexibility.

■ Number of Registers

- Somewhere between 8 and 32 registers appears optimum
- Fewer registers result in more memory references
- More registers do not noticeably reduce memory references

■ Register Length

- Address registers must be long enough to hold the largest address
- Data registers should be able to hold values of most data types

Condition Codes

- Many processors do not use condition codes at all
- A conditional branch instructions specify a comparison to be made and act on the result of the comparison
- **Advantages**
- **Disadvantages**

Condition Codes

- Many processors do not use condition codes at all
- A conditional branch instructions specify a comparison to be made and act on the result of the comparison
- **Advantages**
 - Set by normal arithmetic and data movement instructions
 - Conditional instructions, such as BRANCH are simplified relative to composite instructions, such as TEST AND BRANCH
 - Condition codes facilitate multi-way branches.
 - Condition codes can be saved on the stack during subroutine calls along with other register information
- **Disadvantages**

Condition Codes

- Many processors do not use condition codes at all
- A conditional branch instructions specify a comparison to be made and act on the result of the comparison
- **Advantages**
- **Disadvantages**
 - Condition codes add complexity
 - Condition codes are irregular; they are typically not part of the main data path, so they require extra hardware connections
 - Often condition code machines must add special non-condition-code instructions for special situations anyway, such as bit checking, loop control, and atomic semaphore operations
 - In a pipelined implementation, condition codes require special synchronization to avoid conflicts

Control and Status Registers

- **Program Counter (PC)**
 - Contains the address of an instruction to be fetched
- **Instruction Register (IR)**
 - Contains the instruction most recently fetched
- **Memory Address Register (MAR)**
 - Contains the address of a location in memory
- **Memory Buffer Register (MBR)**
 - Contains a word of data to be written to memory or the word most recently read

Program Status Word (PSW)

- **Sign**

- Contains the sign bit of the result of the last arithmetic operation

- **Zero**

- Set when the result is 0

- **Equal**

- Set if a logical compare result is equality

- **Overflow**

- Used to indicate arithmetic overflow

- **Carry**

- Set if an operation resulted in a carry (addition) into or borrow (sub- traction) out of a high-order bit. Used for multiword arithmetic operations

- **Interrupt Enable/Disable**

- Used to enable or disable interrupts

- **Supervisor**

- Indicates whether the processor is executing in supervisor or user mode

Register Organization Examples

Data registers	
D0	
D1	
D2	
D3	
D4	
D5	
D6	
D7	

Address registers	
A0	
A1	
A2	
A3	
A4	
A5	
A6	
A7'	

Program status	
Program counter	
	Status register

(a) MC68000

General registers

AX	Accumulator
BX	Base
CX	Count
DX	Data

Pointers & index

SP	Stack ptr
BP	Base ptr
SI	Source index
DI	Dest index

Segment

CS	Code
DS	Data
SS	Stack
ES	Extrat

Program status

Flags
Instr ptr

(b) 8086

General Registers

EAX		AX
EBX		BX
ECX		CX
EDX		DX

ESP		SP
EBP		BP
ESI		SI
EDI		DI

Program Status

FLAGS Register
Instruction Pointer

(c) 80386 - Pentium 4



Processor Structure and Function

Lecture Content

- Processor Organization
- **Pipelining & Branch Prediction**
- Intel Pipelining & Processor Structure

Recall: The Instruction Cycle

- **Fetch**

- Read the next instruction from memory into the processor

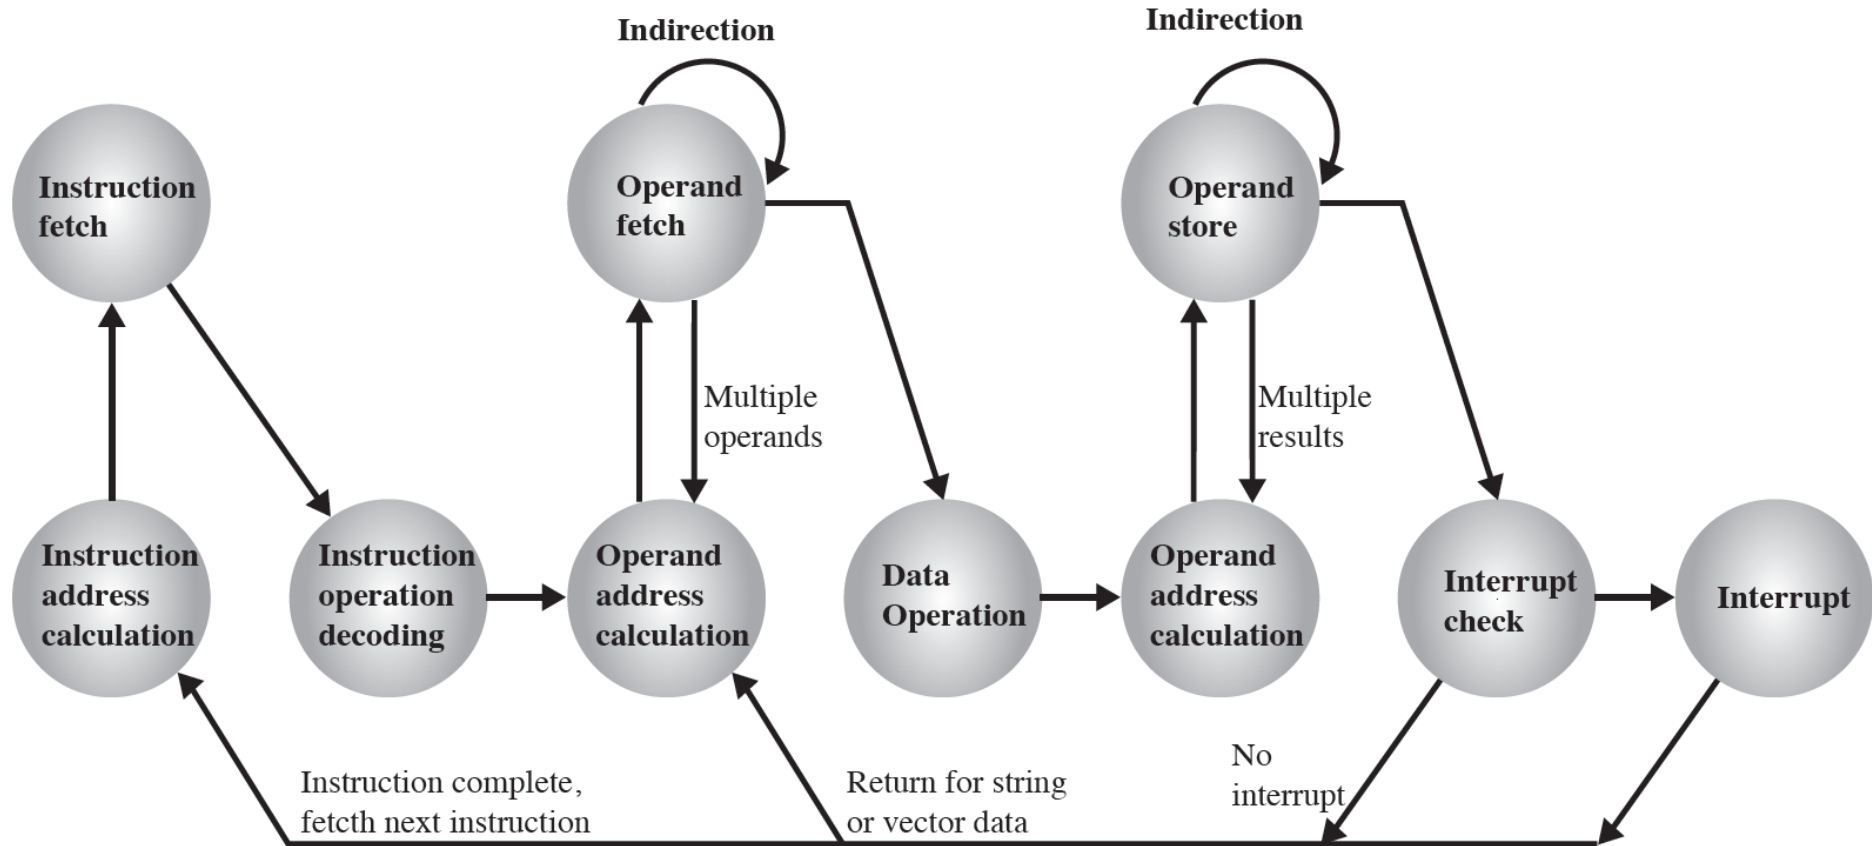
- **Execute**

- Interpret the opcode and perform the indicated operation

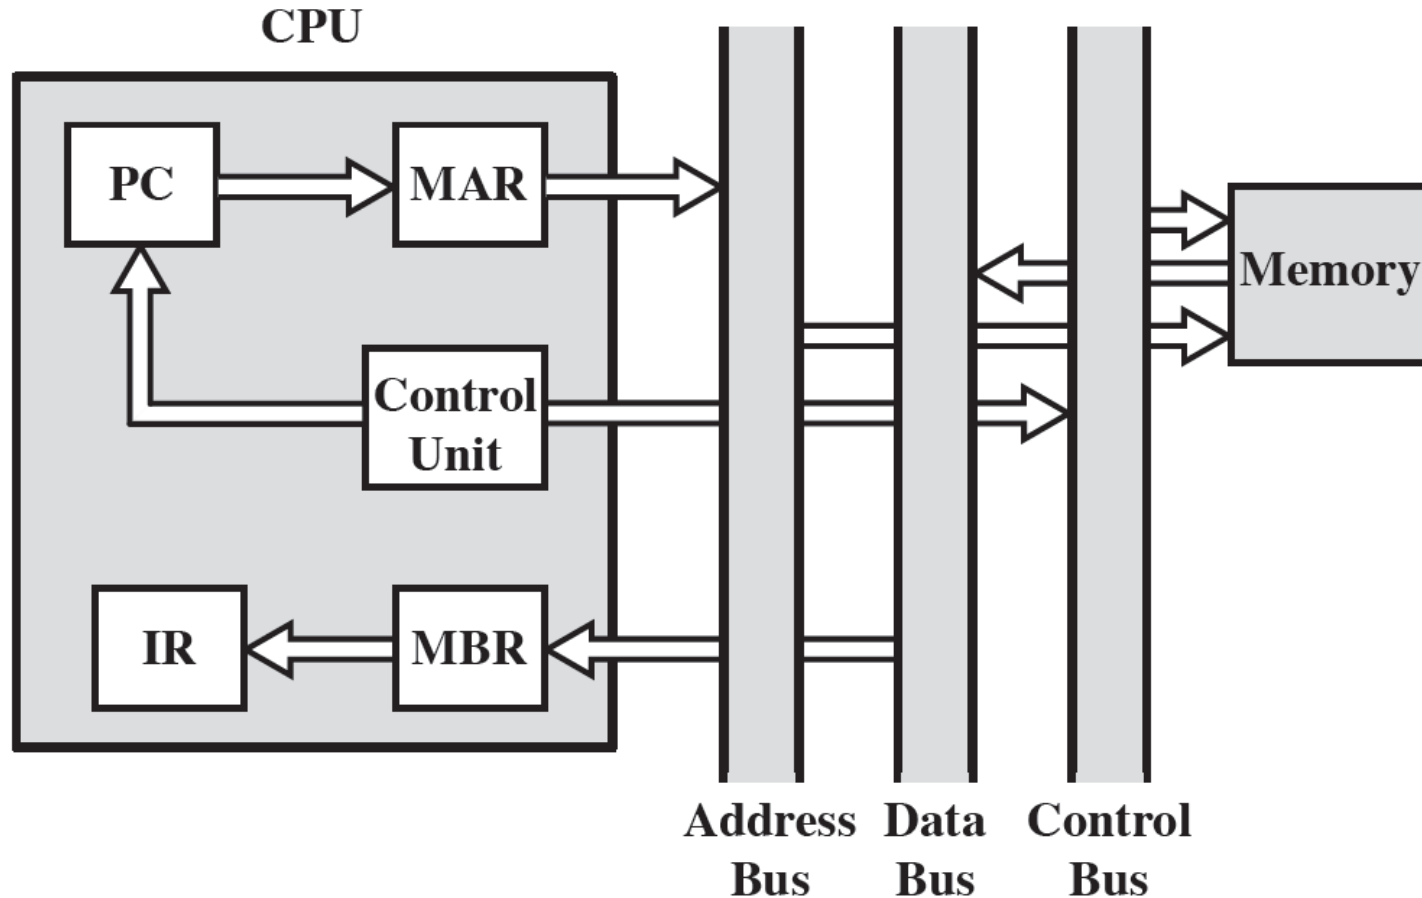
- **Interrupt**

- If interrupts are enabled and an interrupt has occurred, save the current process state and service the interrupt

The Instruction Cycle



Example Data Flow: Fetch Cycle



MBR = Memory buffer register
MAR = Memory address register
IR = Instruction register
PC = Program counter

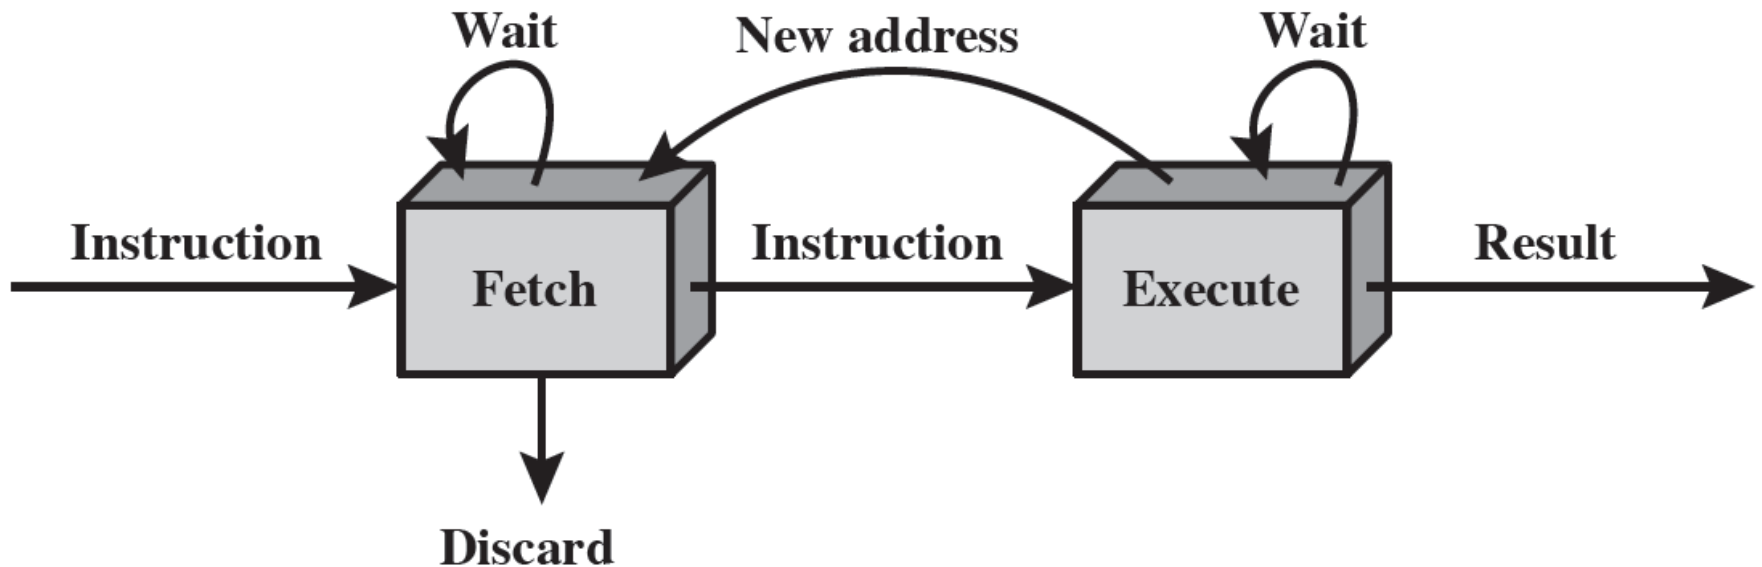
Pipelining: General Idea

- Each step of the instruction cycle only requires parts of the CPU's infrastructure
- Similar to the use of an assembly line in a manufacturing plant
- New inputs are accepted at one end before previously accepted inputs appear as outputs at the other end
- To apply this concept to instruction execution we must recognize that an instruction has a number of stages

A Simple Two Stage Pipeline



(a) Simplified view



More Stages of a Pipeline

■ Fetch Instruction (FI)

- Read the next expected instruction into a buffer

■ Decode Instruction (DI)

- Determine the opcode and the operand specifiers

■ Calculate Operands (CO)

- Calculate the effective address of each source operand
- This may involve displacement, register indirect, indirect, or other forms of address calculation

■ Fetch Operands (FO)

- Fetch each operand from memory
- Operands in registers need not be fetched

■ Execute Instruction (EI)

- Perform the indicated operation and store the result, if any, in the specified destination operand location

■ Write Operand (WO)

- Store the result in memory

Principle of a Pipeline

Time →

	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Instruction 1	FI	DI	CO	FO	EI	WO								
Instruction 2		FI	DI	CO	FO	EI	WO							
Instruction 3			FI	DI	CO	FO	EI	WO						
Instruction 4				FI	DI	CO	FO	EI	WO					
Instruction 5					FI	DI	CO	FO	EI	WO				
Instruction 6						FI	DI	CO	FO	EI	WO			
Instruction 7							FI	DI	CO	FO	EI	WO		
Instruction 8								FI	DI	CO	FO	EI	WO	
Instruction 9									FI	DI	CO	FO	EI	WO

Speed-Up of a Pipeline

- Assuming we have a pipeline with k stages
- Thus, each instruction requires k cycles to finish
- Once, the pipeline is full, we finish one instruction every cycle
- Does this mean we have a speed-up of k ?
- When does this principle not apply? What can happen?

Mess Things Up: A Conditional Branch

	Time →							← Branch Penalty						
	1	2	3	4	5	6	7	8	9	10	11	12	13	14
Instruction 1	FI	DI	CO	FO	EI	WO								
Instruction 2		FI	DI	CO	FO	EI	WO							
Instruction 3			FI	DI	CO	FO	EI	WO						
Instruction 4				FI	DI	CO	FO							
Instruction 5					FI	DI	CO							
Instruction 6						FI	DI							
Instruction 7							FI							
Instruction 15								FI	DI	CO	FO	EI	WO	
Instruction 16									FI	DI	CO	FO	EI	WO

Alternate Display

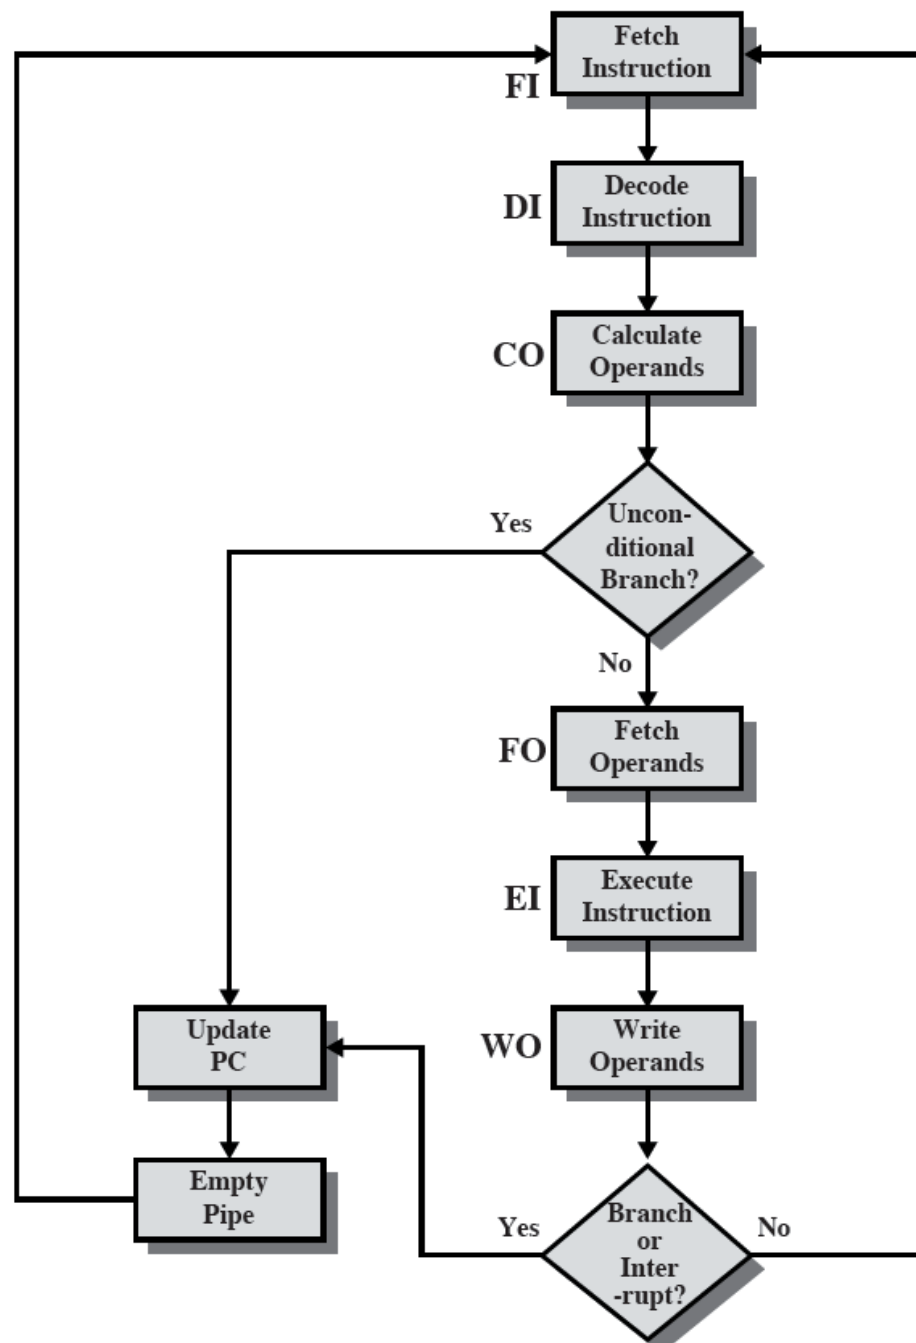
Time
↓

	FI	DI	CO	FO	EI	WO
1	I1					
2	I2	I1				
3	I3	I2	I1			
4	I4	I3	I2	I1		
5	I5	I4	I3	I2	I1	
6	I6	I5	I4	I3	I2	I1
7	I7	I6	I5	I4	I3	I2
8	I8	I7	I6	I5	I4	I3
9	I9	I8	I7	I6	I5	I4
10		I9	I8	I7	I6	I5
11			I9	I8	I7	I6
12				I9	I8	I7
13					I9	I8
14						I9


(a) No branches

	FI	DI	CO	FO	EI	WO
1	I1					
2	I2	I1				
3	I3	I2	I1			
4	I4	I3	I2	I1		
5	I5	I4	I3	I2	I1	
6	I6	I5	I4	I3	I2	I1
7	I7	I6	I5	I4	I3	I2
8	I15					I3
9	I16	I15				
10		I16	I15			
11			I16	I15		
12				I16	I15	
13					I16	I15
14						I16


(b) With conditional branch



Speed-Up



Speed-Up



Pipeline Hazards

- Occur when the pipeline, or some portion of the pipeline, must stall because conditions do not permit continued execution
- Also referred to as a pipeline bubble
- There are three types of hazards:
 - **Resource**
 - **Data**
 - **Control**

Resource Hazard

- A resource hazard occurs when two or more instructions that are already in the pipeline need the same resource
- The result is that the instructions must be executed in serial rather than parallel for a portion of the pipeline
- A resource hazard is sometimes referred to as a structural hazard

Resource Hazard

		Clock cycle								
		1	2	3	4	5	6	7	8	9
Instrucion	I1	FI	DI	FO	EI	WO				
	I2		FI	DI	FO	EI	WO			
	I3			FI	DI	FO	EI	WO		
	I4				FI	DI	FO	EI	WO	

(a) Five-stage pipeline, ideal case

		Clock cycle								
		1	2	3	4	5	6	7	8	9
Instrucion	I1	FI	DI	FO	EI	WO				
	I2		FI	DI	FO	EI	WO			
	I3			Idle	FI	DI	FO	EI	WO	
	I4					FI	DI	FO	EI	WO

(b) I1 source operand in memory

Data Hazards

- **Read After Write (RAW)**, or true dependency
- **Write After Read (WAR)**, or antidependency
- **Write After Write (WAW)**, or output dependency

Data Hazards

- **Read After Write (RAW)**, or true dependency
 - An instruction modifies a register or memory location
 - Succeeding instruction reads data in memory or register location
 - Hazard occurs if the read takes place before write operation is complete
- **Write After Read (WAR)**, or antidependency
- **Write After Write (WAW)**, or output dependency

Data Hazards

- **Read After Write (RAW)**, or true dependency
- **Write After Read (WAR)**, or antidependency
 - An instruction reads a register or memory location
 - Succeeding instruction writes to the location
 - Hazard occurs if the write operation completes before the read operation takes place
- **Write After Write (WAW)**, or output dependency

Data Hazards

- **Read After Write (RAW)**, or true dependency
- **Write After Read (WAR)**, or antidependency
- **Write After Write (WAW)**, or output dependency
 - Two instructions both write to the same location
 - Hazard occurs if the write operations take place in the reverse order of the intended sequence

Example For a Data Hazard

		Clock cycle									
		1	2	3	4	5	6	7	8	9	10
add %rbx,%rax		FI	DI	FO	EI	WO					
sub %rax,%rcx			FI	DI	Idle		FO	EI	WO		
I3				FI			DI	FO	EI	WO	
I4							FI	DI	FO	EI	WO

Control Hazard

- Also known as a branch hazard
- Occurs when the pipeline makes the wrong decision on a branch prediction
- Brings instructions into the pipeline that must subsequently be discarded
- Dealing with Branches:

Control Hazard

- Also known as a branch hazard
- Occurs when the pipeline makes the wrong decision on a branch prediction
- Brings instructions into the pipeline that must subsequently be discarded
- Dealing with Branches:
 - **Multiple Streams**
 - **Prefetch Branch Target**
 - **Loop Buffer**
 - **Branch Prediction**
 - **Delayed Branch**

Multiple Streams

- At each branch, a pipeline has to choose the next instruction to fetch
- A wrong decision is costly
- A brute-force approach:
 - Replicate the initial portions of the pipeline
 - Pipeline can fetch two instructions, making use of two streams
- Drawbacks:
 - With multiple pipelines there are contention delays for access to the registers and to memory
 - Additional branch instructions may enter the pipeline before the original branch decision is resolved

Prefetch Branch Target

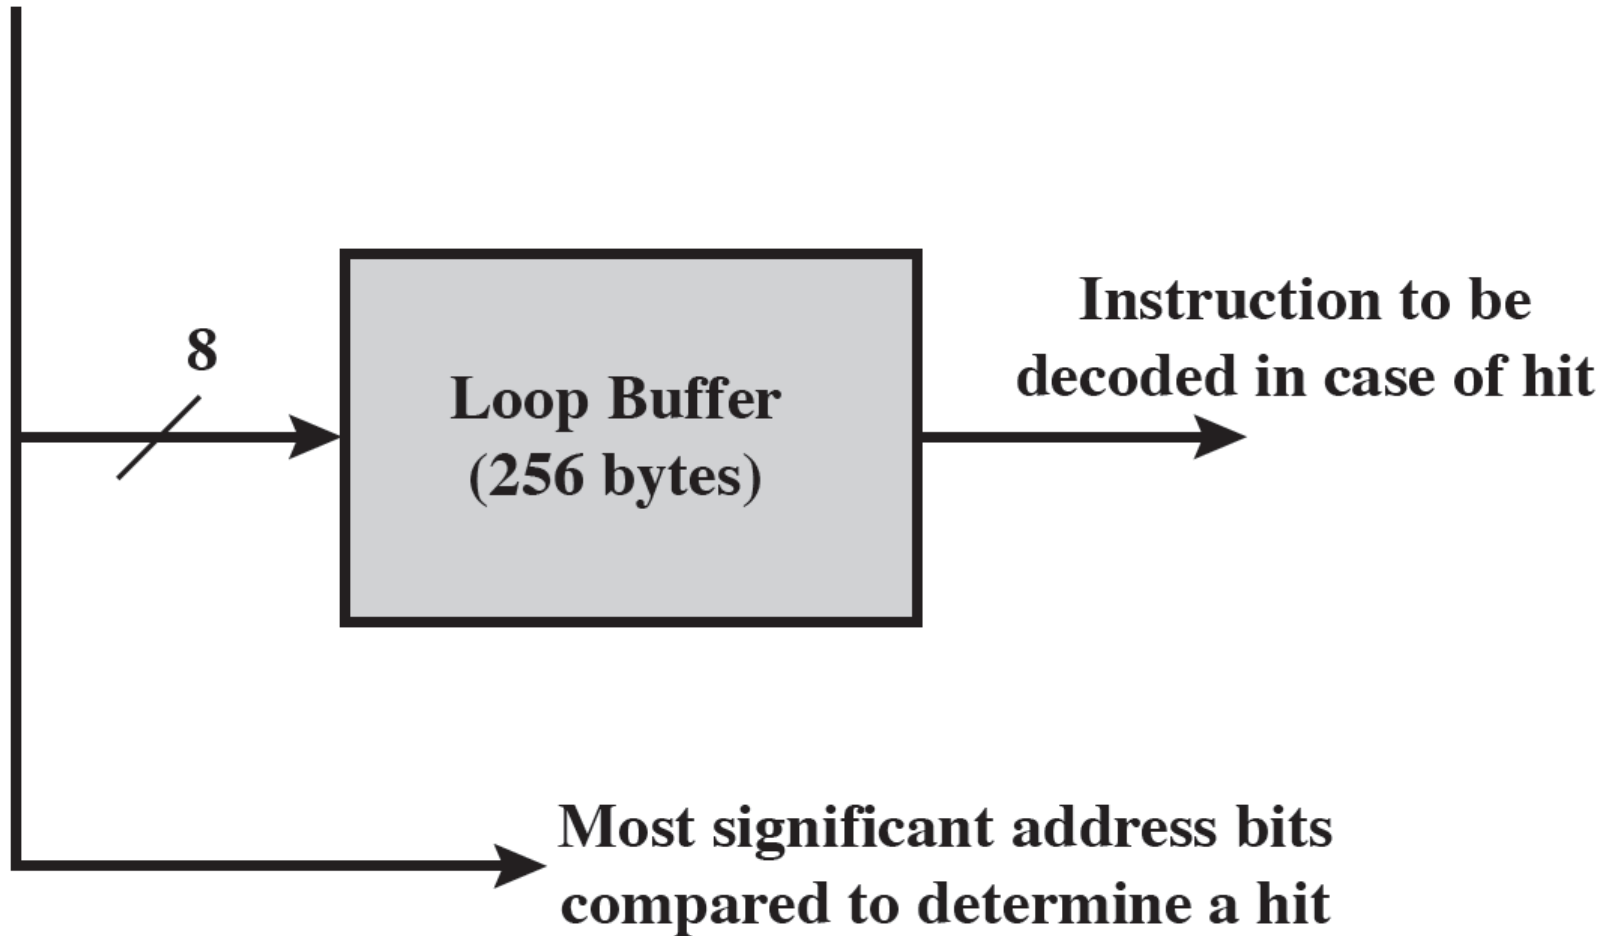
- When a conditional branch is recognized, the target of the branch is prefetched, in addition to the instruction following the branch
- Target is then saved until the branch instruction is executed
- If the branch is taken, the target has already been prefetched
- IBM 360/91 uses this approach

Loop Buffer

- Small, very-high speed memory containing the n most recently fetched instructions, in sequence
- Benefits:
 - Instructions fetched in sequence will be available without the usual memory access time
 - If a branch occurs to a target just a few locations ahead of the address of the branch instruction, the target will already be in the buffer
- This strategy is particularly well suited to dealing with loops
- Similar in principle to a cache dedicated to instructions
- Differences:
 - The loop buffer only retains instructions in sequence
 - Is much smaller in size and hence lower in cost

Loop Buffer

Branch address



Branch Prediction

- **Static Approaches**
 - Predict never taken
 - Predict always taken
 - Predict by opcode
- **Dynamic Approaches**
 - Taken/not taken switch
 - Branch history table

Static Branch Prediction

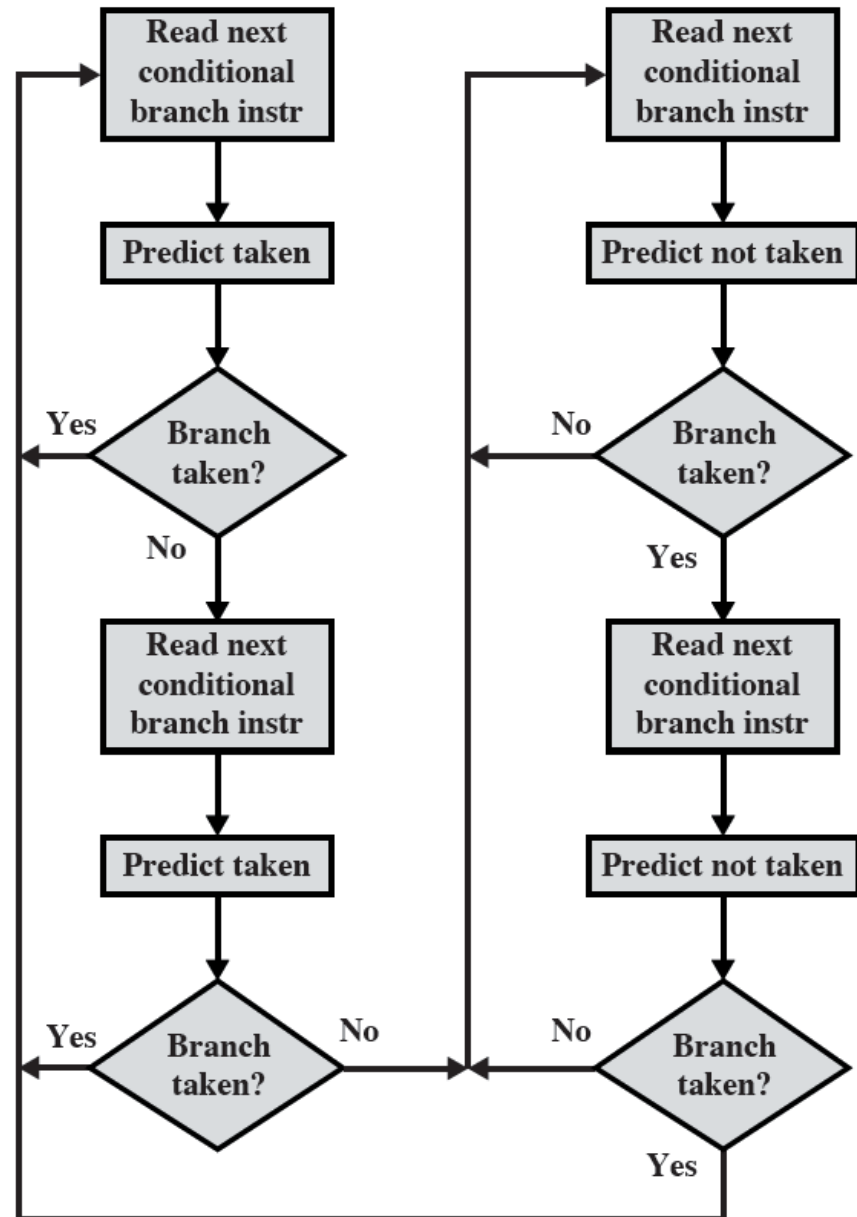
■ Predict Always/Never Taken

- Continue to fetch instructions in sequence or fetch from the branch target
- The predict-never-taken approach is the most popular of all the branch prediction methods
- Analyzing program behavior shows that conditional branches are taken more than 50% of the time
- However, in a paged machine, prefetching the branch target is more likely to cause a page fault than prefetching the next instruction in sequence

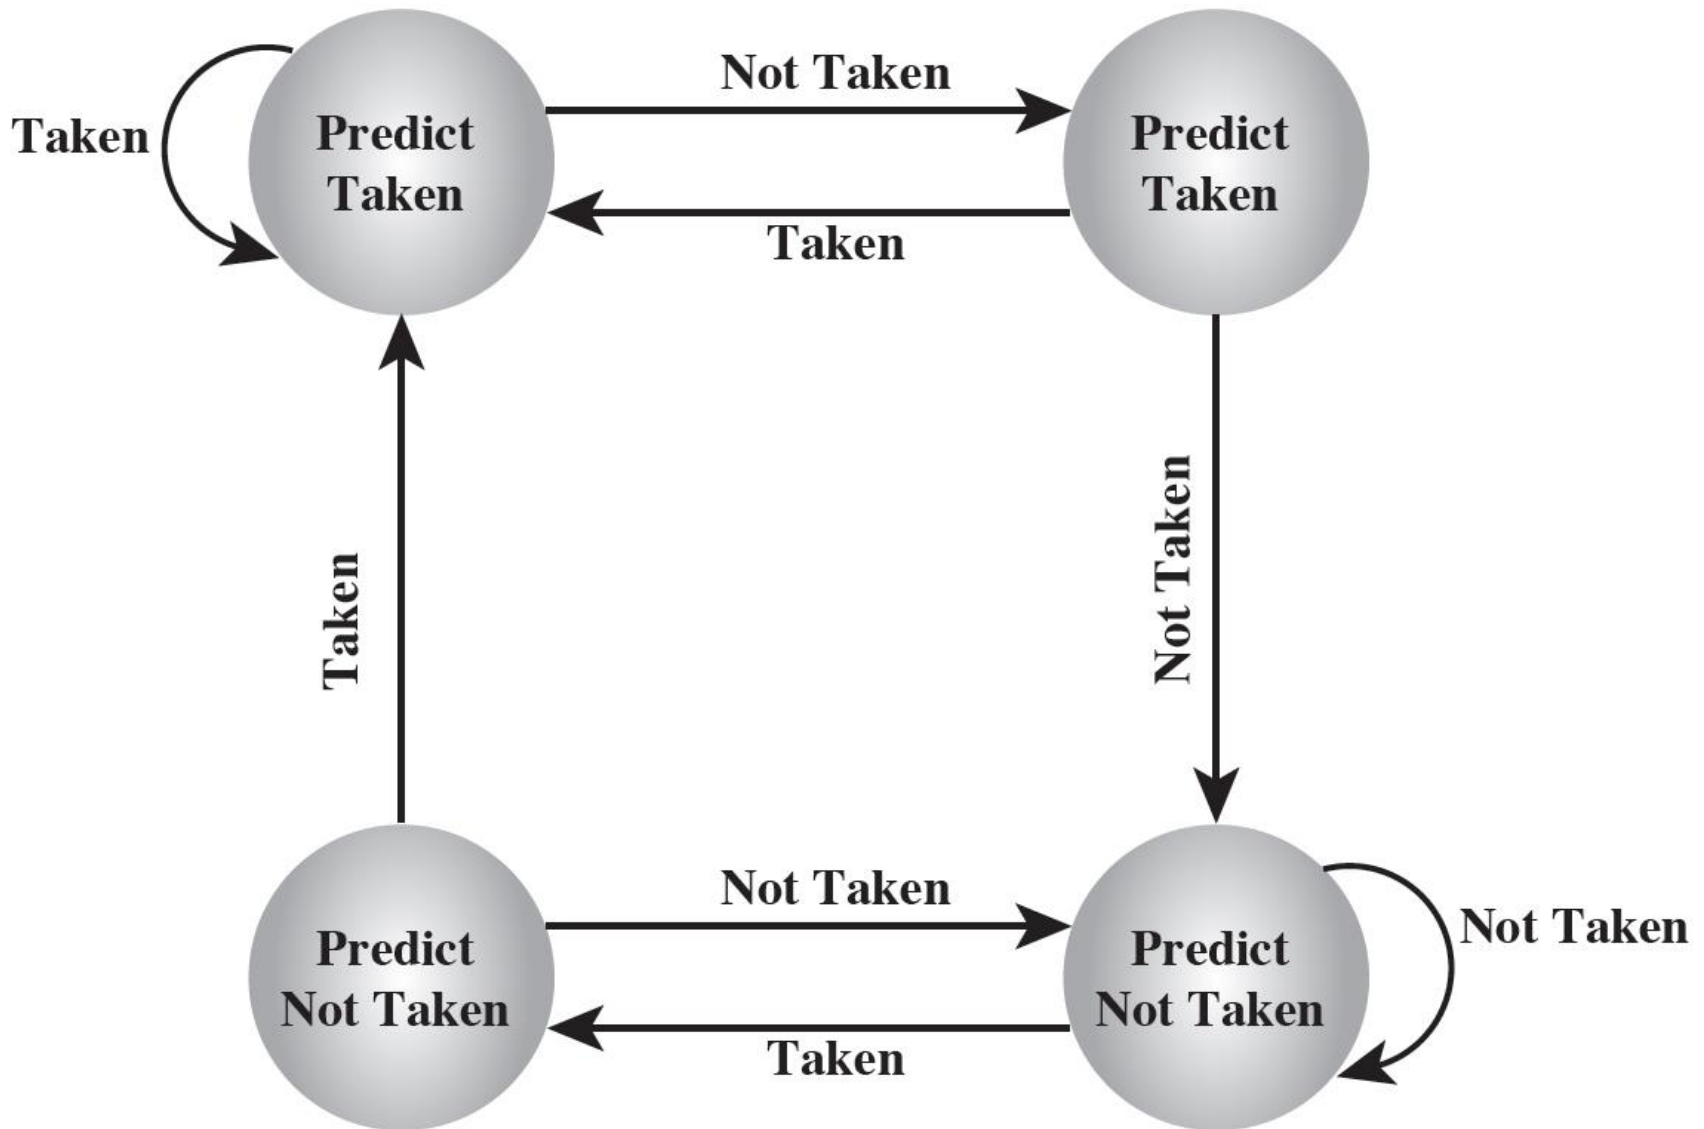
■ Opcode Dependent

- The processor assumes that the branch will be taken for certain branch opcodes and not for others
- Success rates of greater than 75% with this strategy.


Branch Prediction Flow Chart



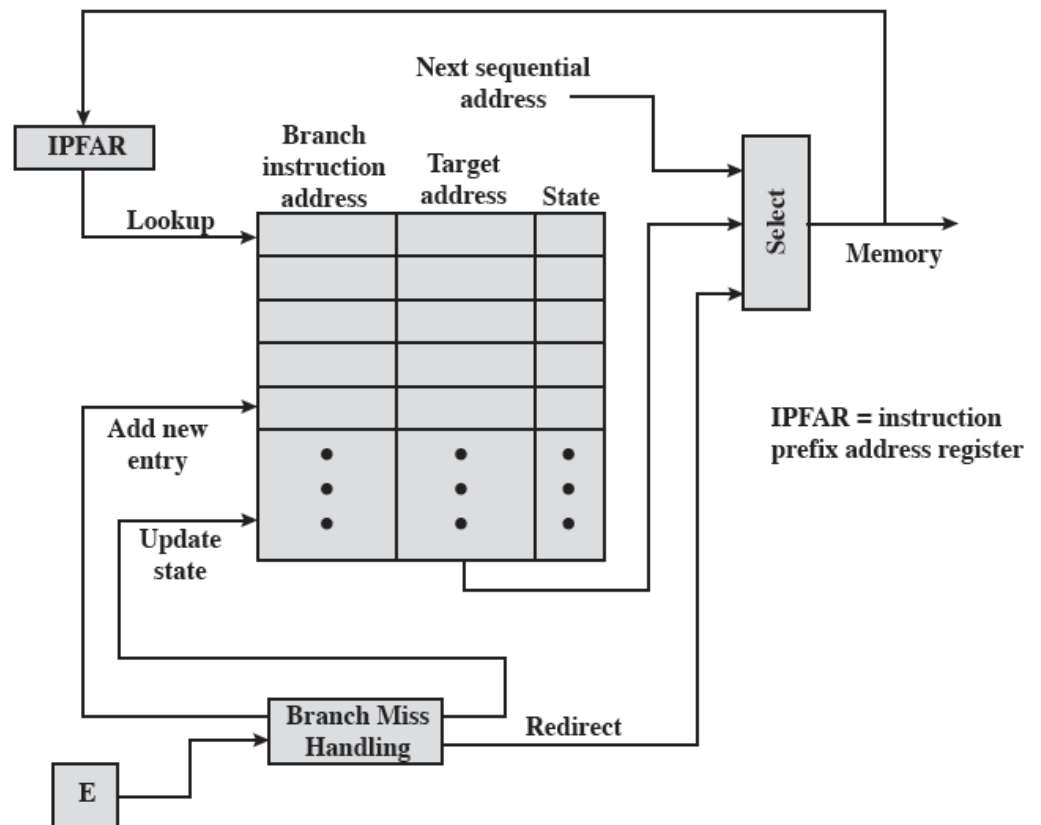
Branch Prediction State Diagram



Static vs. Dynamic Branch Prediction



(a) Predict never taken strategy





Processor Structure and Function

Lecture Content

- Processor Organization
- Pipelining & Branch Prediction
- **Intel Pipelining & Processor Structure**

Intel 80486 Pipeline

■ Fetch

- Operates independently of the other stages to keep the prefetch buffers full

■ Decode Stage 1

- All opcode and addressing-mode information is decoded in the D1 stage
- 3 bytes of instruction are passed to the D1 stage from the prefetch buffers

■ Decode Stage 2

- Expands each opcode into control signals for the ALU
- Also controls the computation of the more complex addressing modes

■ Execute

- Stage includes ALU operations, cache access, and register update


■ Write Back

- Updates registers and status flags modified during the preceding execute stage


80486 Pipeline Examples



(a) No Data Load Delay in the Pipeline



(b) Pointer Load Delay



(c) Branch Instruction Timing


Integer Registers

32 Bit			
Type	Number	Length (bits)	Purpose
General	8	32	General-purpose user registers
Segment	6	16	Contain segment selectors
EFLAGS	1	32	Status and control bits
Instruction Pointer	1	32	Instruction pointer
64 Bit			
Type	Number	Length (bits)	Purpose
General	16	64	General-purpose user registers
Segment	6	16	Contain segment selectors
RFLAGS	1	64	Status and control bits
Instruction Pointer	1	64	Instruction pointer

Floating Point & Additional Registers

Type	Number	Length (bits)	Purpose
Numeric	8	80	Hold floating-point numbers
Control	1	16	Control bits
Status	1	16	Status bits
Tag Word	1	16	Specifies contents of numeric registers
Instruction Pointer	1	48	Points to instruction interrupted by exception
Data Pointer	1	48	Points to operand interrupted by exception

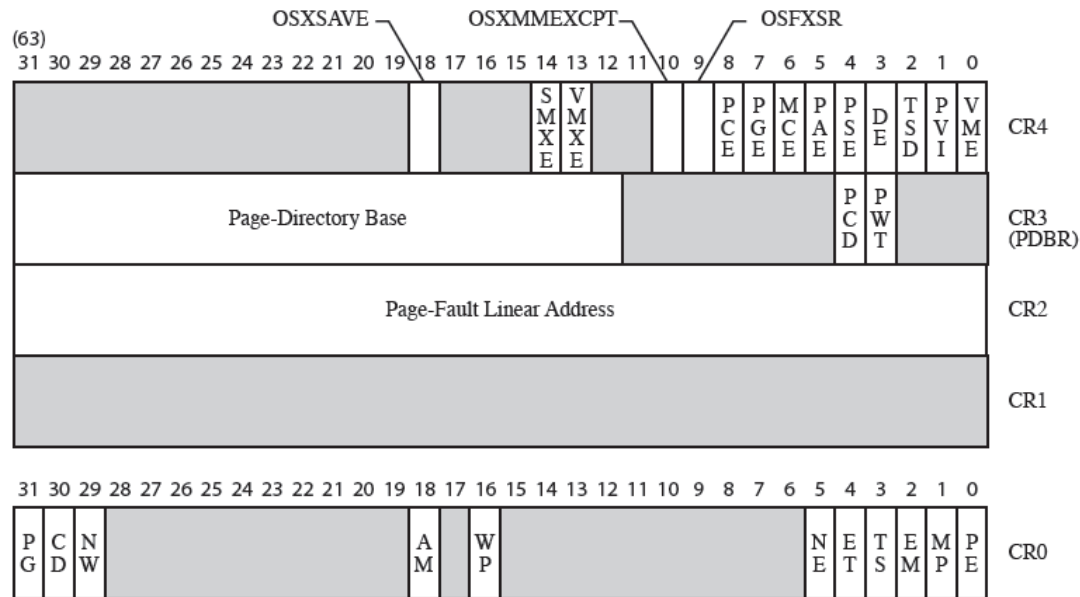
EFLAGS Register



ID = Identification flag
 VIP = Virtual interrupt pending
 VIF = Virtual interrupt flag
 AC = Alignment check
 VM = Virtual 8086 mode
 RF = Resume flag
 NT = Nested task flag
 IOPL = I/O privilege level
 OF = Overflow flag

DF = Direction flag
 IF = Interrupt enable flag
 TF = Trap flag
 SF = Sign flag
 ZF = Zero flag
 AF = Auxiliary carry flag
 PF = Parity flag
 CF = Carry flag

Status Registers



shaded area indicates reserved bits

OSXSAVE	=	XSAVE enable bit	PCD	=	Page-level Cache Disable
SMXE	=	Enable Safer mode extensions	PWT	=	Page-level Writes Transparent
VMXE	=	Enable virtual machine extensions	PG	=	Paging
OSXMMEXCPT	=	Support unmasked SIMD FP exceptions	CD	=	Cache Disable
OSFXSR	=	Support FXSAVE, FXSTOR	NW	=	Not Write Through
PCE	=	Performance Counter Enable	AM	=	Alignment Mask
PGE	=	Page Global Enable	WP	=	Write Protect
MCE	=	Machine Check Enable	NE	=	Numeric Error
PAE	=	Physical Address Extension	ET	=	Extension Type
PSE	=	Page Size Extensions	TS	=	Task Switched
DE	=	Debug Extensions	EM	=	Emulation
TSD	=	Time Stamp Disable	MP	=	Monitor Coprocessor
PVI	=	Protected Mode Virtual Interrupt	PE	=	Protection Enable
VME	=	Virtual 8086 Mode Extensions			

x86 Extension Registers

- **Multi Media Extensions (MMX)**
 - 8 Integer Registers, 64 Bit (mapped registers)
 - 24 new instruction
- **Streaming SIMD Extensions (SSE)**
 - 8 new floating point registers (128 Bit)
 - With amd64 architecture 8 additional 128 Bit registers
 - 70 new Instructions
- **SSE2-4**
 - Additional instructions



Mapping of the MMX Registers

Floating-Point

Tag



Floating-Point Registers



MMX Registers