

Software Architecture

Henrik Møller Hansen

Agenda

- About me and my work
- Some general thoughts on software architecture
- Case walkthrough – Flight plan filing



About me

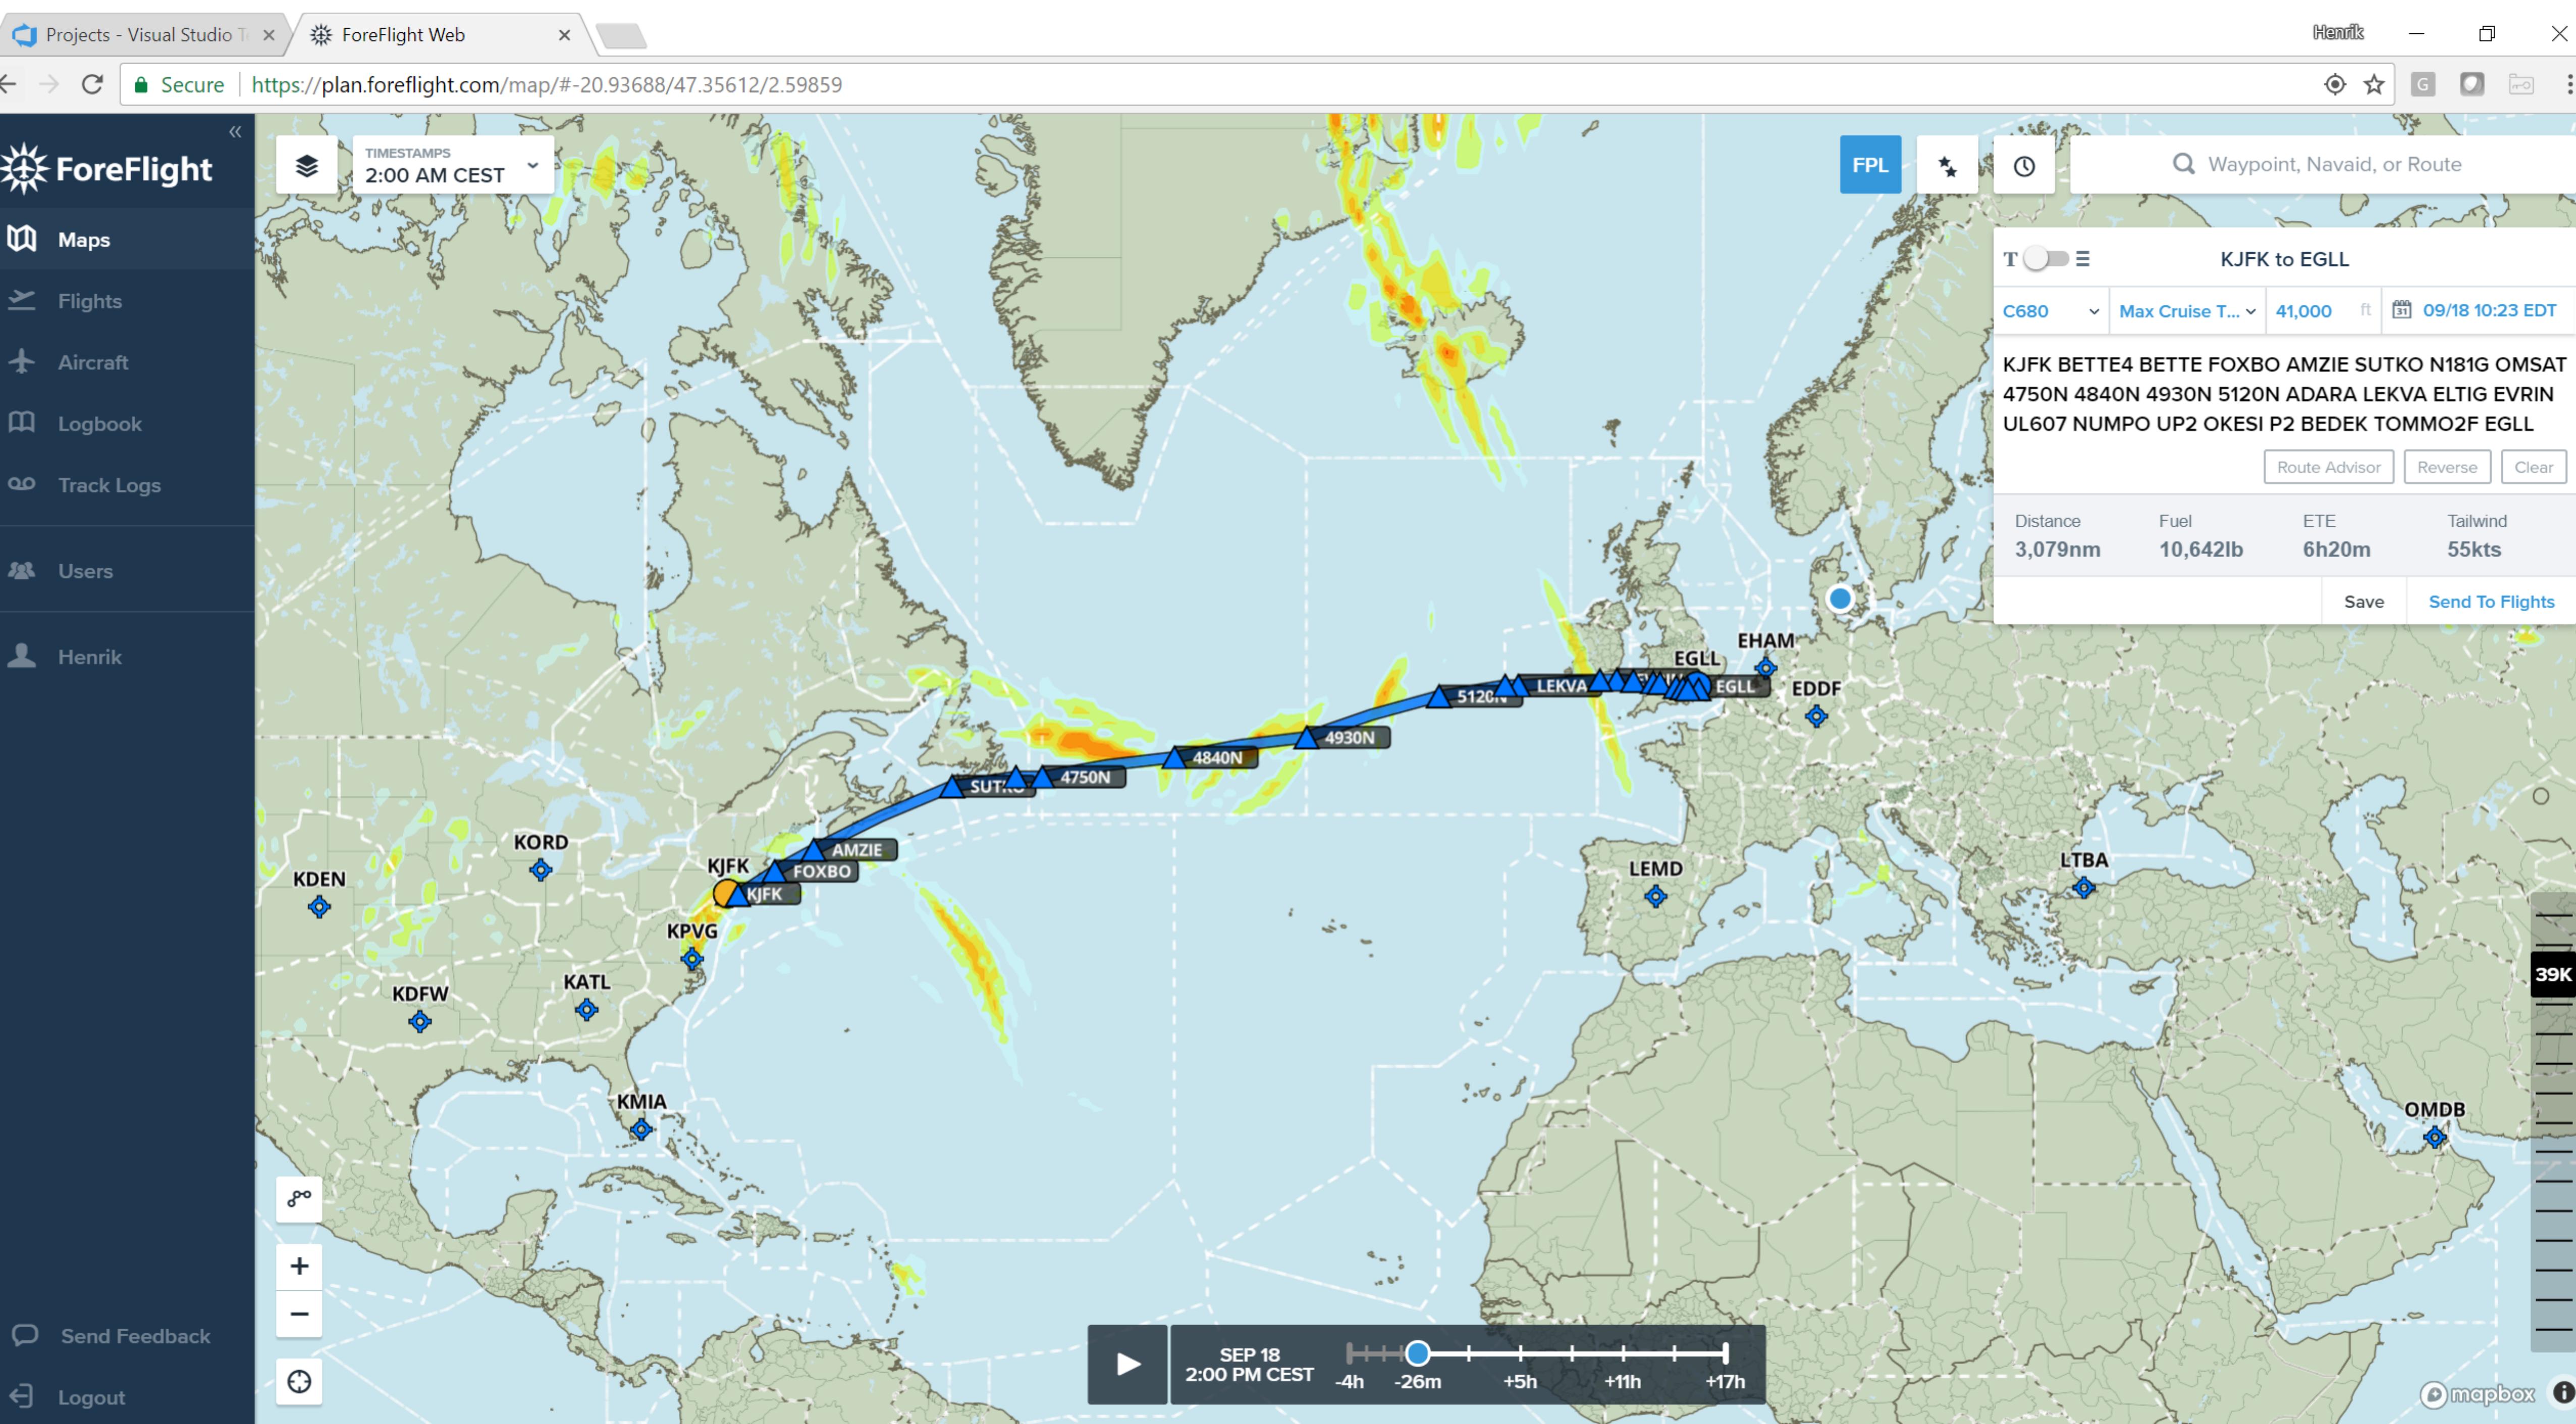
- Education
 - Bogense skole, 1996-2006
 - Kold College HTX, 2006-2009
 - SDU, B.Eng, IT/Comm Engineer, 2009-2013
 - SDU, MS.c, Software Engineering 2013-2016
- Work
 - Student-dev/intern @ AIR SUPPORT 2011-2013
 - Lead-developer @ AIR SUPPORT 2013-2014
 - Head of Development @ AIR SUPPORT 2014-2015
 - CTO @ AviationCloud, 2012-Now
 - AviationCloud aquired by ForeFlight 2016

What does AC do?

- Build systems for *flight planning*
 - Route optimization
 - Fuel calculation algorithms
 - Air Traffic Control Communications systems
 - Runs in the cloud on 40+ servers
 - 13 people in Denmark, 120+ globally



ForeFlight Corporate Jet @ Houston, TX, USA



Whats an API?

- Application programmable interface.
- A well defined interface that allows other systems (or component within the same system) to connect
- Most often this means an HTTP endpoint where you can submit/get some data as JSON or XML.
- Relevant link: <https://swagger.io/>



AVIATION CLOUD

Everything needs an API

- Internal APIs
 - Connect server with mobile apps
 - Connect one tier of servers to another
- External APIs
 - Reddit: API allows you to retrieve and create posts, retrieve user data etc
 - Google Maps API: Allow you to generate static maps, retrieve driving directions etc
 - Twitter API: Allows you to retrieve and post tweets
 - Essentially, most modern system has some kind of API

AviationClouds product: APIs for flight planning

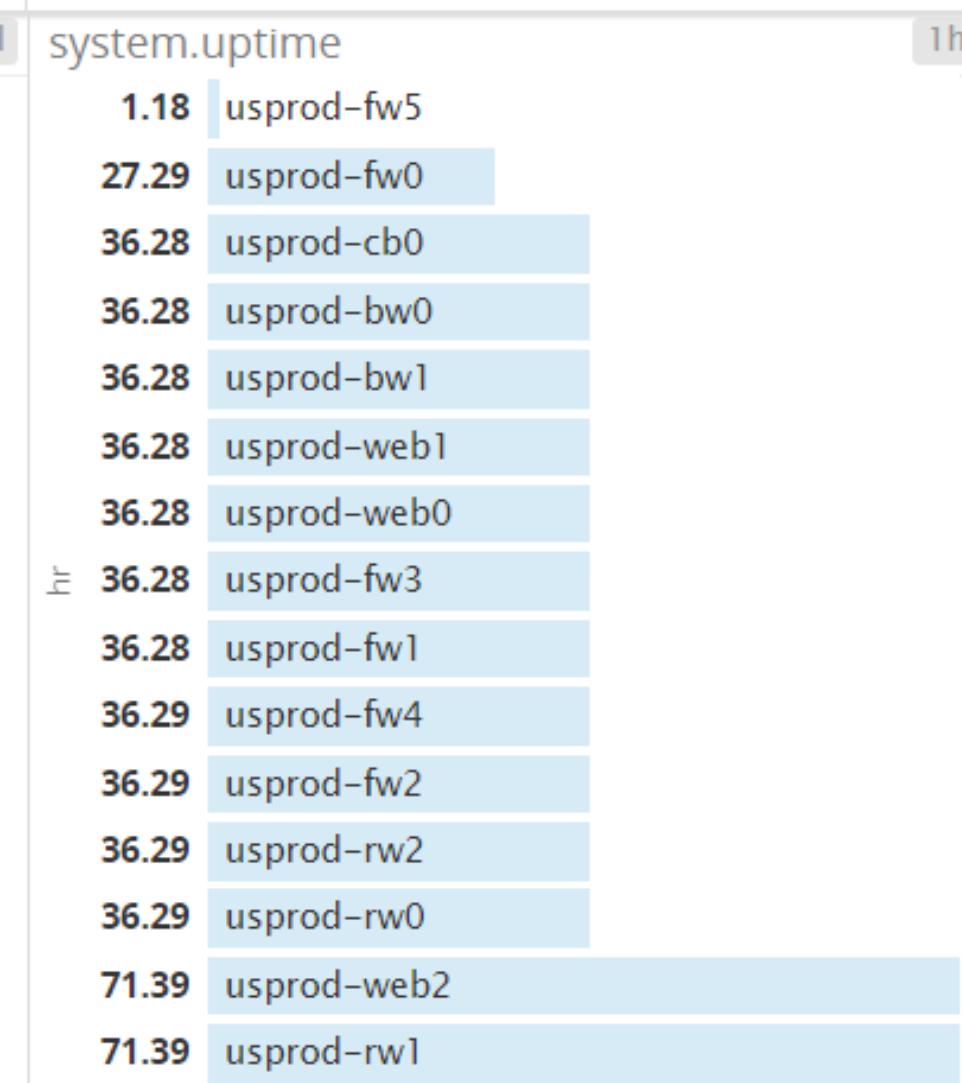
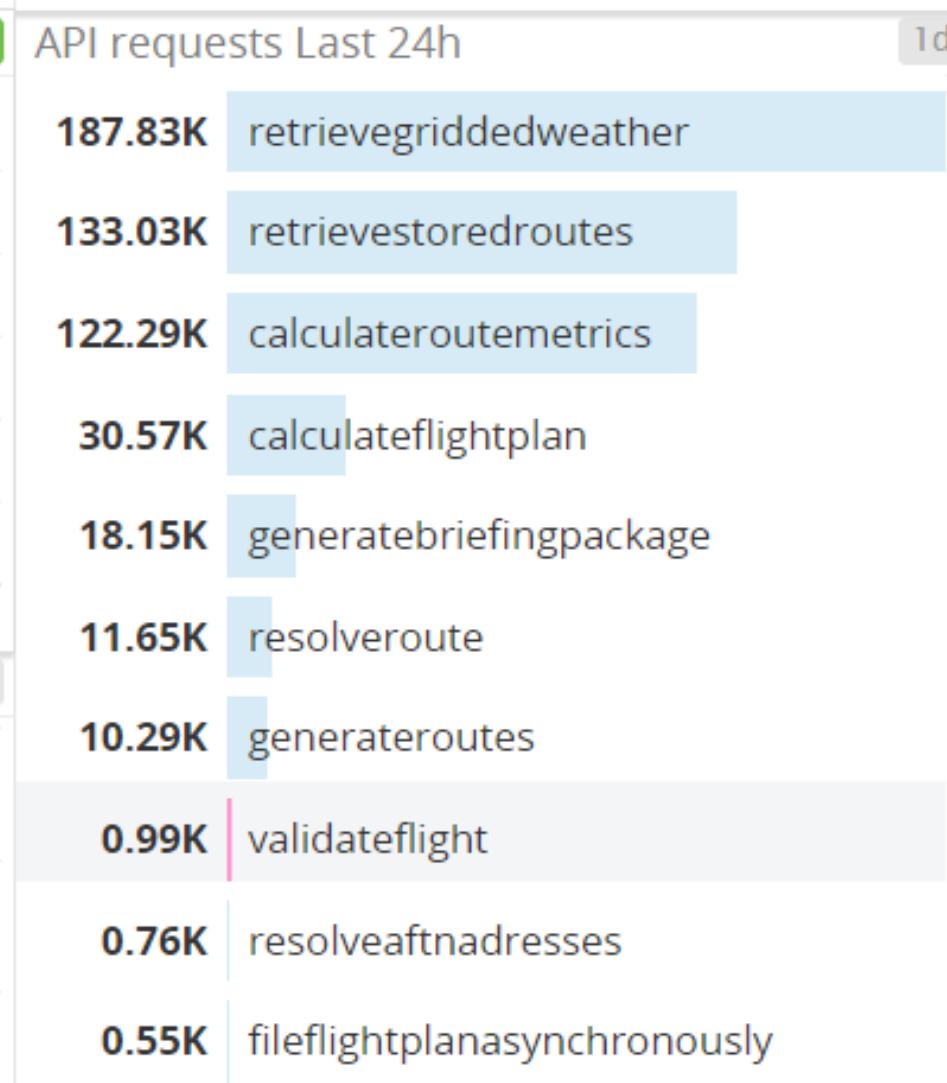
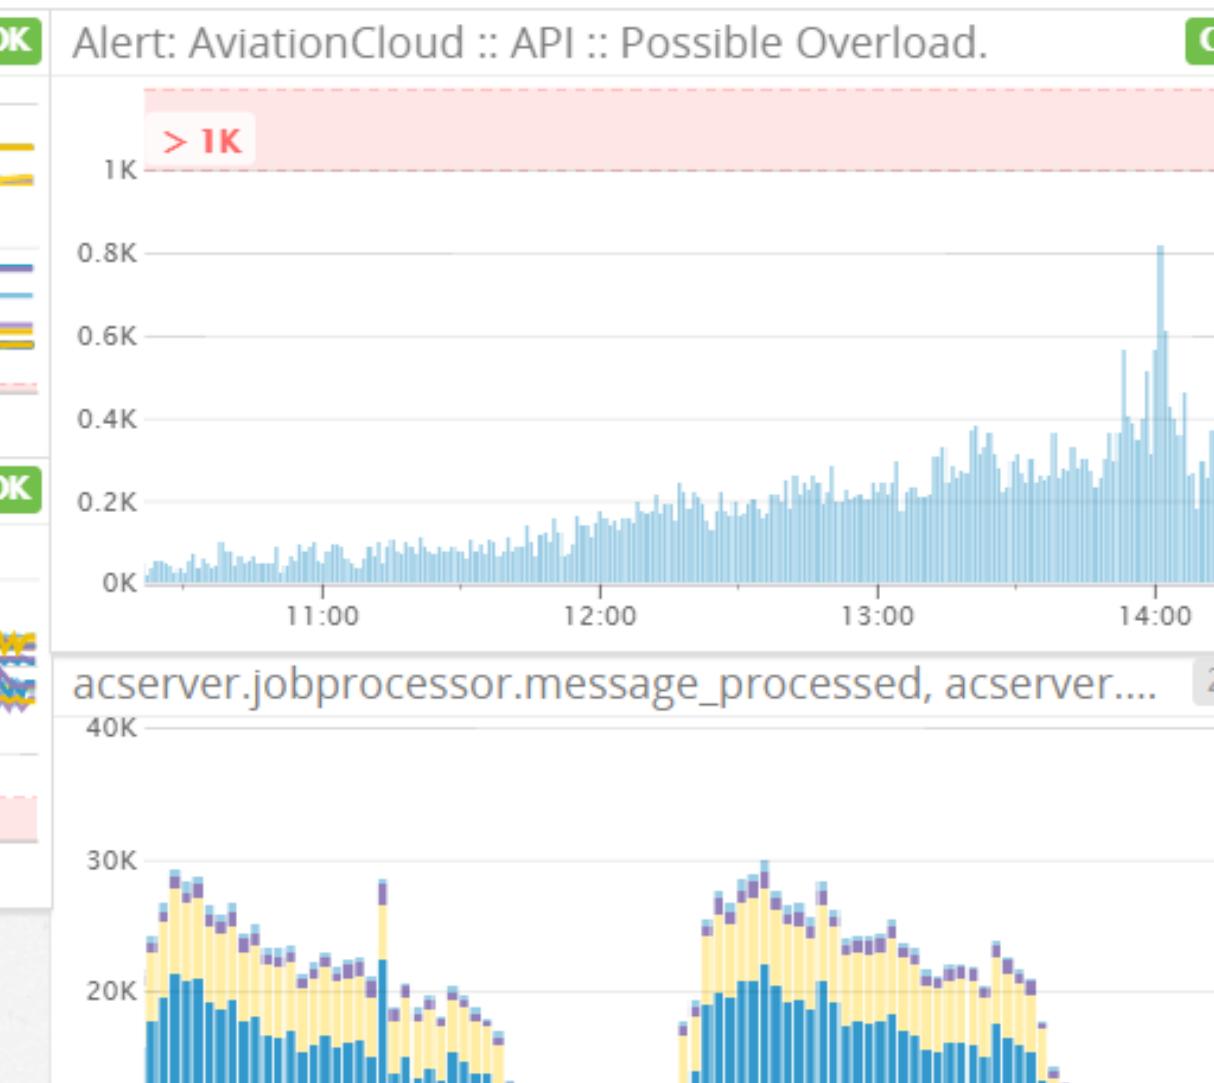
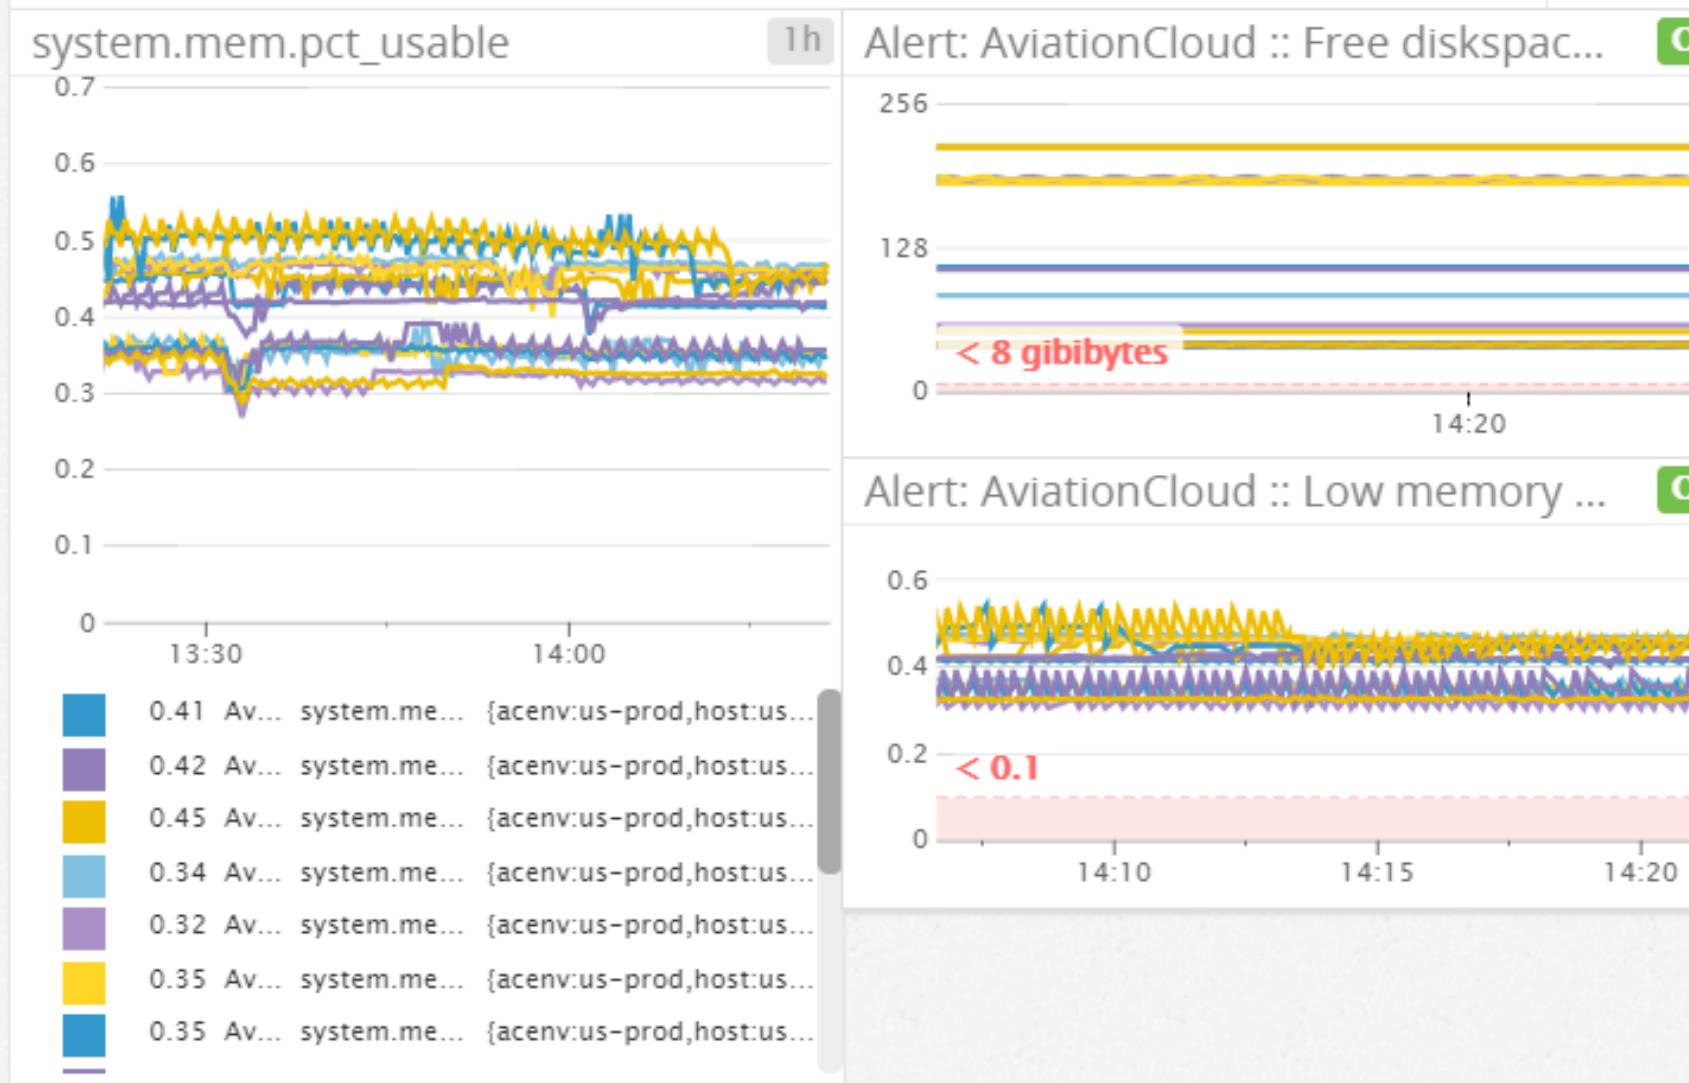
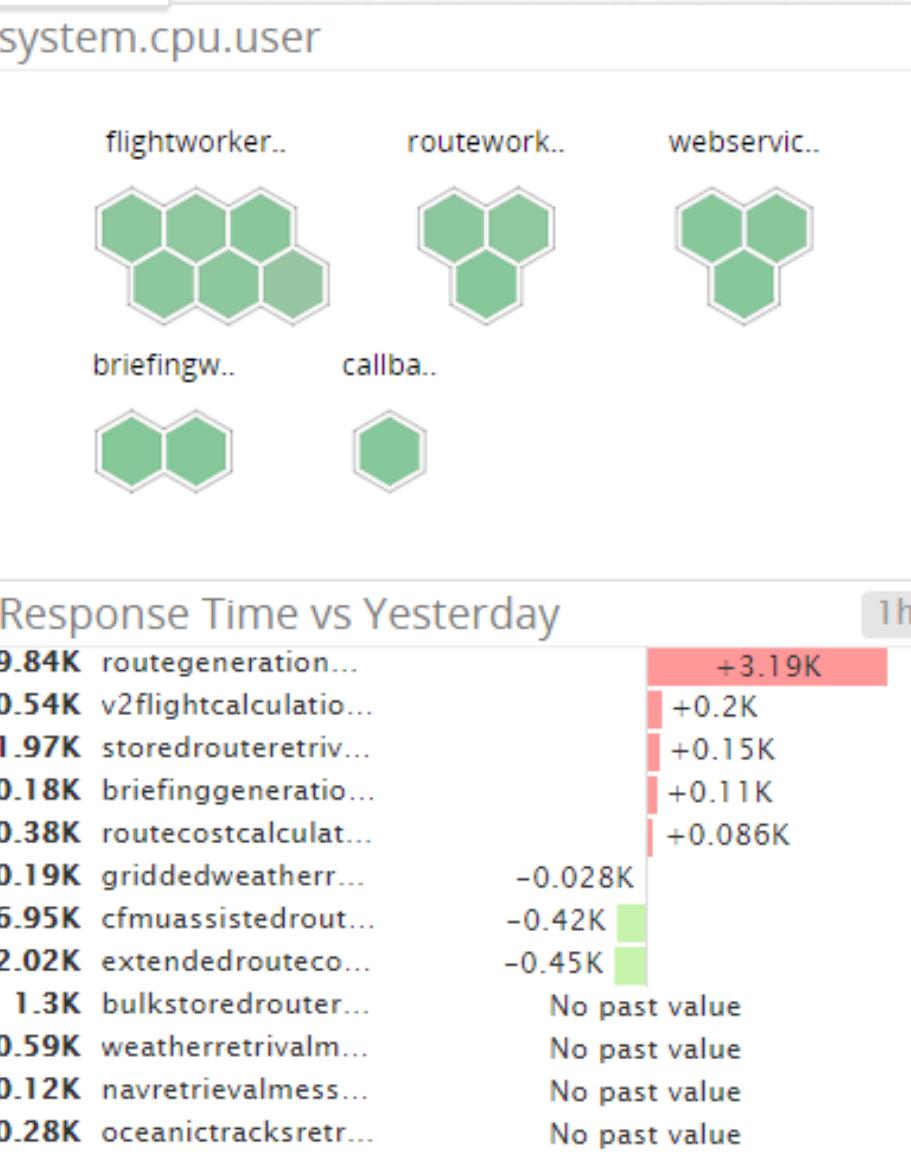
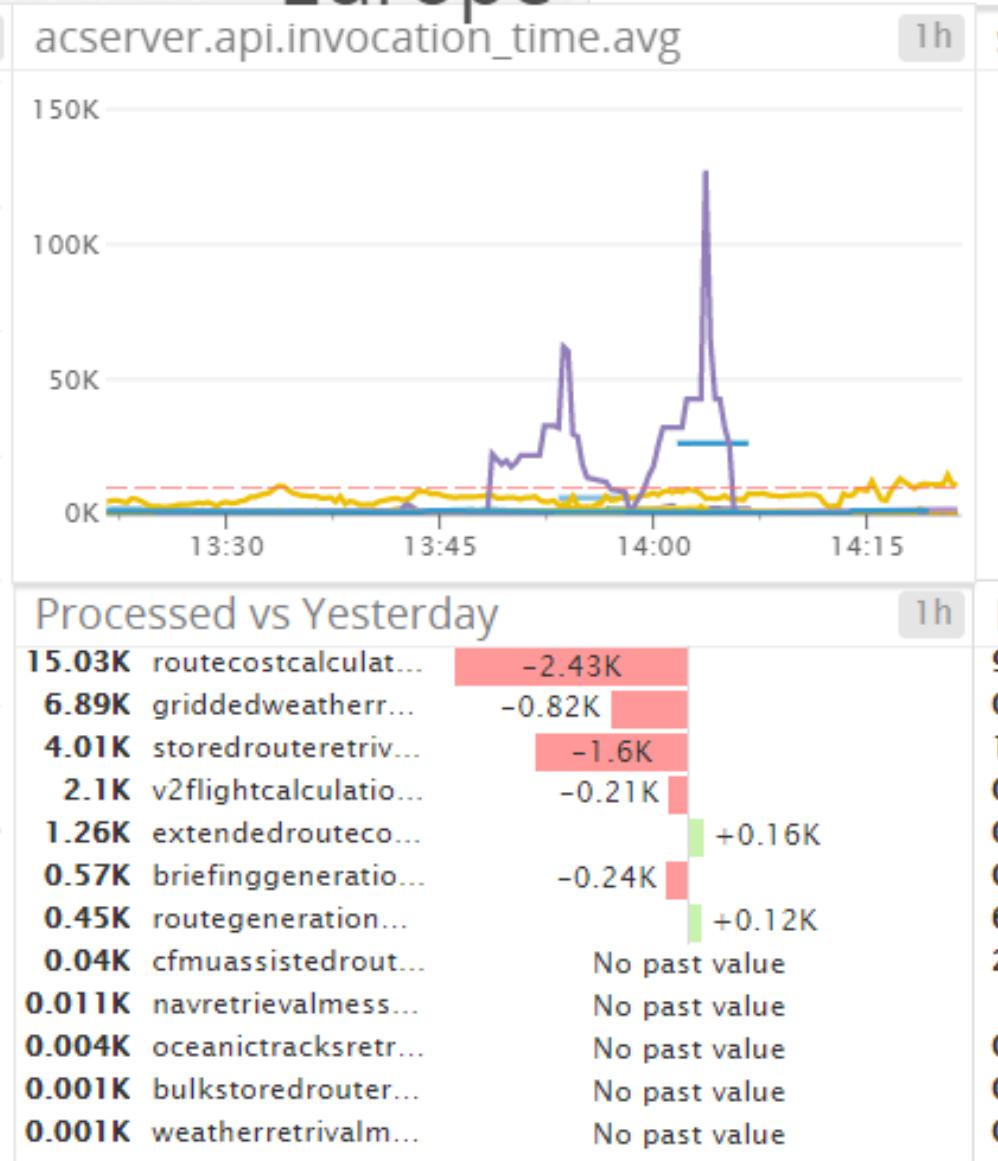
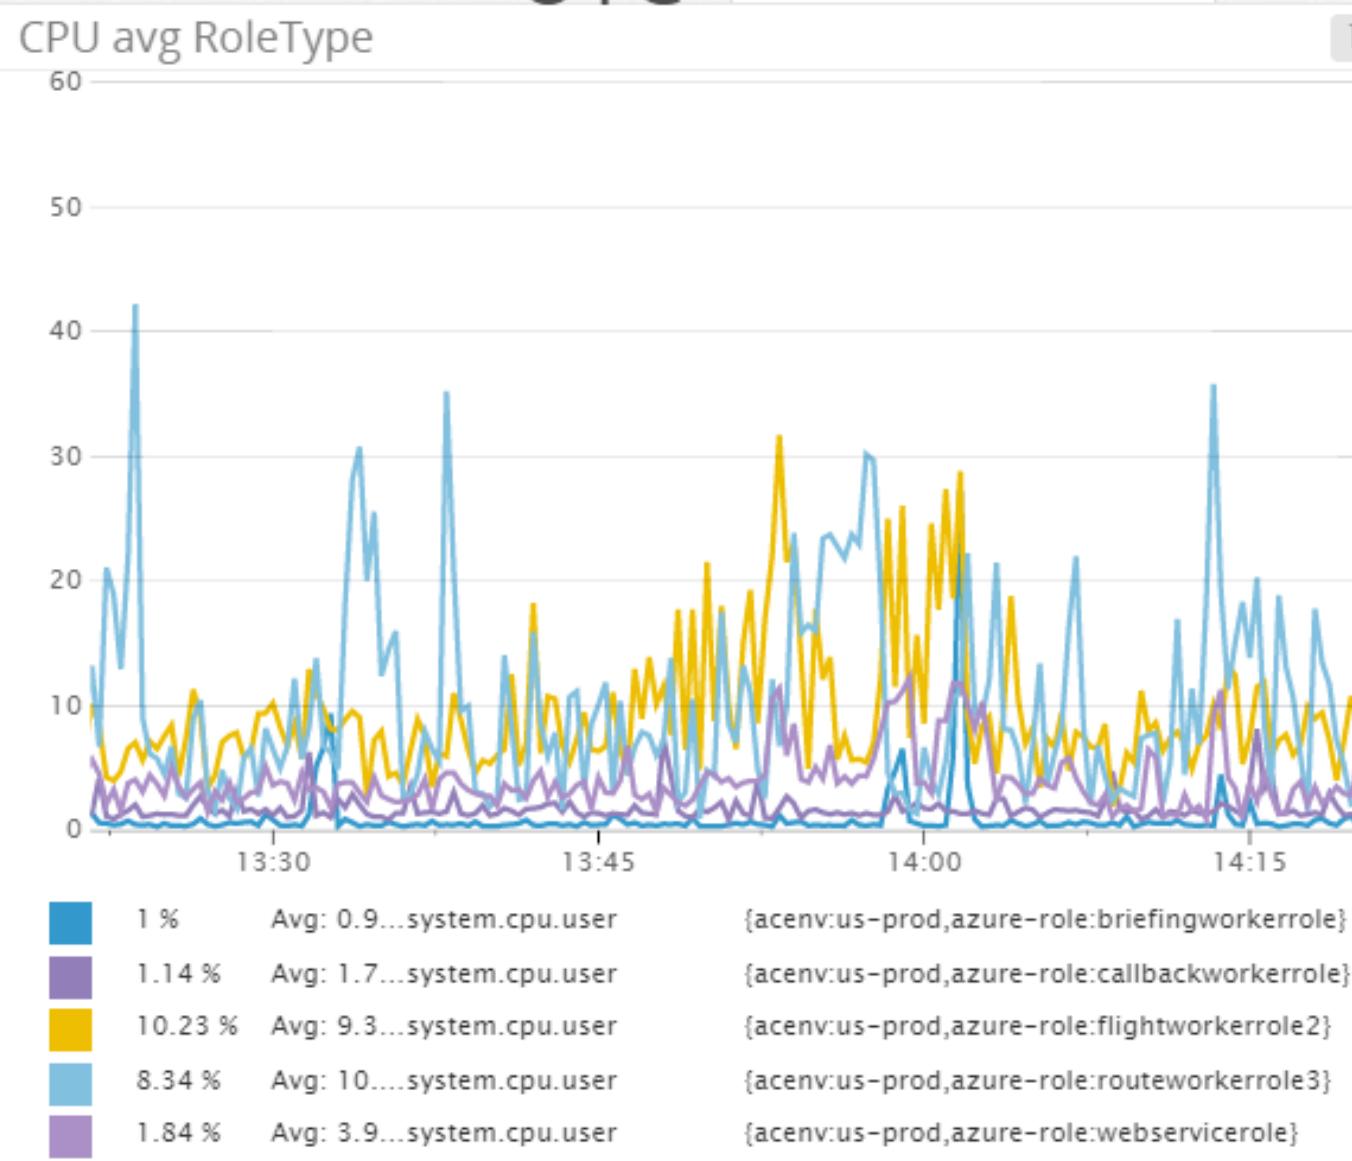
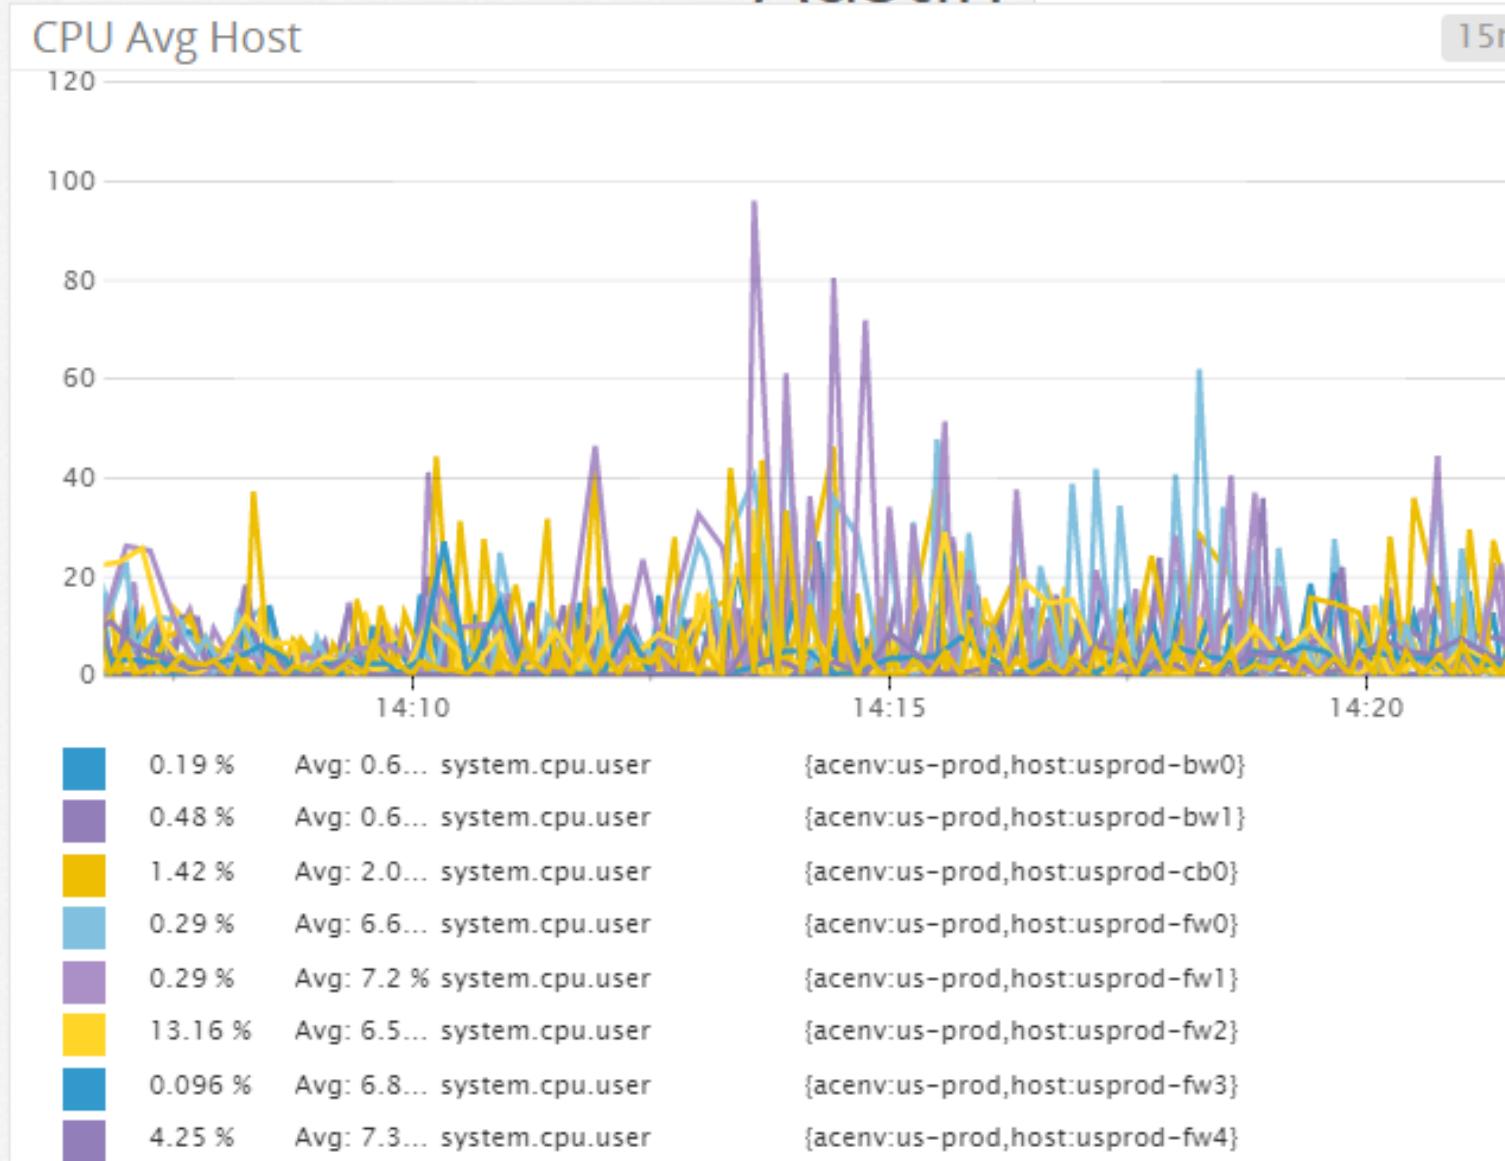
- AviationCloud builds algorithms and systems for flight planning and sells them as APIs
- As we don't have much user interface work, a lot of our work is focused on non-functional requirements
 - High uptime
Our system is a dependency and fundamental part of other systems, so when we go down, they go down.
 - High robustness
We do not control the other systems – our system must be well tested and robust against different types of input
 - Generalization and adaption
As we do not know the ultimate use-cases of our system and algorithms, things must be generalized enough to be able to adapt to a wide range of use-cases
 -

Austin

07:21:56 CDT

UTC

Europe

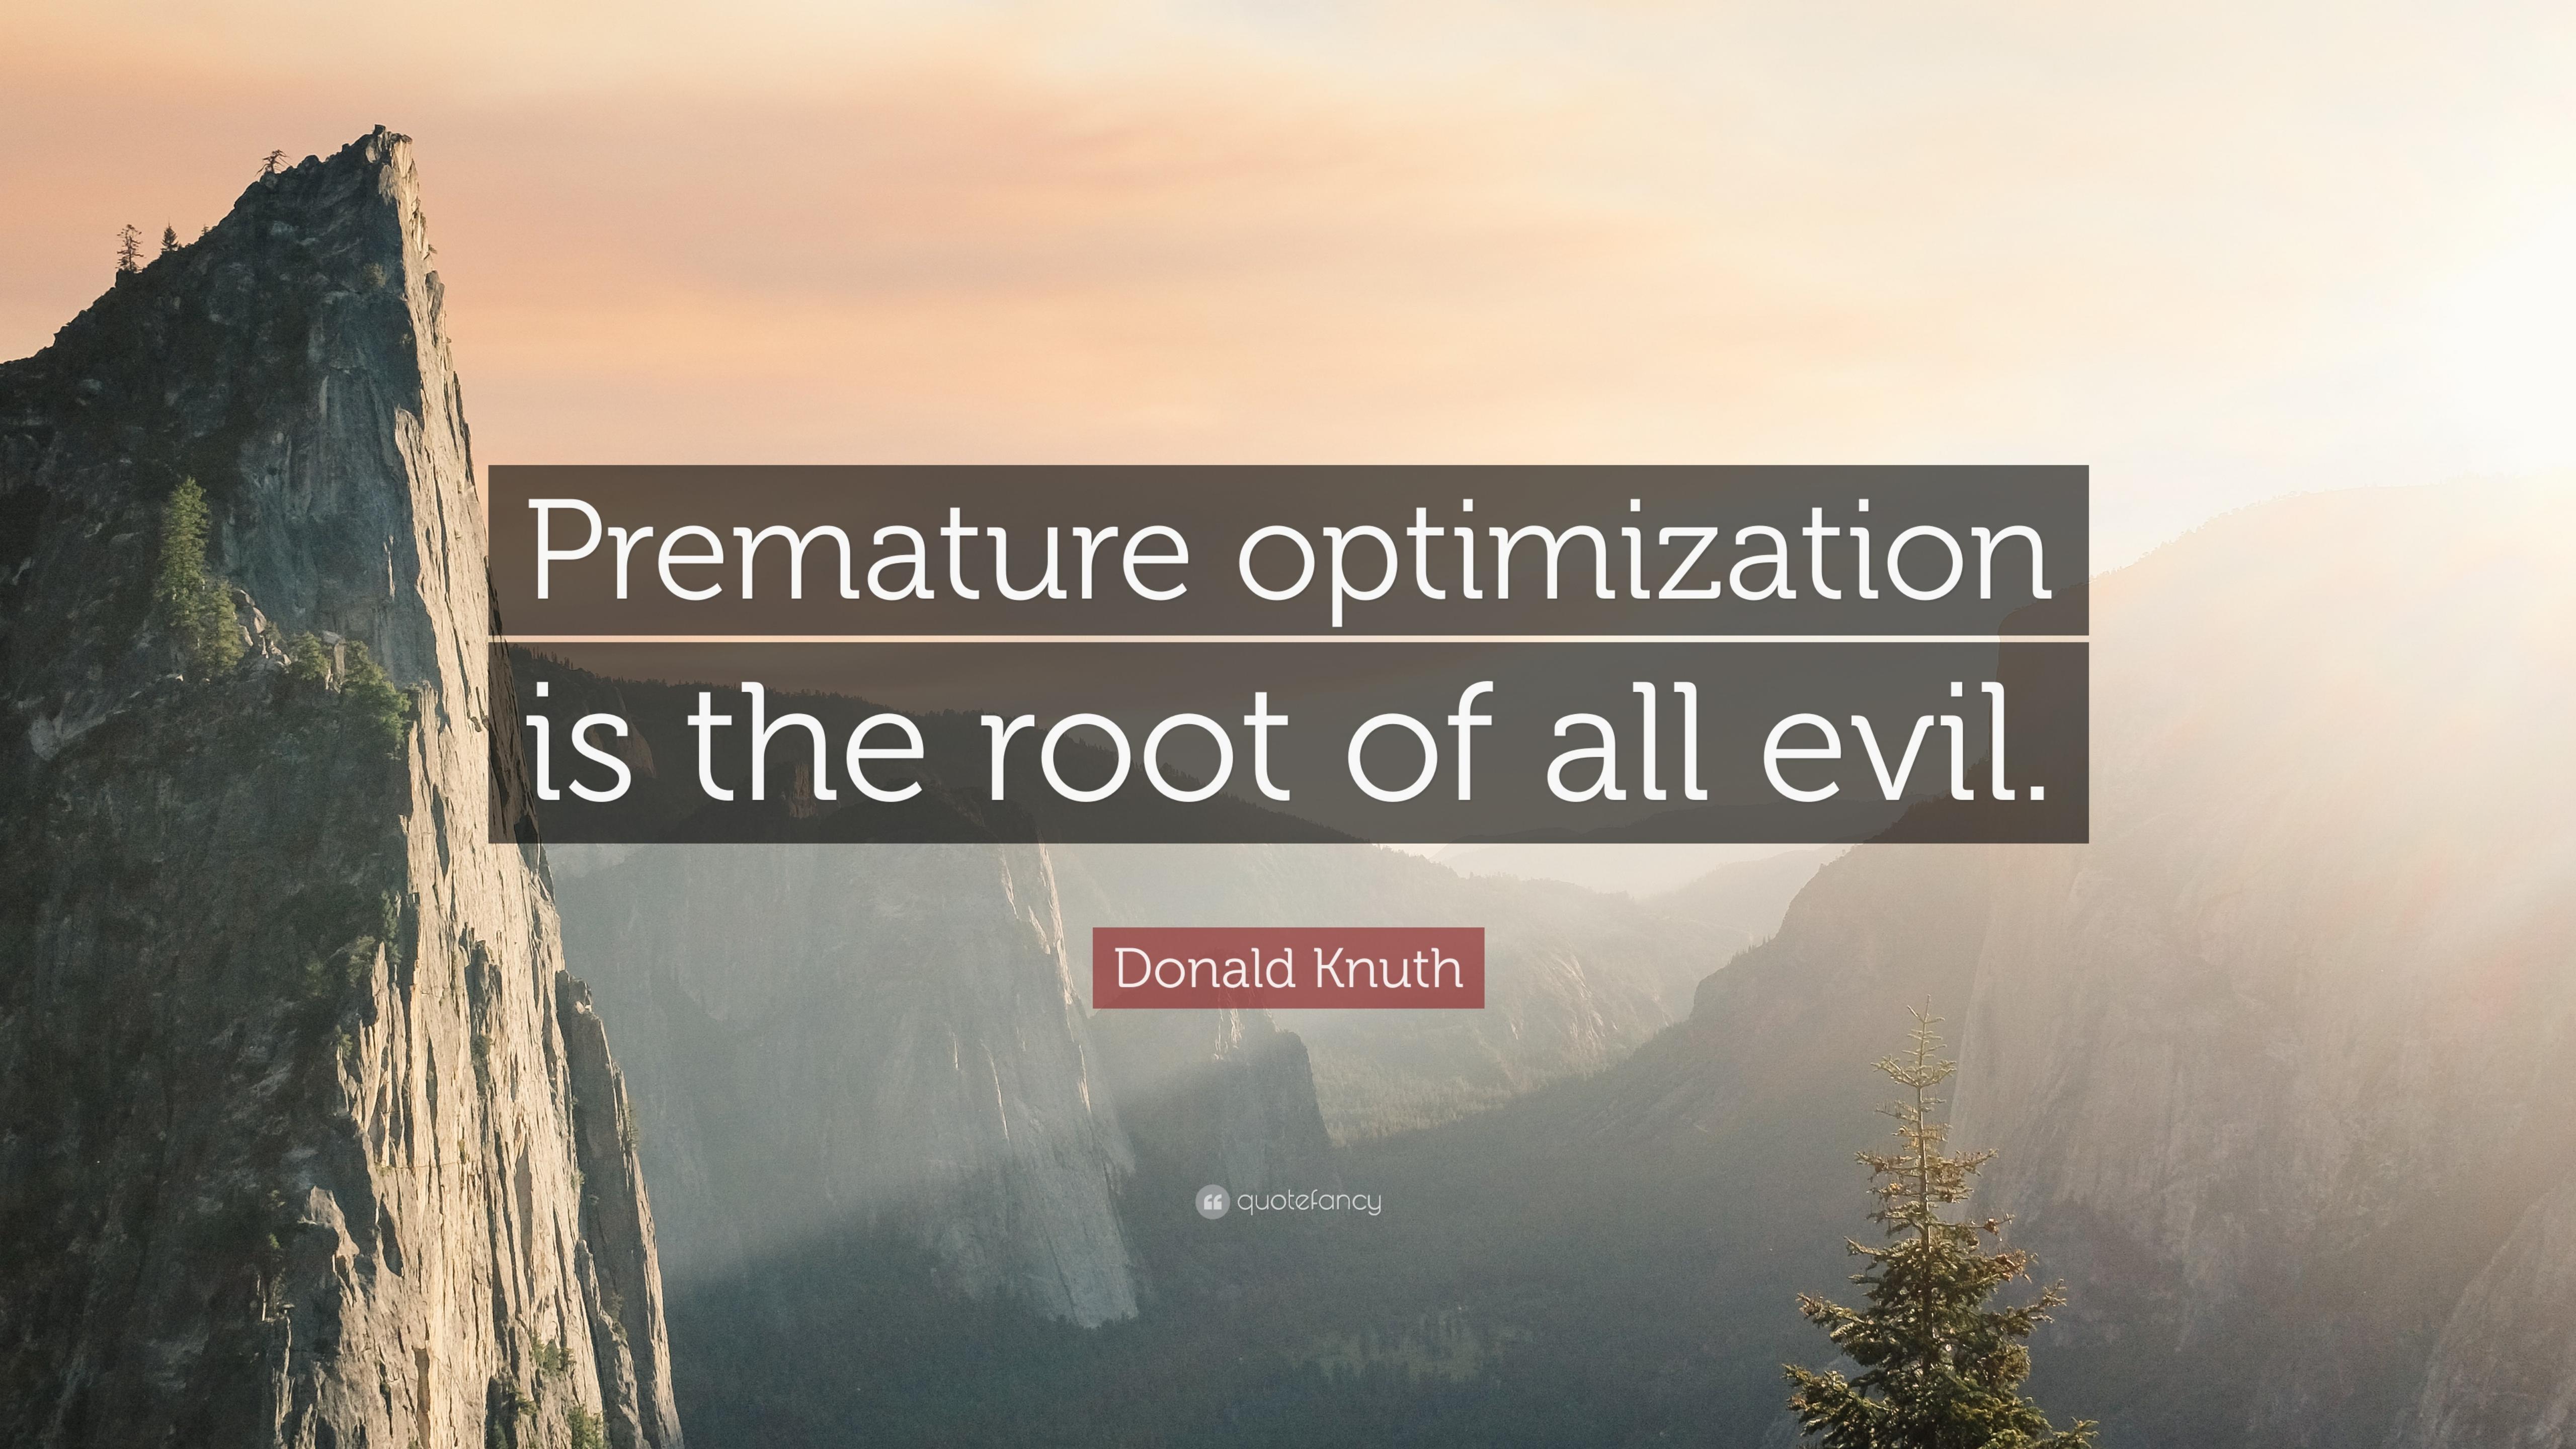
12:21:56 UTC
2:21:56 pm CEST

Software architecture is a mixed field

- It's easy to throw around a lot of buzzwords without having any idea of what you're talking about
- If neglected, it can lead to unmanageable technical debt, leading to unmaintainable and unreliable systems
- If overdone, it leads to inefficient development, ultimately making those systems fail to meet their functional requirements
- Therefore, two guiding principles ... :)

Talk is cheap. Show me the code.

Linus Torvalds



Premature optimization
is the root of all evil.

Donald Knuth

Flight plan filing – An introduction

- Back in the 1960's the aviation industry was booming because of the invention of jet aircrafts
- Countries needed a way to communicate about flights – and this was before the internet
- The solution: Aeronautical Fixed Telecommunication Network, AFTN



The Boeing 707
Launched in 1958 marked the beginning of the jet-age

An AFTN message

- AFTN: Maximum, max 69 characters per line, max 1104 characters
- Below is an example of a flight plan message (FPL) describing a flight from Malage to Copenhagen
- These messages are distributed to every country along the route of flight and all relevant airports

ICAO FPL: (FPL-TAILNO-IG
-SR22/L-GLORUVXWY/N
-LEMG1402
-N0174F170 BLN1L BLN/N0172F160 G5 CJN R10 PPN R299 BEGUY/N0174F170
DCT SOVOS DCT BASTO DCT SECHE DCT BUGUS DCT TUGLI DCT LASEV DCT
RISUN/N0172F160 G21 GIVOR V36 AKELU N852 DIK/N0174F170 N853 ARCKY
Z907 ELDAR Z727 HMM M170 STADE L23 HAM P605 GESKA Z711 MONAK MONAK1D
-EKCH0835
-PBN/B2 DOF/170918 EET/LFBB0252 LFFF0432 LFEE0459 EDGG0555 EBBU0556
EDGG0609 EBBU0610 EDGG0620 EBBU0620 EDGG0622 EDWW0701 EKDK0805)

Flight plan filing system

- Our system generates messages which gets sent to air traffic control, we denote these **outgoing messages**
- We (might) receive replies back to these messages conveying important information about the flight. We call these **incoming messages**
- **Outgoing messages** has an internationally standardized format
- **Incoming messages** messages do not

Some bad situations

- If the system is down → Pilots cannot submit flight plan → Aircraft cannot take off → \$\$\$\$\$
- If the system fails to send an outgoing message → Aircraft might be rejected entrance into an airspace → \$\$\$
- If the system loses an incoming message → Pilot might not get important information → \$\$\$ or worse



AVIATION CLOUD

Functional and non-functional requirements

- Functional requirements describes what the system should do from a user point of view
 - The system should allow the user to input X and save that information
- Non functional requirements describe attributes about the system that are not directly visible to the user
 - The system should be able to handle at least X requests per second
 - Non-functional requirements are often some of the main drivers when designing software architecture

Functional requirements

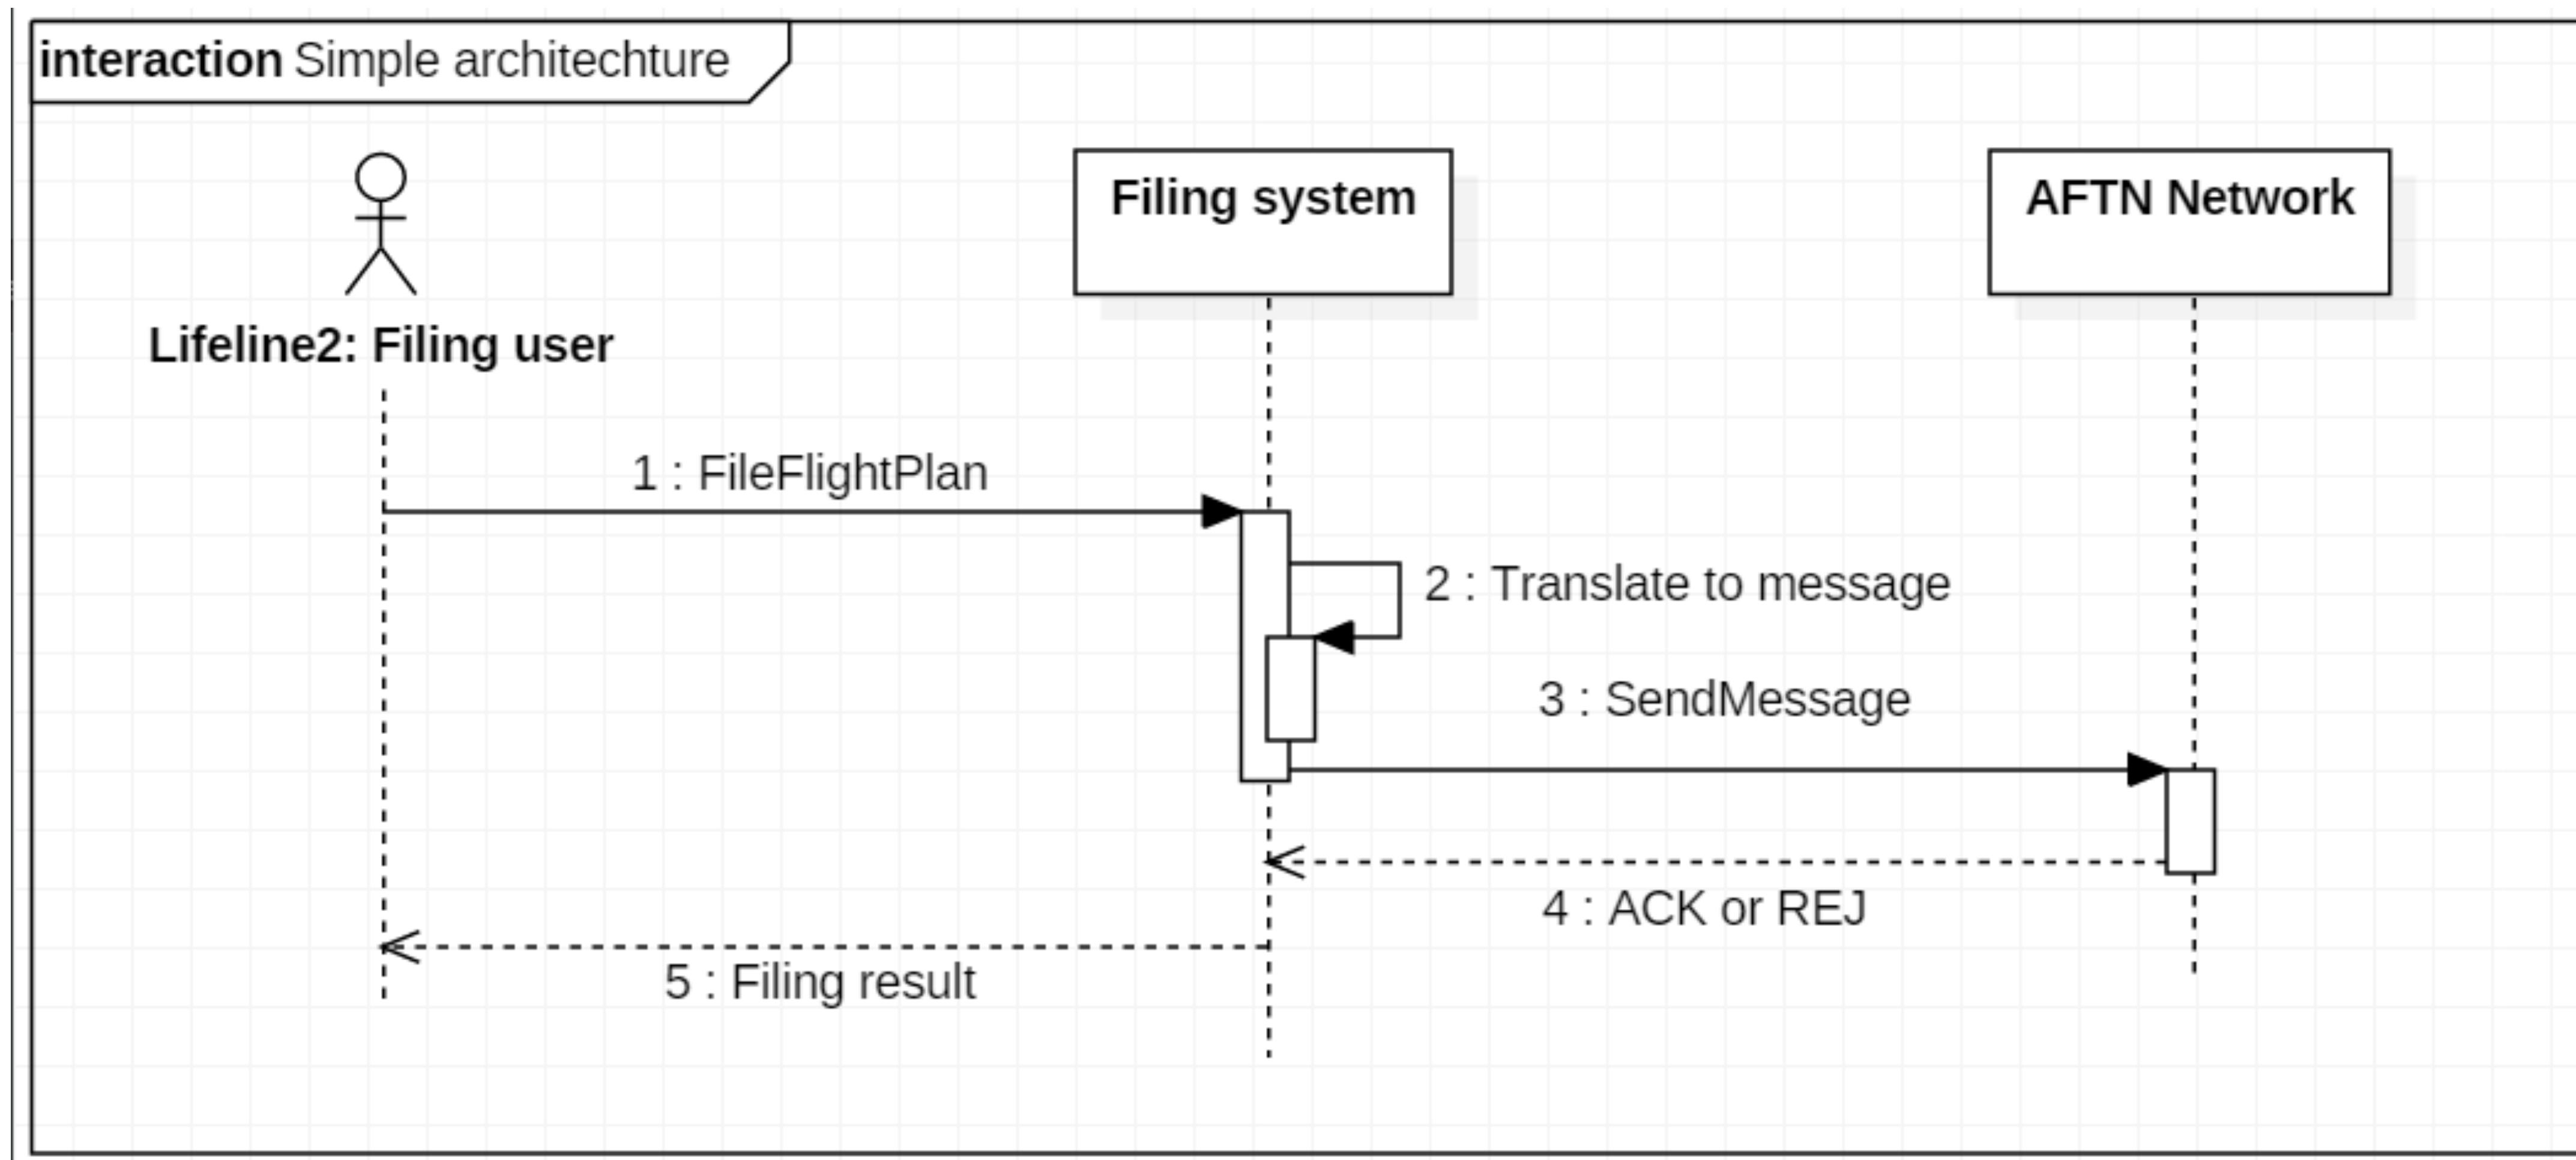
- The user must be able to submit a flight plan (shown on right) to the system
- The user must be able to update an existing flight plan
- The user must be able to cancel a submitted flight plan
- The user must be able to view an existing flight plan
- The user must be able to view all messages related to a flight plan

PRIORITY <<=FF		ADDRESSEE(S) KZNYZQZX EUCHZMFP EUCBZMFP CZQMZQZX	
FILING TIME		ORIGINATOR	
SPECIFIC IDENTIFICATION OF ADDRESSEE(S) AND (OR) ORIGINATOR			
3 MESSAGE TYPE <<= (FPL		7 AIRCRAFT IDENTIFICATION - O,Y,A,C,C,	
9 NUMBER - 1		TYPE OF AIRCRAFT G,L,F,5	
13 DEPARTURE AERODROME - KTEB		WAKE TURBULENCE CAT / M	
15 CRUISING SPEED - N0490		TIME 1,4,0,0	
LEVEL F390		ROUTE DCT SANTT J42 BOS DCT BAXIE DCT SCOTS N315A ELSIR/M085F390 DCT 50N050W DCT 52N040W DCT 53N030W DCT 54N020W DCT DOGAL/N0486F390 DCT BEKET DCT KORAK DCT DONEB DCT WST DCT OP DCT LIFFEY UL975 WAL LOREL4F	
TOTAL EET 16 DESTINATION AERODROME - EGGS HR MIN 0,54,3 ALTN AERODROME ⇒ EGCC 2ND ALTN AERODROME ⇒ 			
18 OTHER INFORMATION - PBN/B2 DOF/170809 EET/KZBW0002 CZQM0042 CZQX0123 EGGX0322 EISN0424 EGTT0505			
SUPPLEMENTARY INFORMATION (NOT TO BE TRANSMITTED IN FPL MESSAGES)			
19 ENDURANCE - E / 0,64,8		PERSONS ON BOARD ⇒ P / T,B,N	
SURVIVAL EQUIPMENT ⇒ X / X		EMERGENCY RADIO ⇒ R / UHF VHF ELT	
POLAR ⇒ X / X		JACKETS ⇒ X / X	
DESERT ⇒ X		LIGHT ⇒ X	
MARITIME ⇒ X		FLUORES ⇒ X	
JUNGLE ⇒ X		UHF ⇒ X	
DINGHIES ⇒ X / 		VHF ⇒ X	
NUMBER ⇒ X / 		ELT ⇒ X	
CAPACITY ⇒ X / 			
COVER ⇒ X / 			
COLOUR ⇒ X / 			
AIRCRAFT COLOUR AND MARKINGS A / 			
REMARKS ⇒ X / 			
PILOT-IN-COMMAND C / 			
FILED BY		ADDITIONAL REQUIREMENTS	

Non-functional requirements

- #NF1: The system must support an uptime of at least 99.9 %
- #NF2: The system must ensure message integrity – that is, outgoing and incoming messages may never be lost
- #NF3: The system must continue to be able to accept flight plans even when AFTN network connectivity is unavailable
- #NF4: There can only be one single connection to AFTN, and that connection must be routed through a single IP address on a server sitting on a protected network
- #NF5: Only one message can be sent to AFTN at a time

Design #1



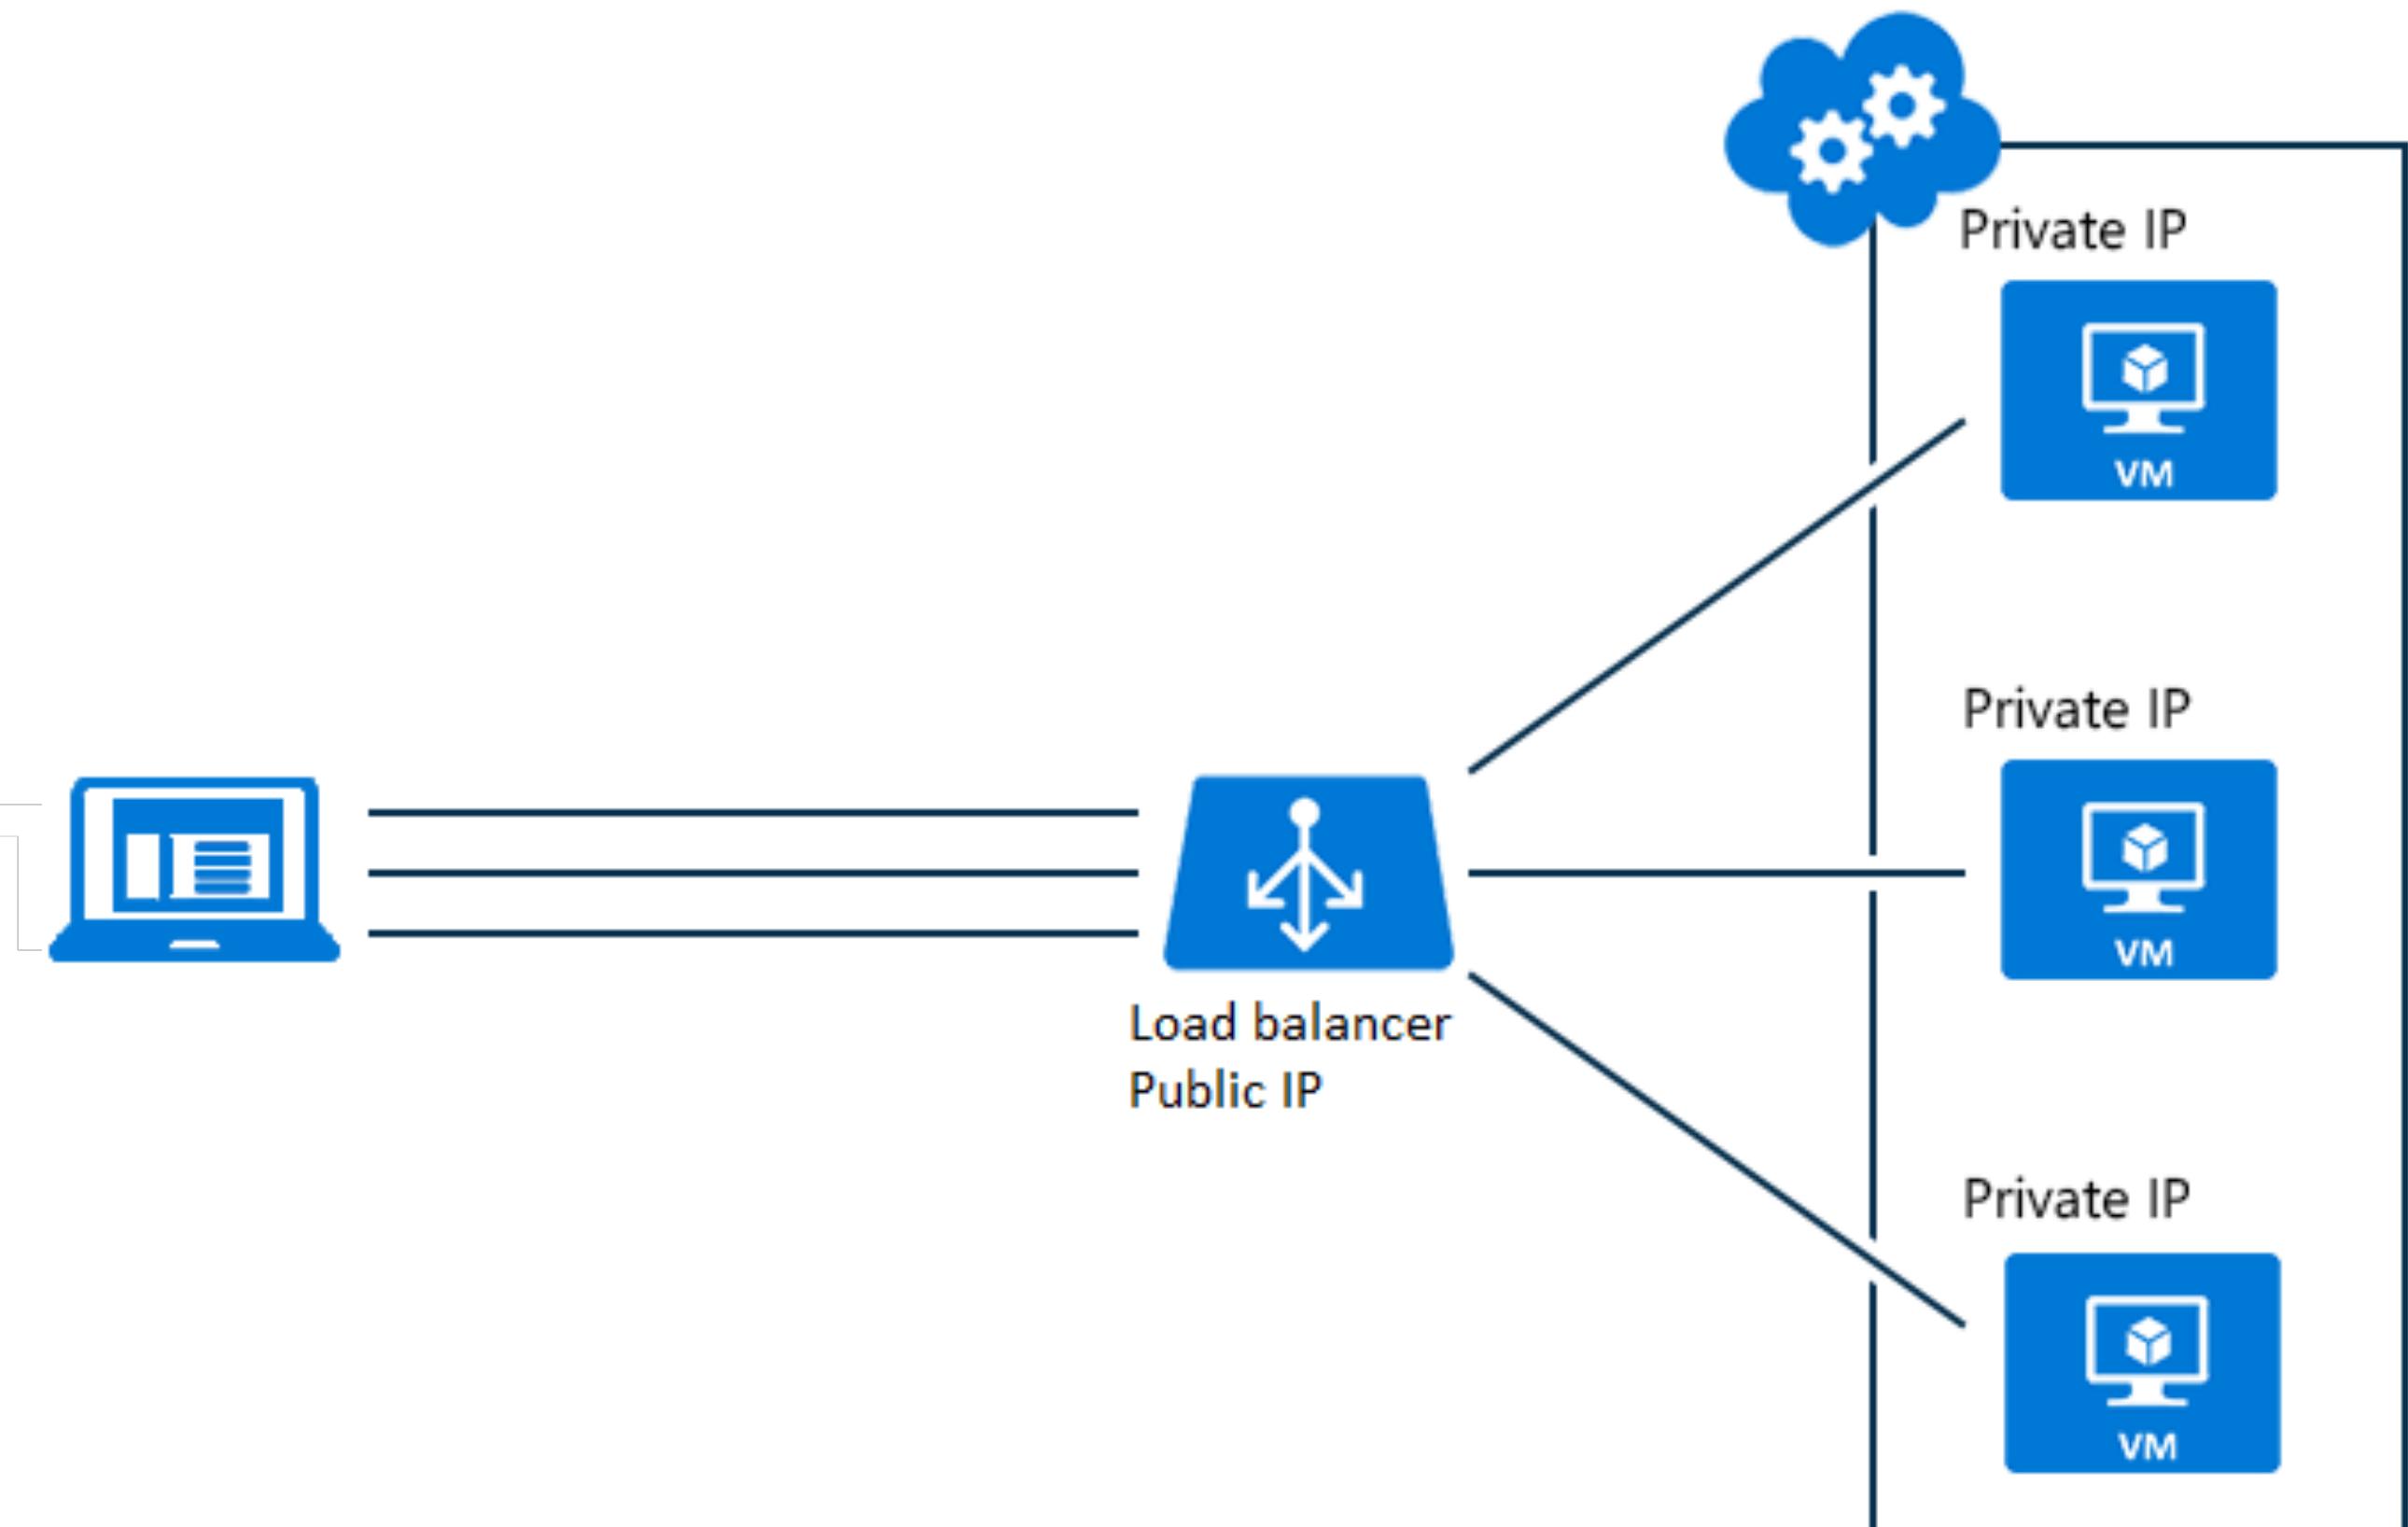
AVIATION CLOUD

Consequences of this design

- #NF5 dictates we can only send one message at a time to AFTN, so with this design, we can only allow one user to file a flight plan at a time
 - Result: Low throughput
- #NF4 dictates we can only make one connection from one single server with a dedicated IP address sitting on a protected network
 - If the network is protected, how can the user connect?
- #NF1 requires 99.9 % uptime. If user connect to a single server, how can we ensure that?
 - Software updates, OS updates, physical hardware maintenance etc...

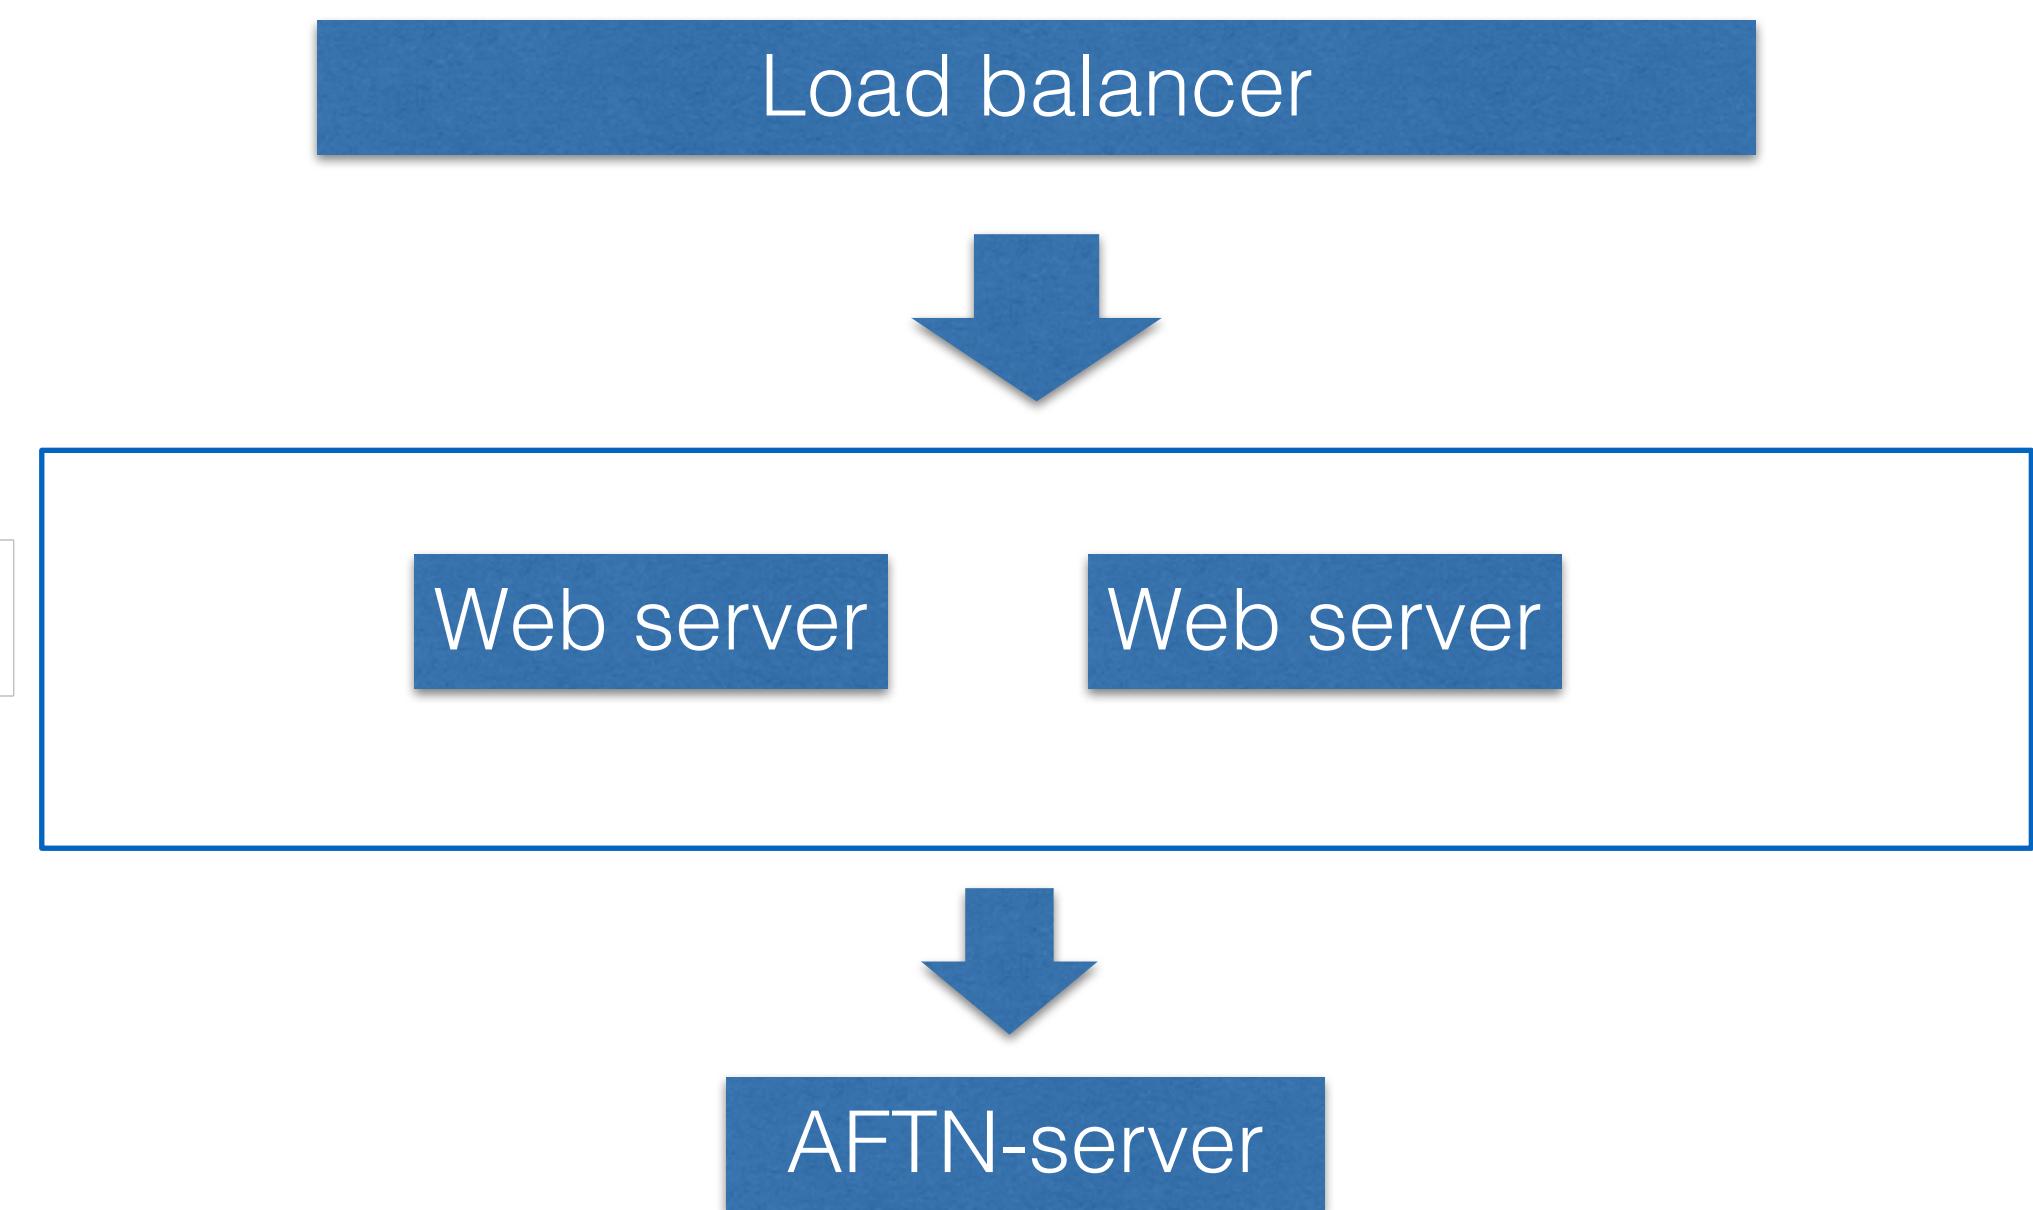
Solving availability – Load balancers

- A load balancer distributes load between servers
- They also help ensure uptime.
One server can be disconnected from the load balancer during maintenance, connected once done, then the other is updated – aka "Rolling updates"



Design #2 – Multi-tier

- **Users (systems) connect to the load balancer**
- **Load balancer forward the request to a web server**
- **The web server sends the AFTN message by invoking an API on AFTN-server**

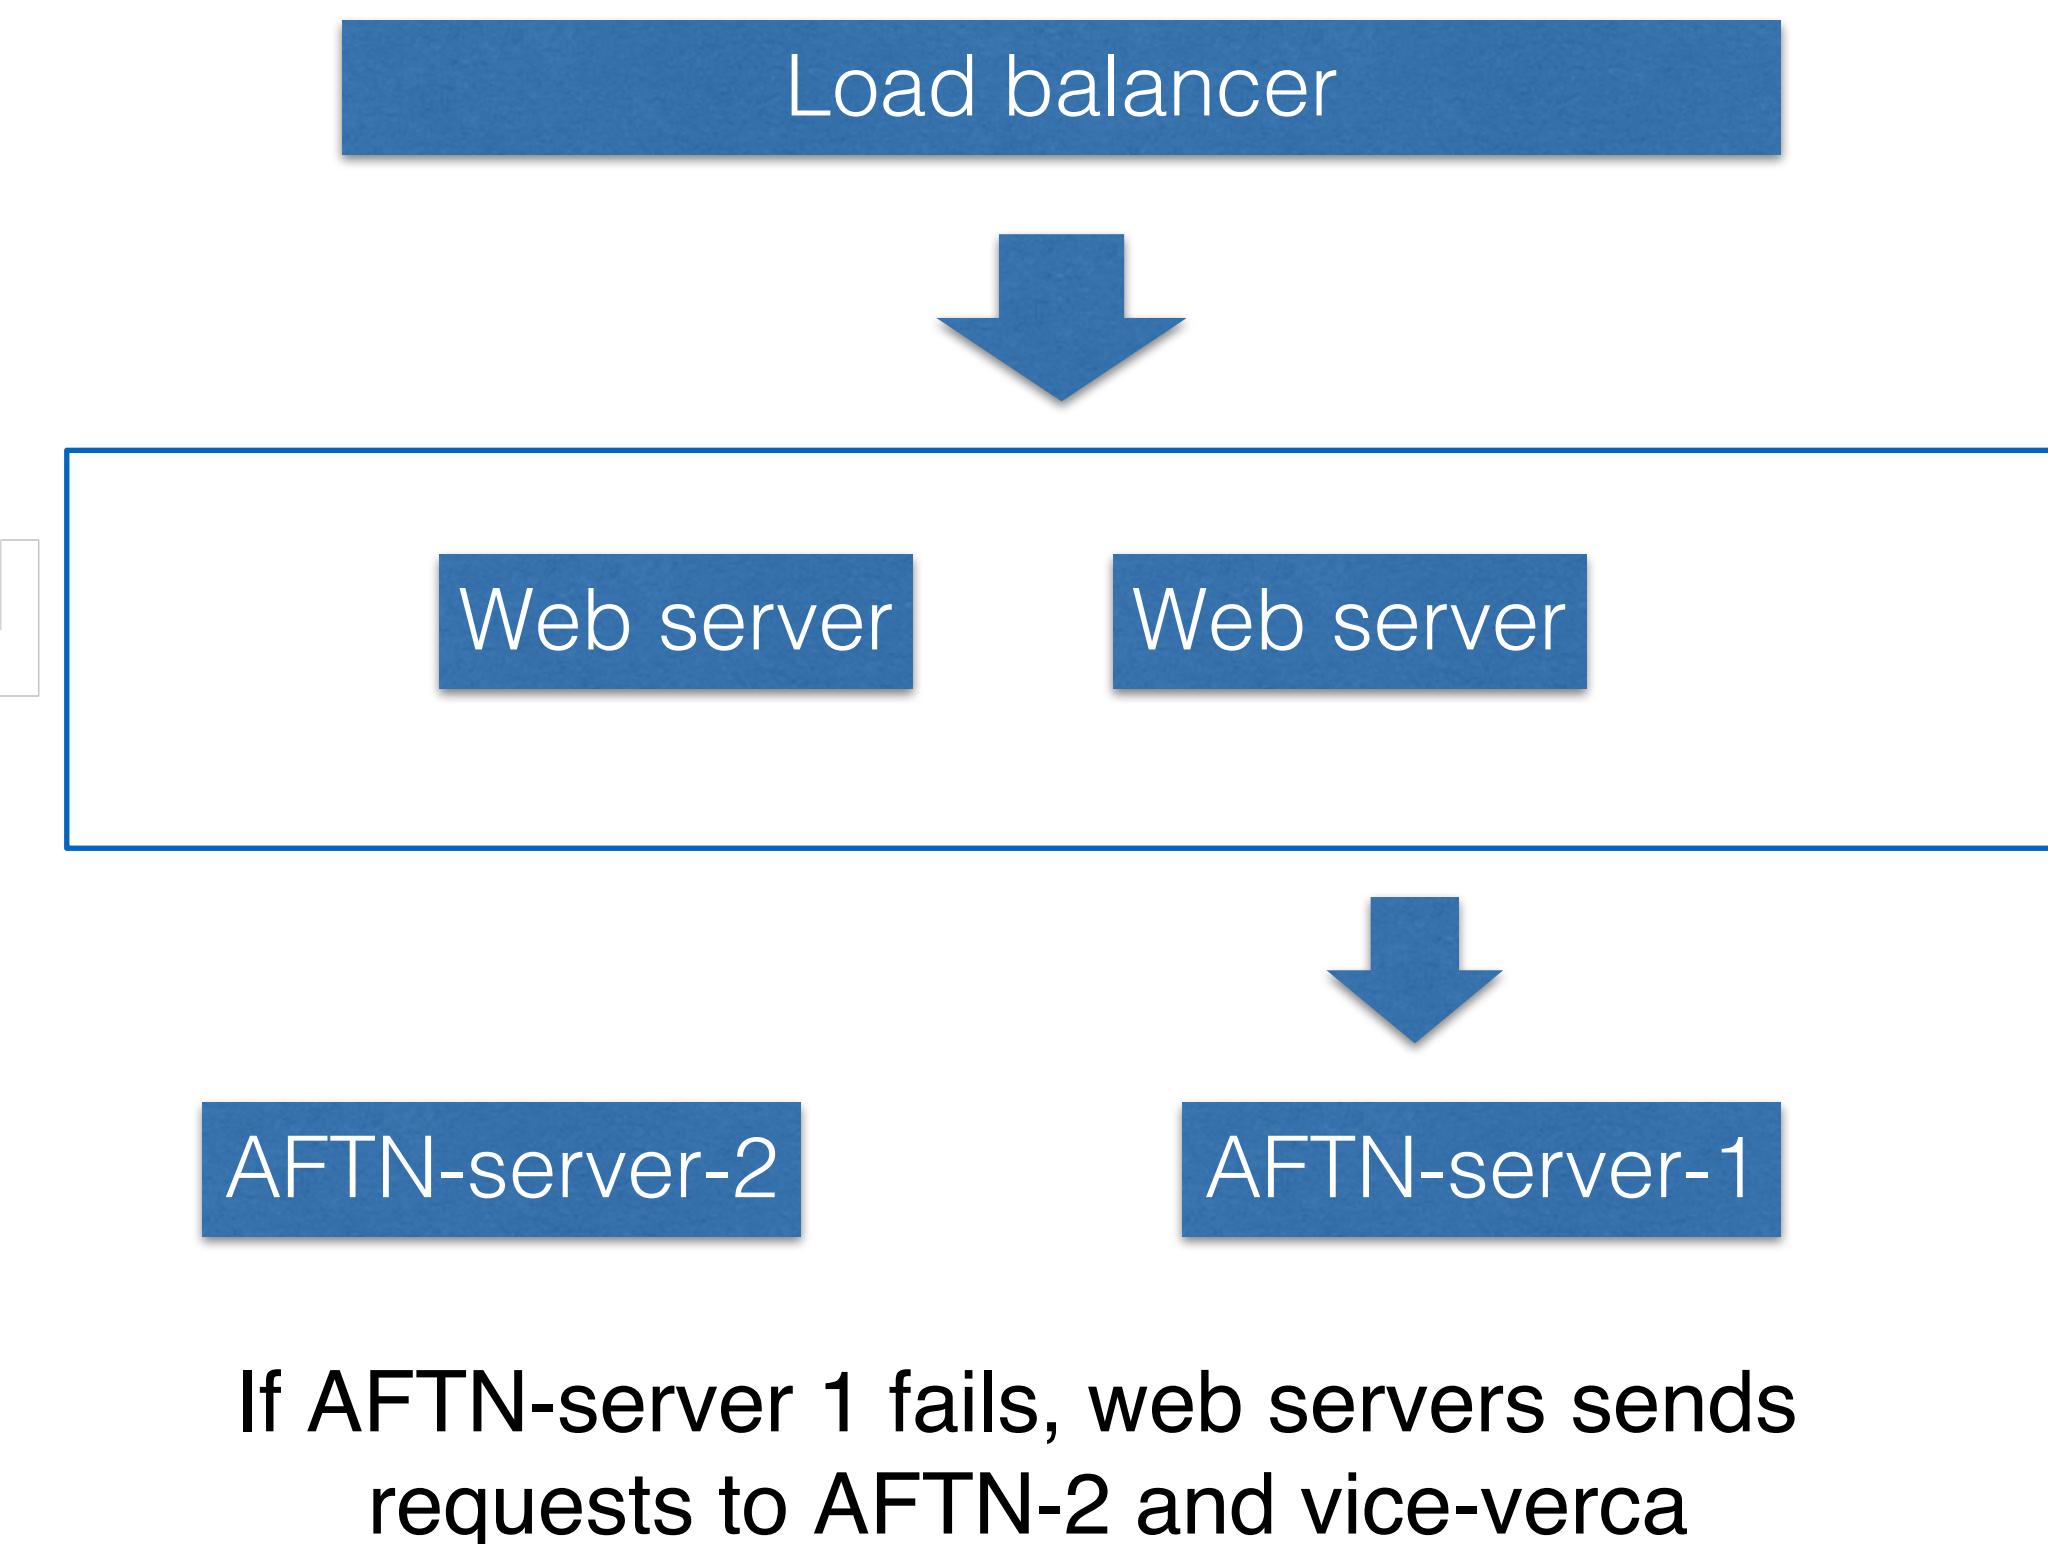


Design #2 - Issues

- Still limited throughput
- Difficult to implement the sent-one-message-at-a-time #NF5
- We solved the availability issue at the web-layer, but what if AFTN-server needs maintenance?
 - We can't do load balancing, since only one active AFTN connection is allowed

Design #3 – Multi-tier + hot-standby AFTN

- A hot standby is a server that is inactive, but is ready to take over at any moment when it discovers the other server(s) is down
- In practice:
 - Active server registers a "ping"
 - Non-active server looks for the ping
 - When non-active server sees missing ping, it makes itself active
 - Issue: Race conditions, but that's a story for later ;-)



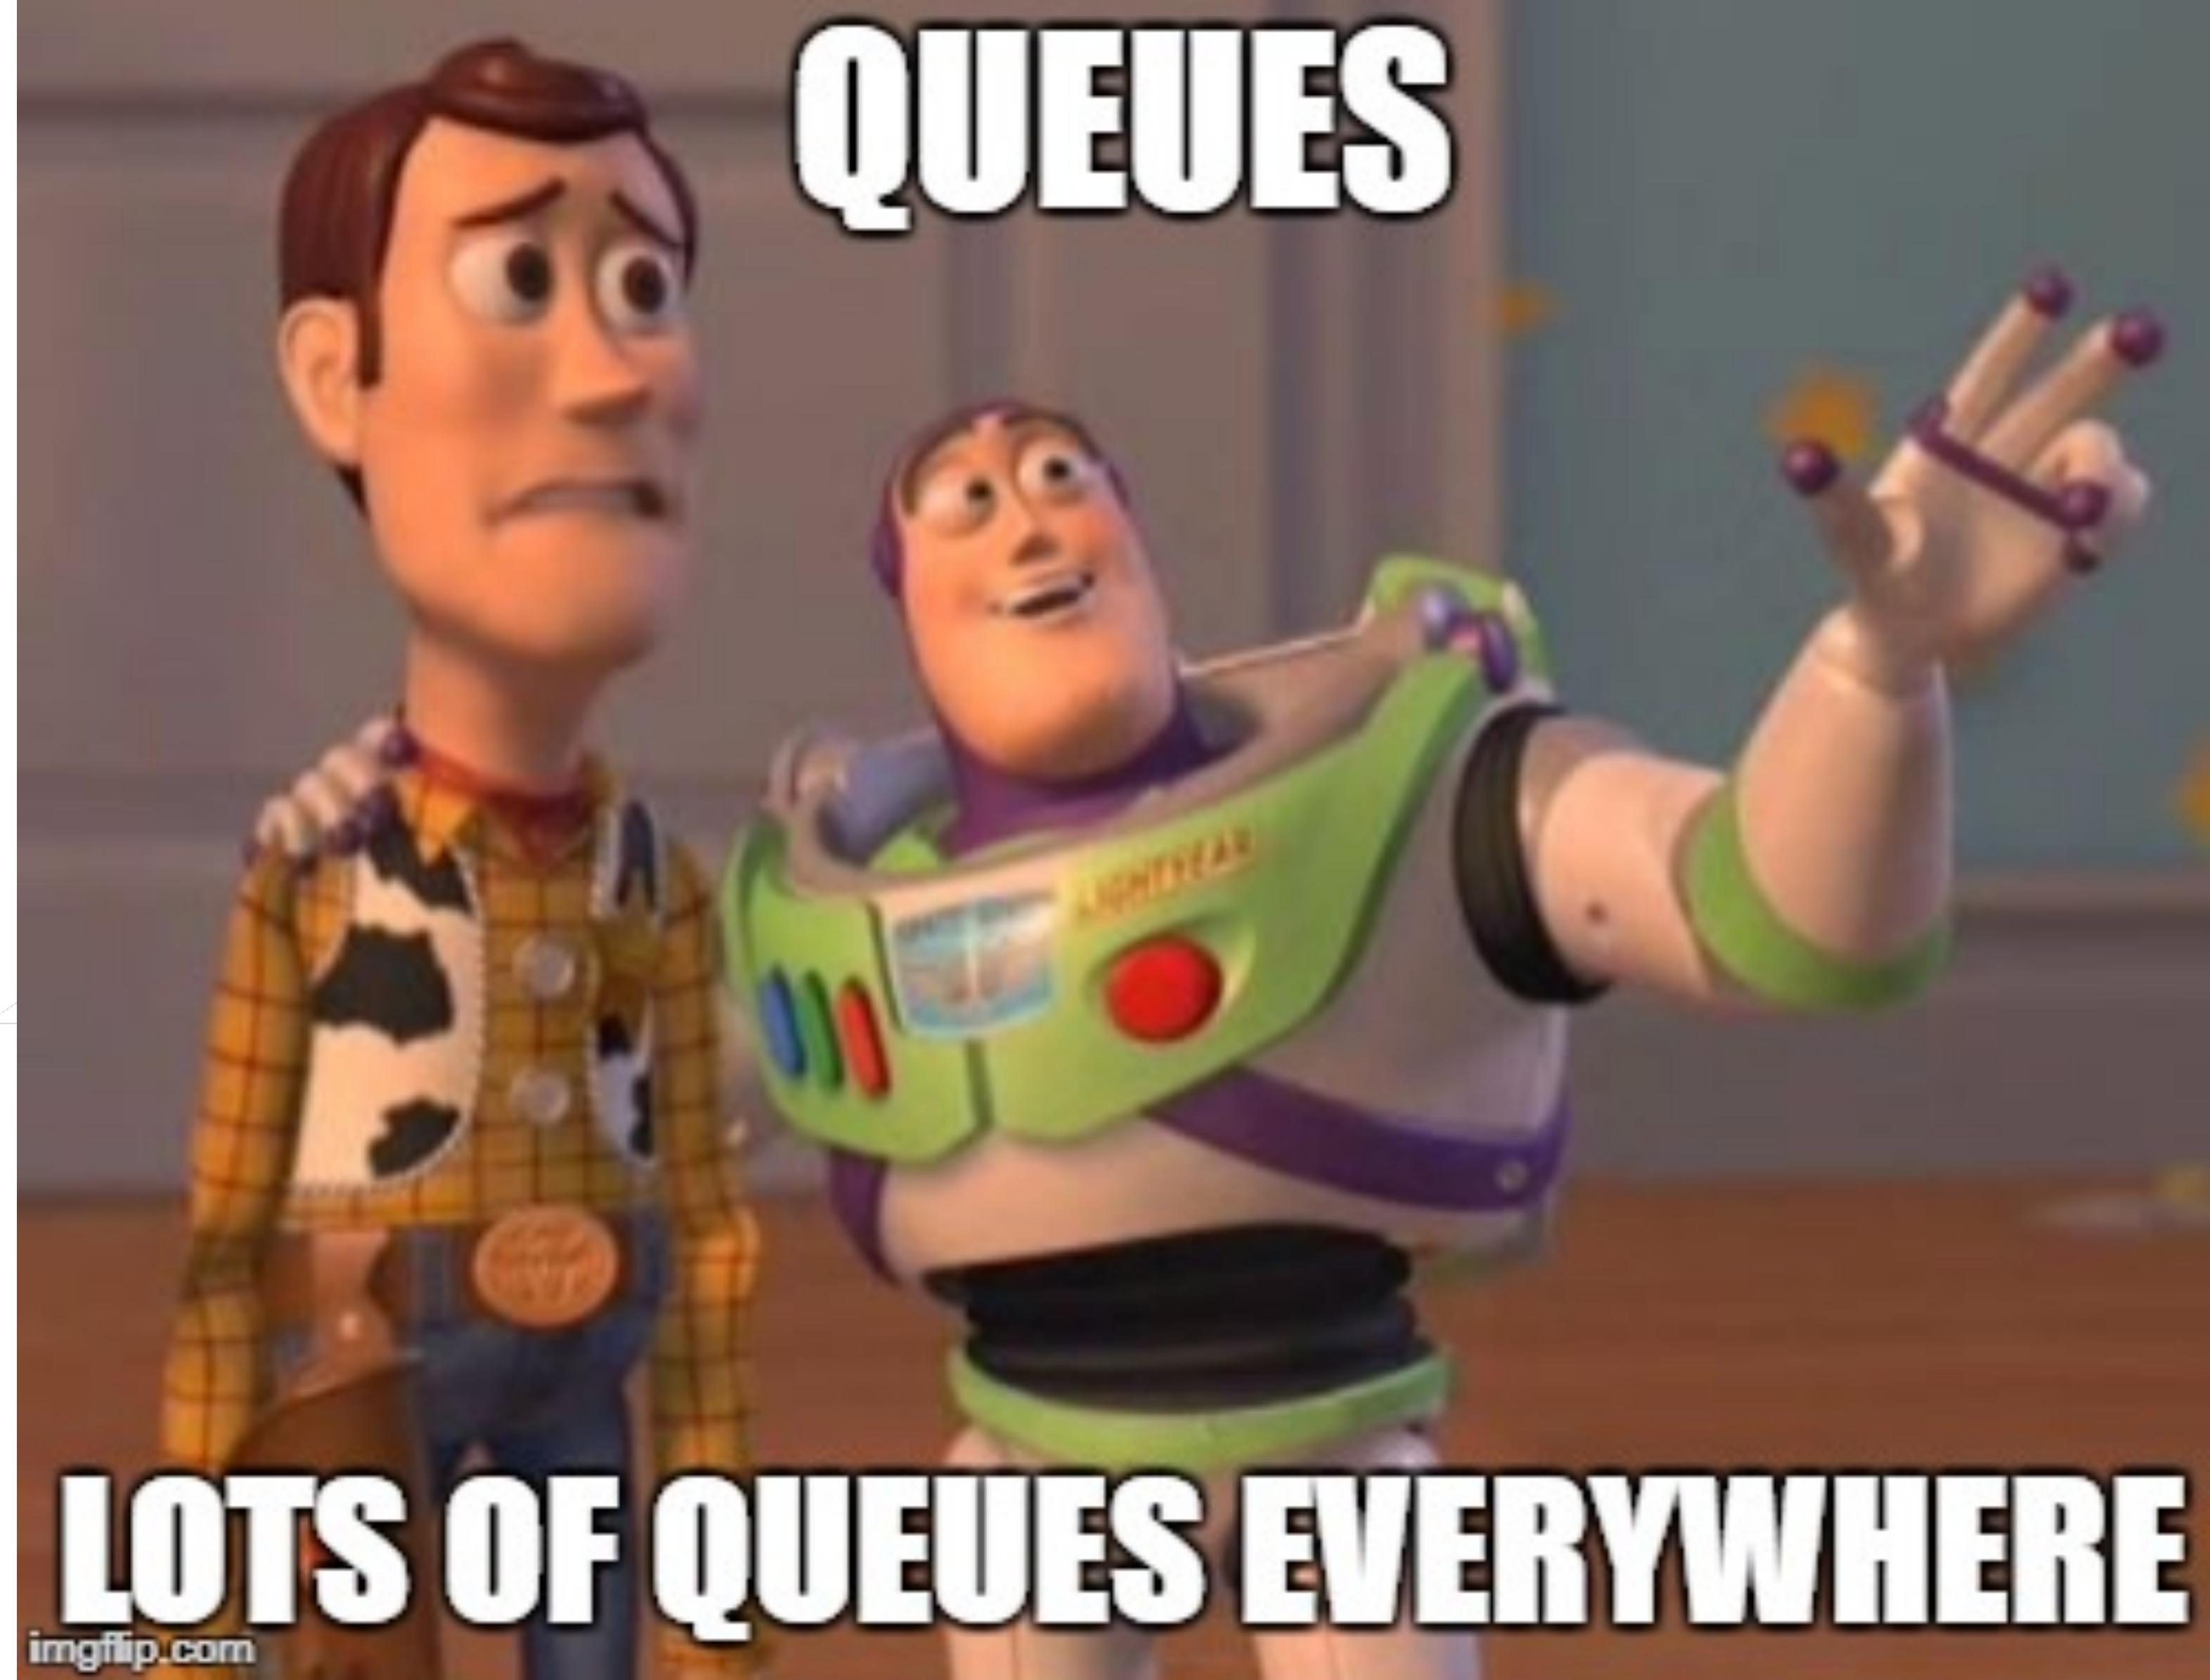
Issues – Design #3

- Things are getting better on the availability side. Using load-balancing and hot-standby we should be able to ensure that our own services keeps running
- We still have some major issues:
 - How can we continue to accept flight plans if AFTN is unavailable?
 - How do we ensure message integrity?

Queues

- Queue is a basic data structure that everyone should know
- They come as in-memory variant (i.e. `Queue<T>` in Java/C#)
- And as *distributed* and *persistent* queues



A scene from Toy Story featuring Woody (left) and Buzz Lightyear (right). Woody is a cowboy doll with a brown vest and a drawstring belt. Buzz is a space ranger with a green and purple suit and a red button on his chest. They are standing in a room with wooden floors and walls. The word "QUEUEUES" is overlaid in large white letters at the top, and "LOTS OF QUEUES EVERYWHERE" is overlaid at the bottom.

QUEUEUES

LOTS OF QUEUES EVERYWHERE

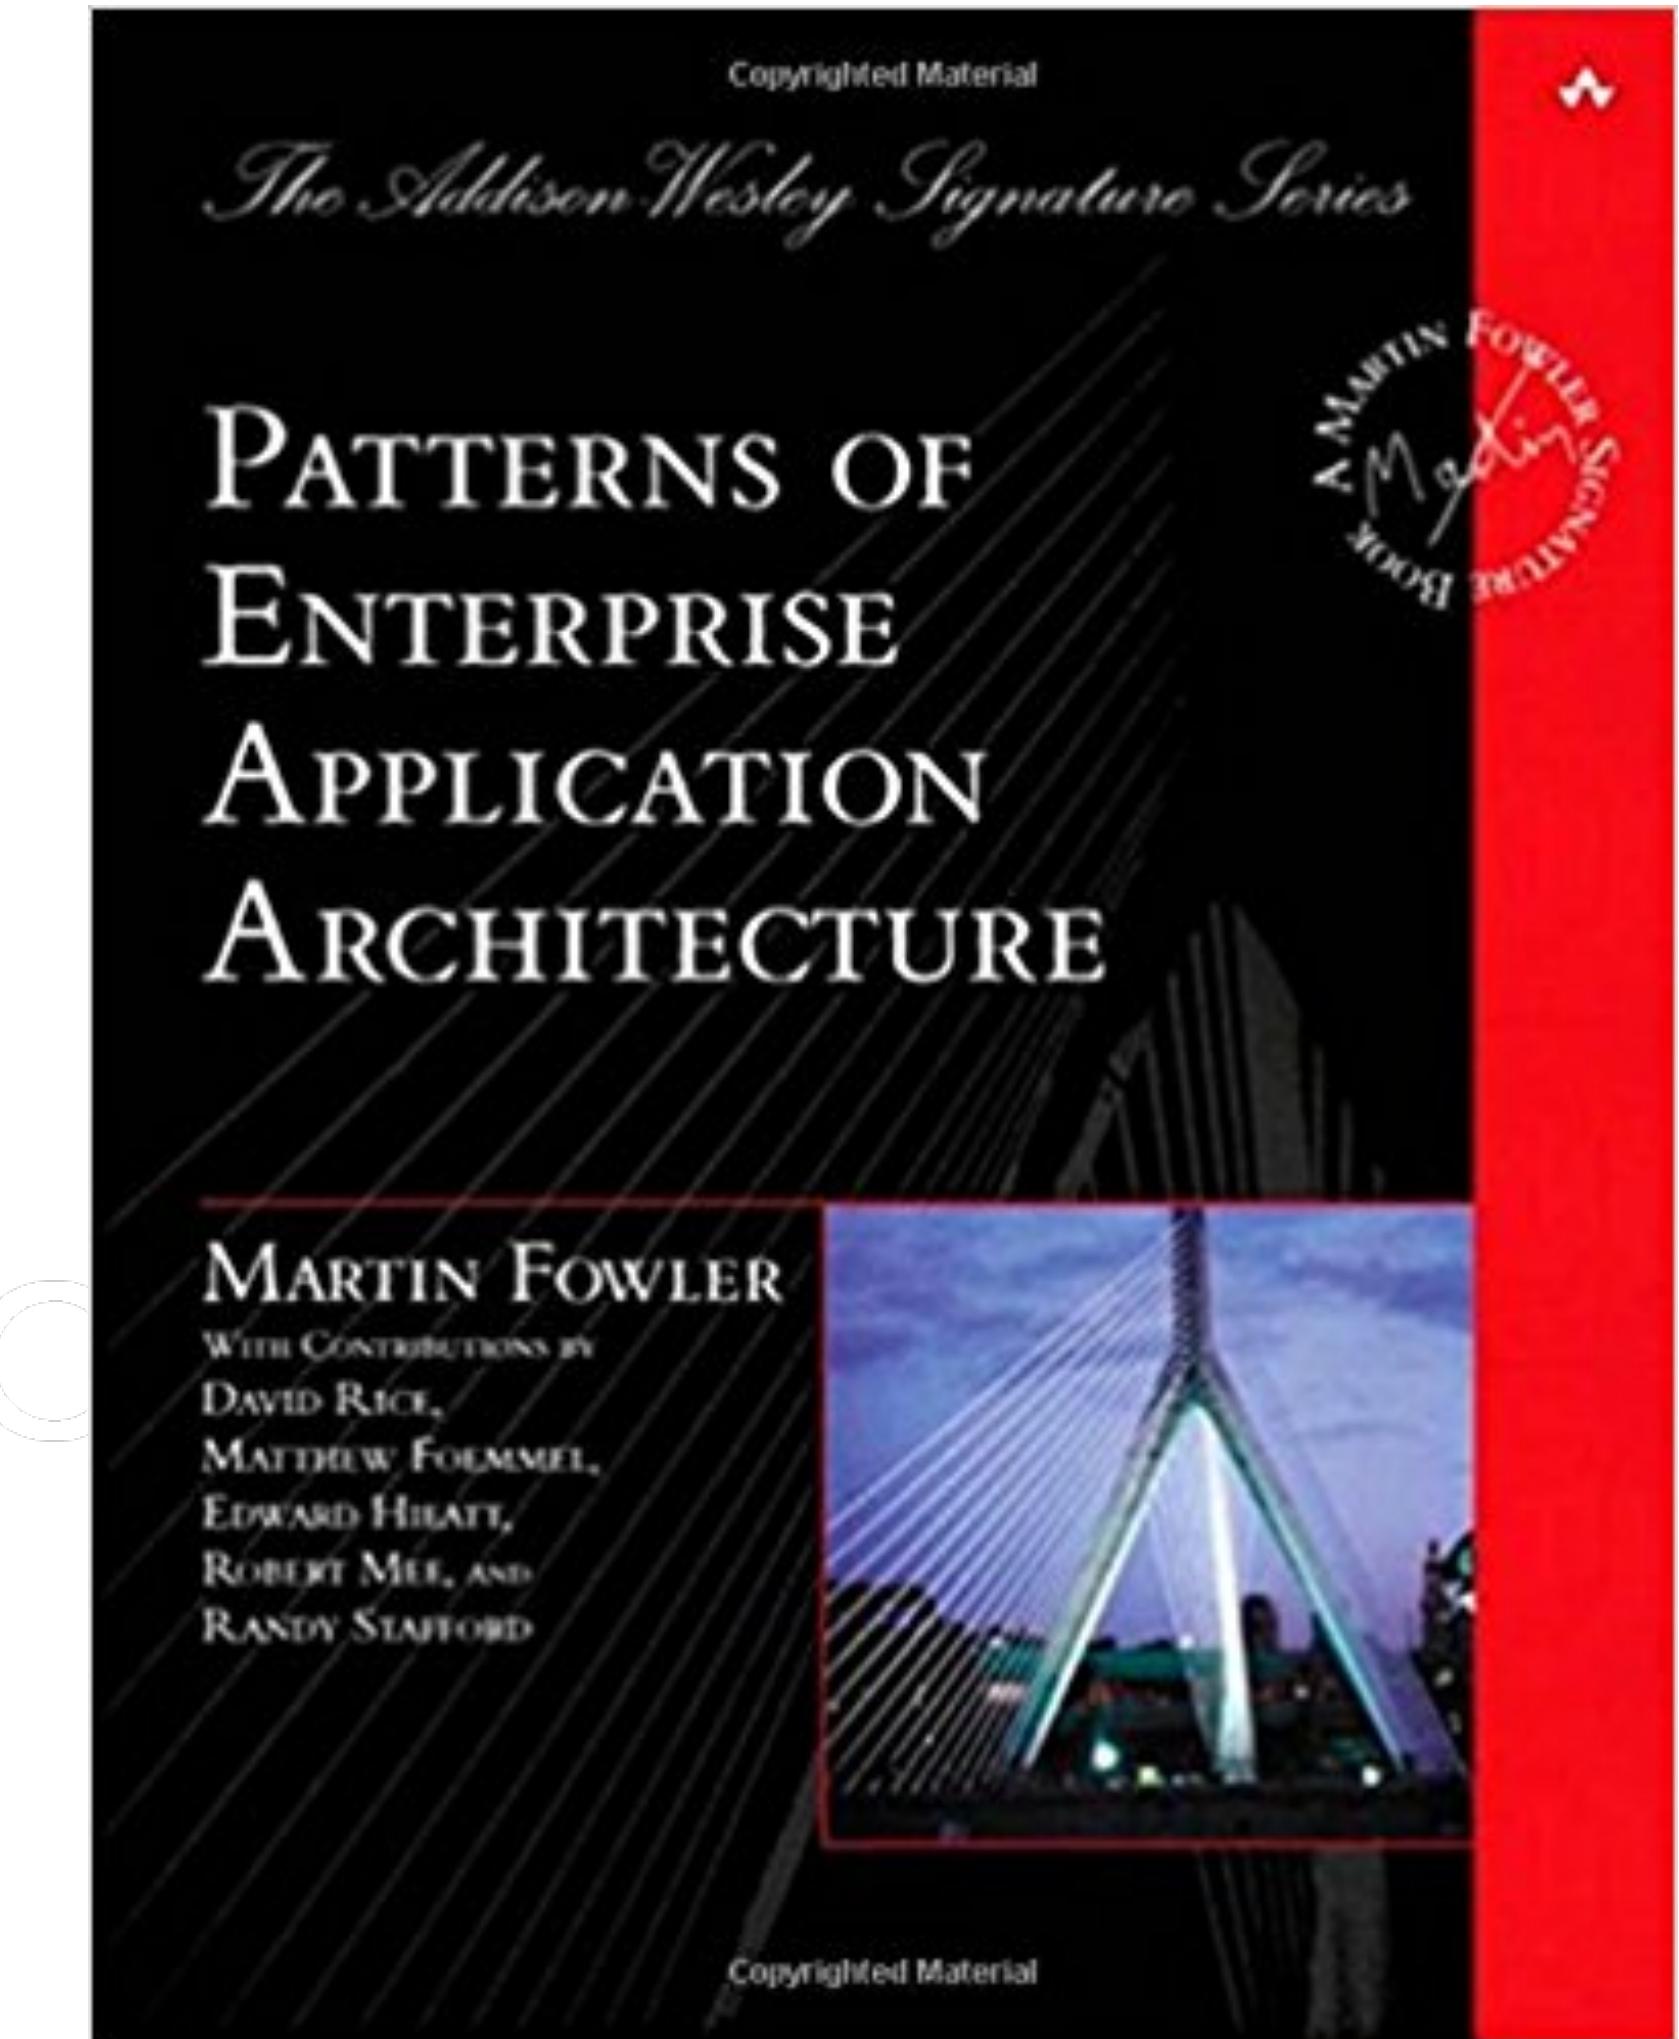
imgflip.com



AVIATION CLOUD

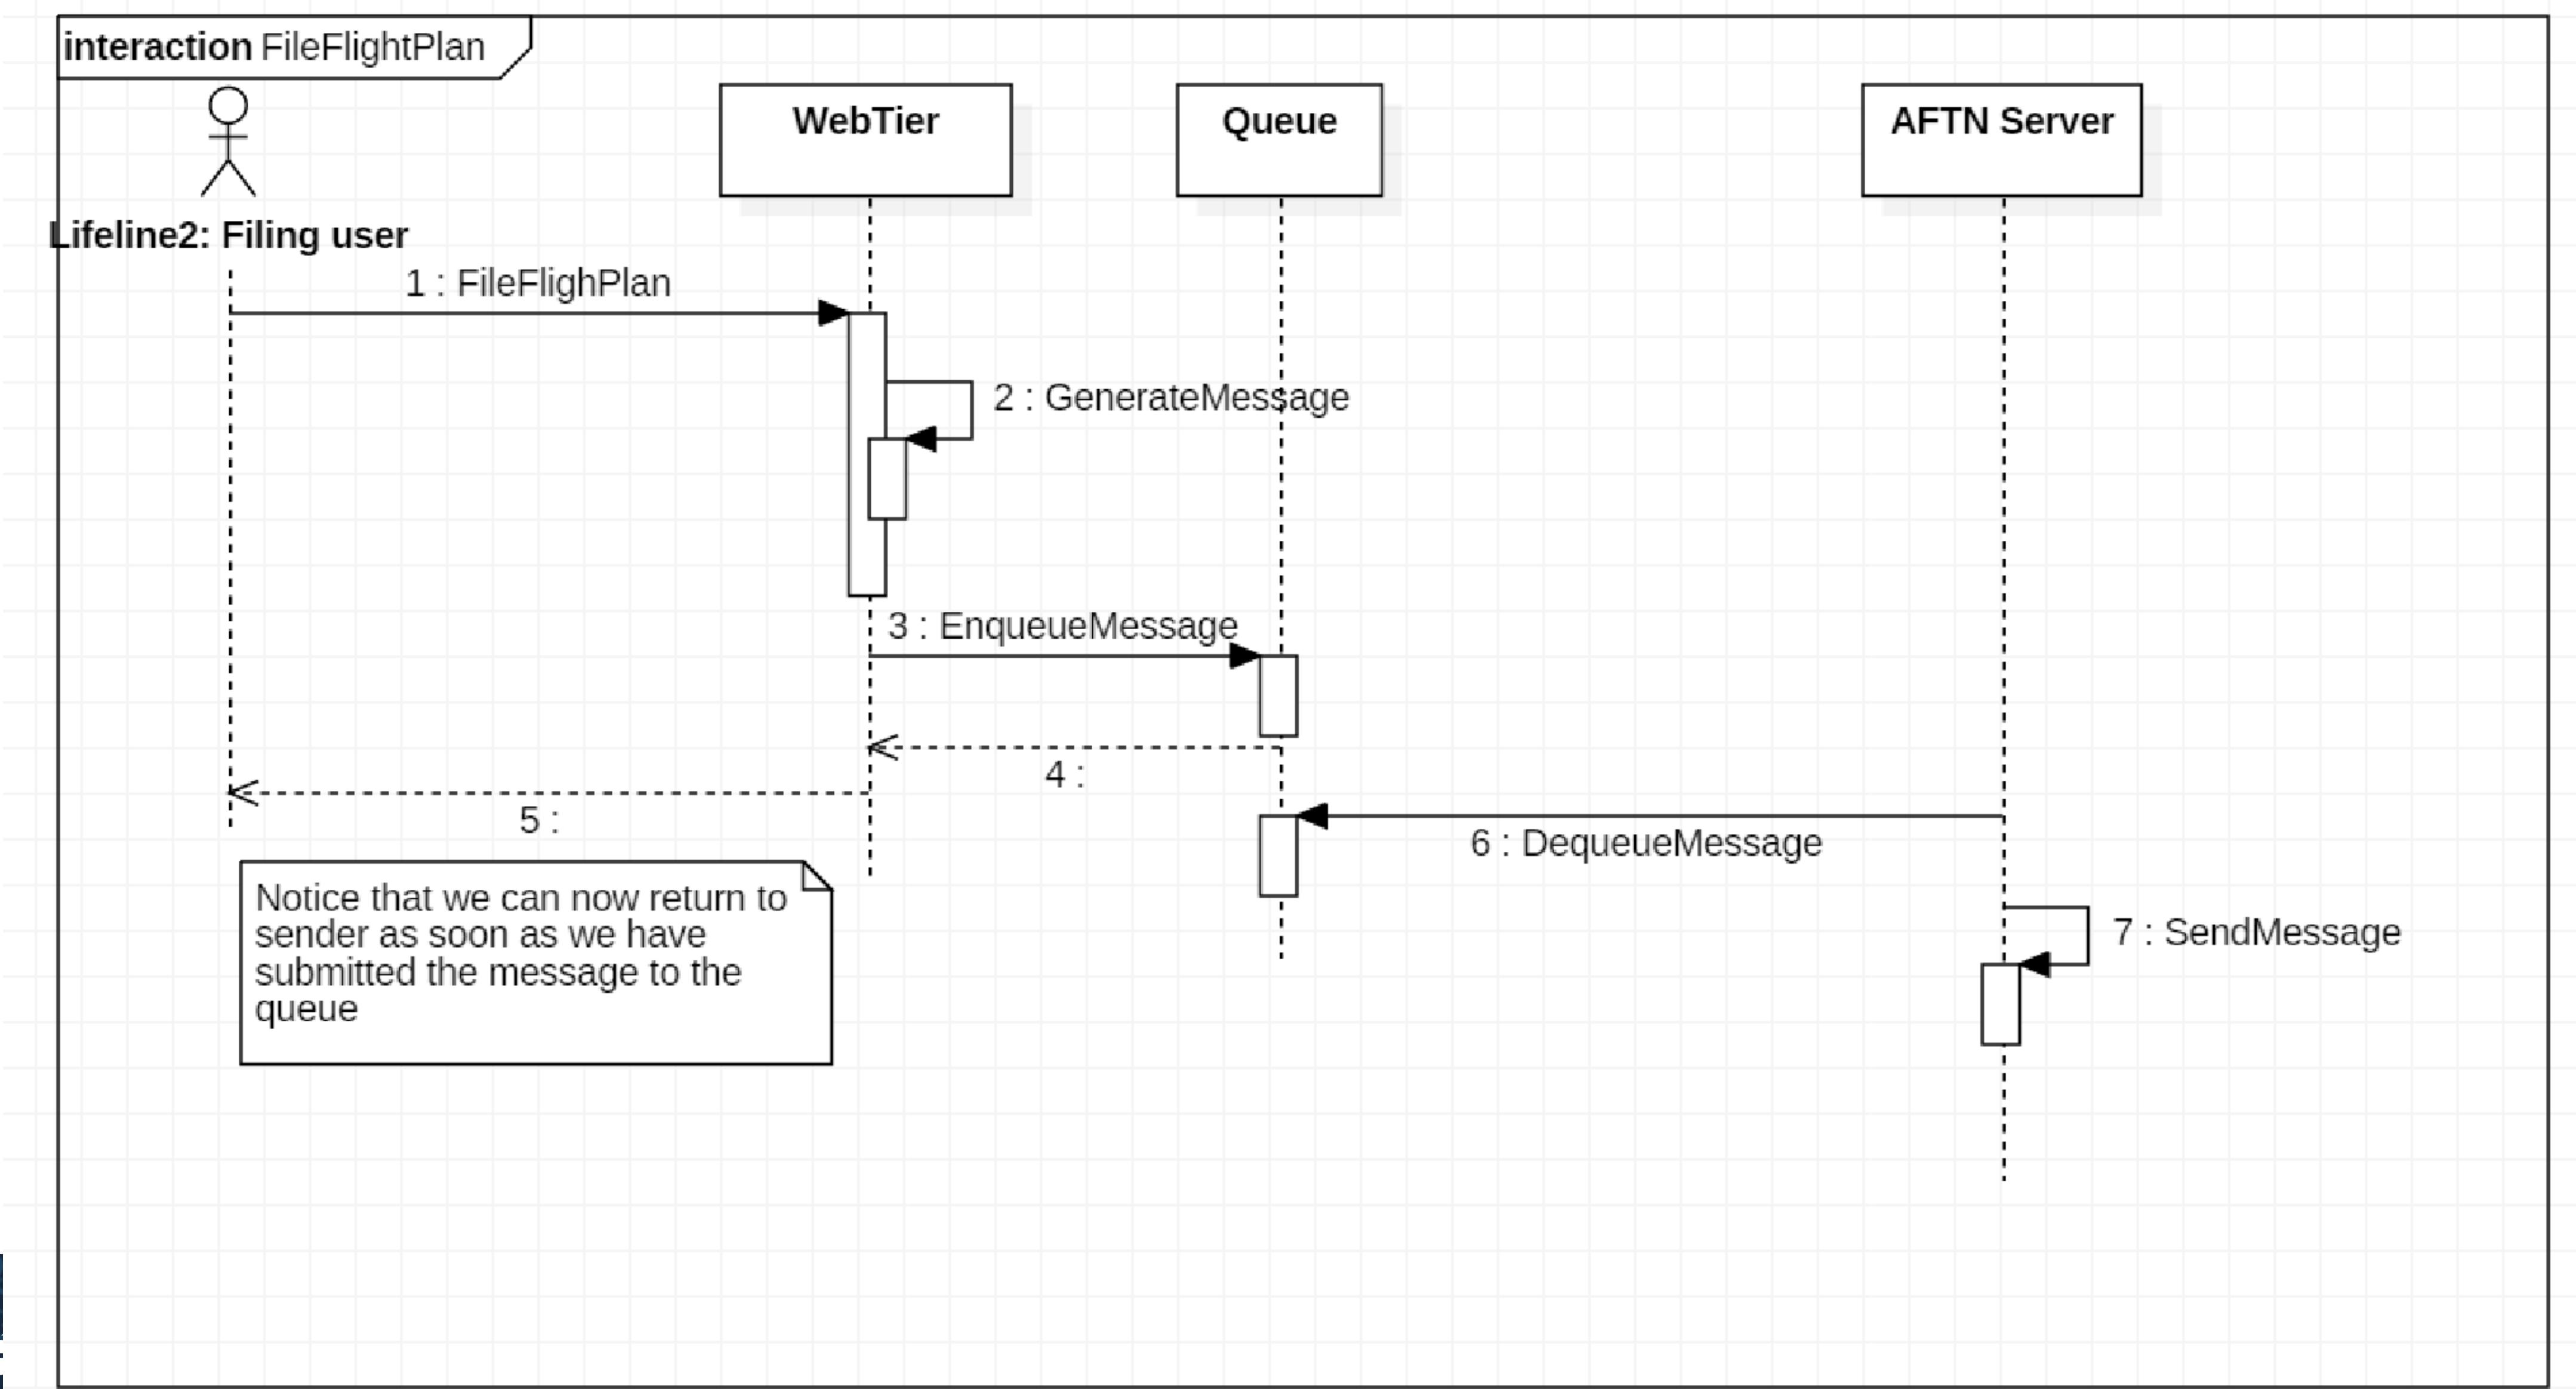
Distributed queues

- A distributed queue is a service running on a server with its only purpose of hosting a queue
- Other servers can enqueue or dequeue messages
- Distributed queues are often persistent, meaning that even if the server hosting the queue service becomes unavailable, the state will be restored once restarted



Explains some of the patterns around using queues as an architectural building block

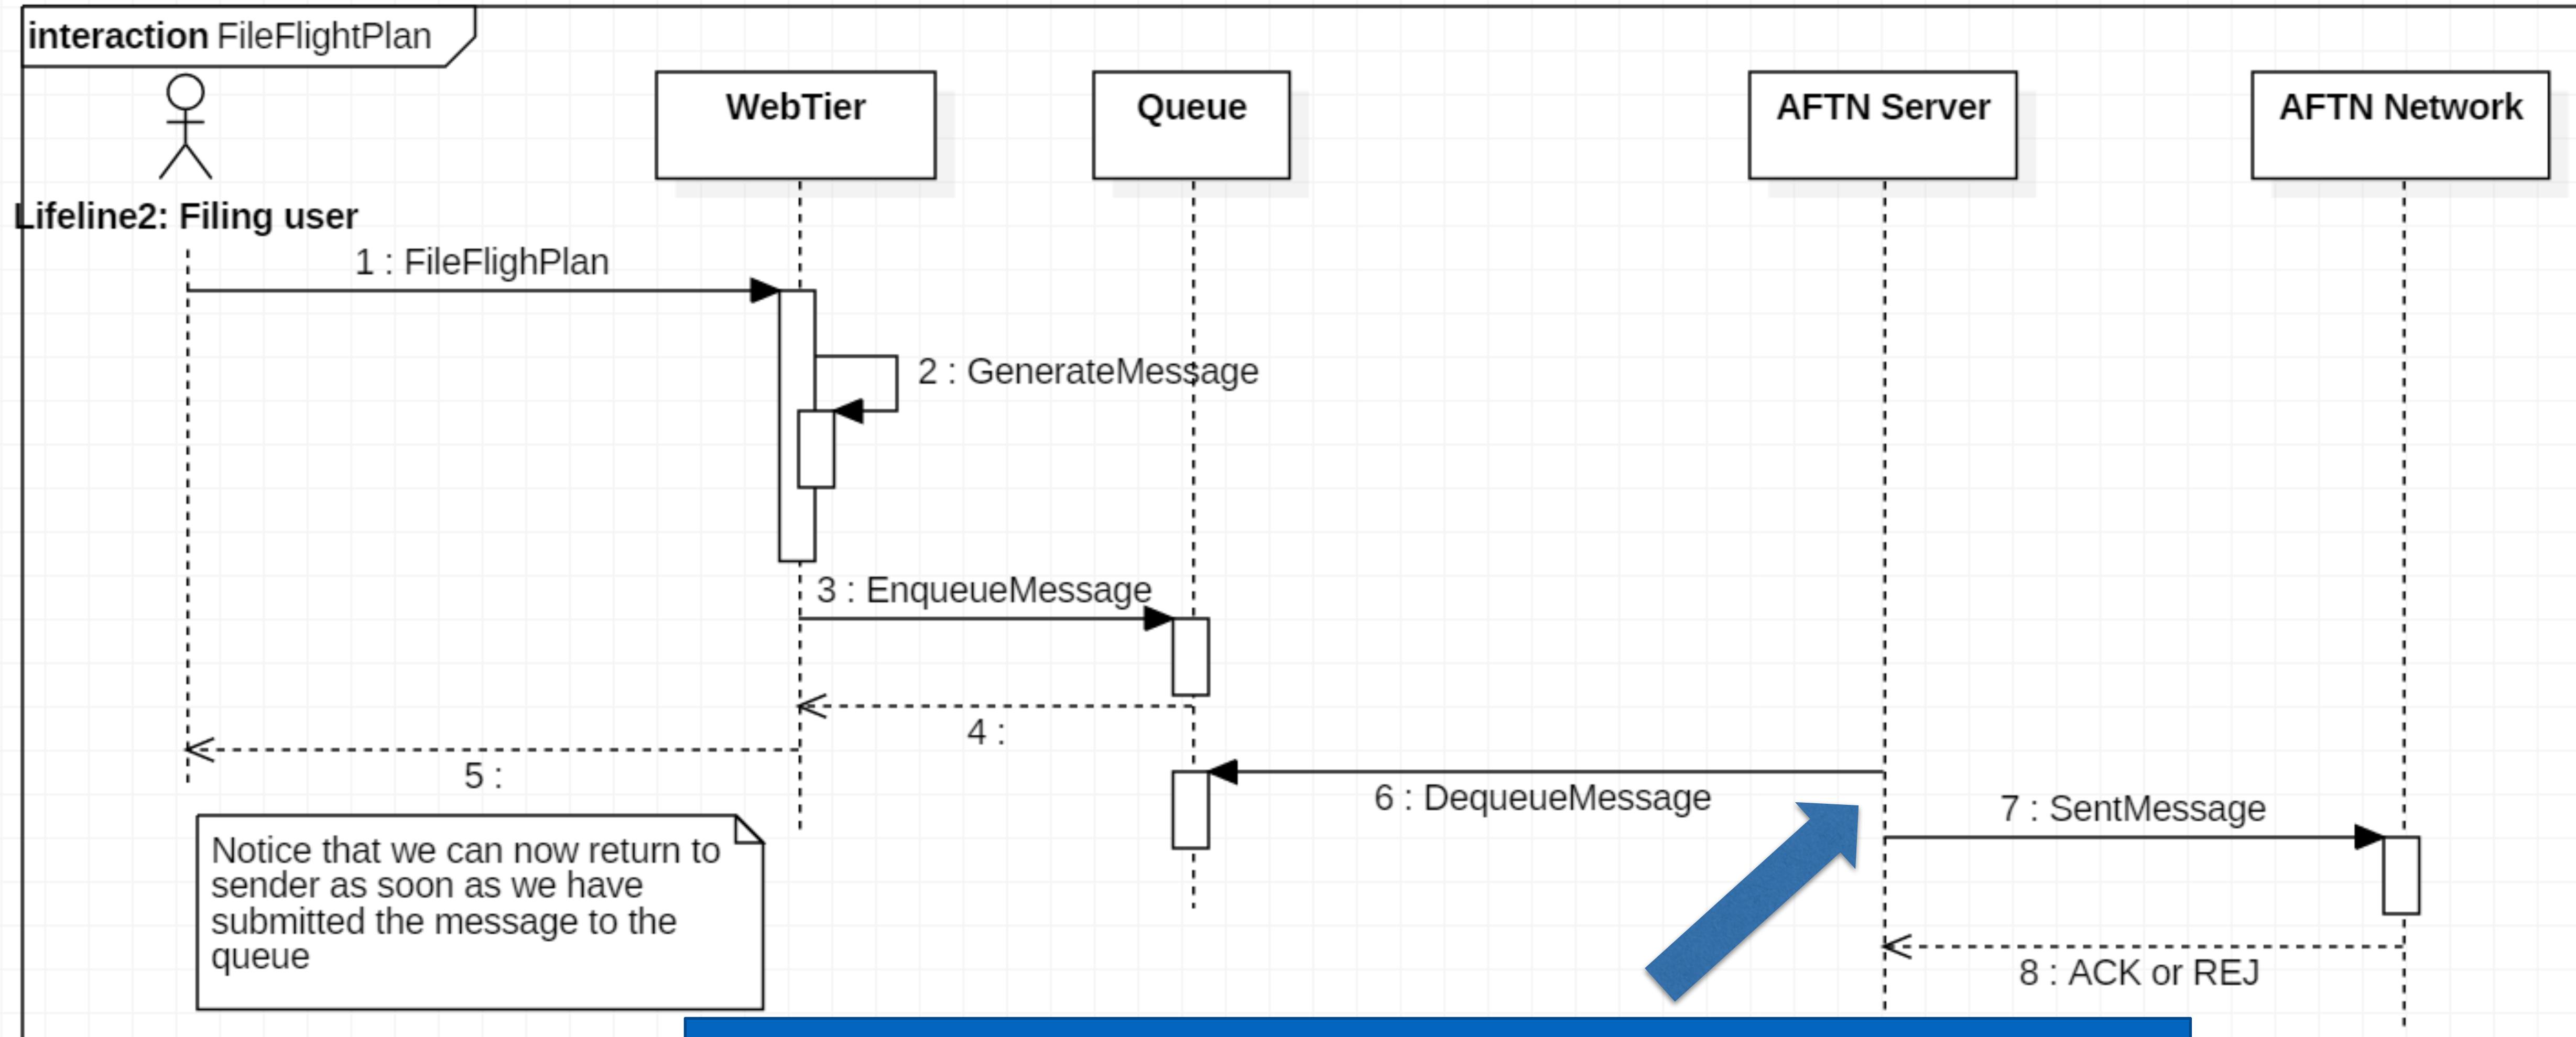
Design #4 – Queue based



Design #4

- We just made some big advances:
 - Messages can now be delivered to the system even when AFTN connectivity is down
 - Coordinating that only a single message can be sent is now easy because of the queue structure (first in, first out!)
 - If servers crashes, the messages are kept in the queue... Or are they?
 - The only external resource the AFTN server needs to be able to access is the queue service, meeting the security requirements

Issues with design #4



What happens if the server crashes after dequeuing the message, but before sending the message?

EVERYTIME YOU LOOSE AN AFTN MESSSAGE



A KITTEN DIES

imgflip.com

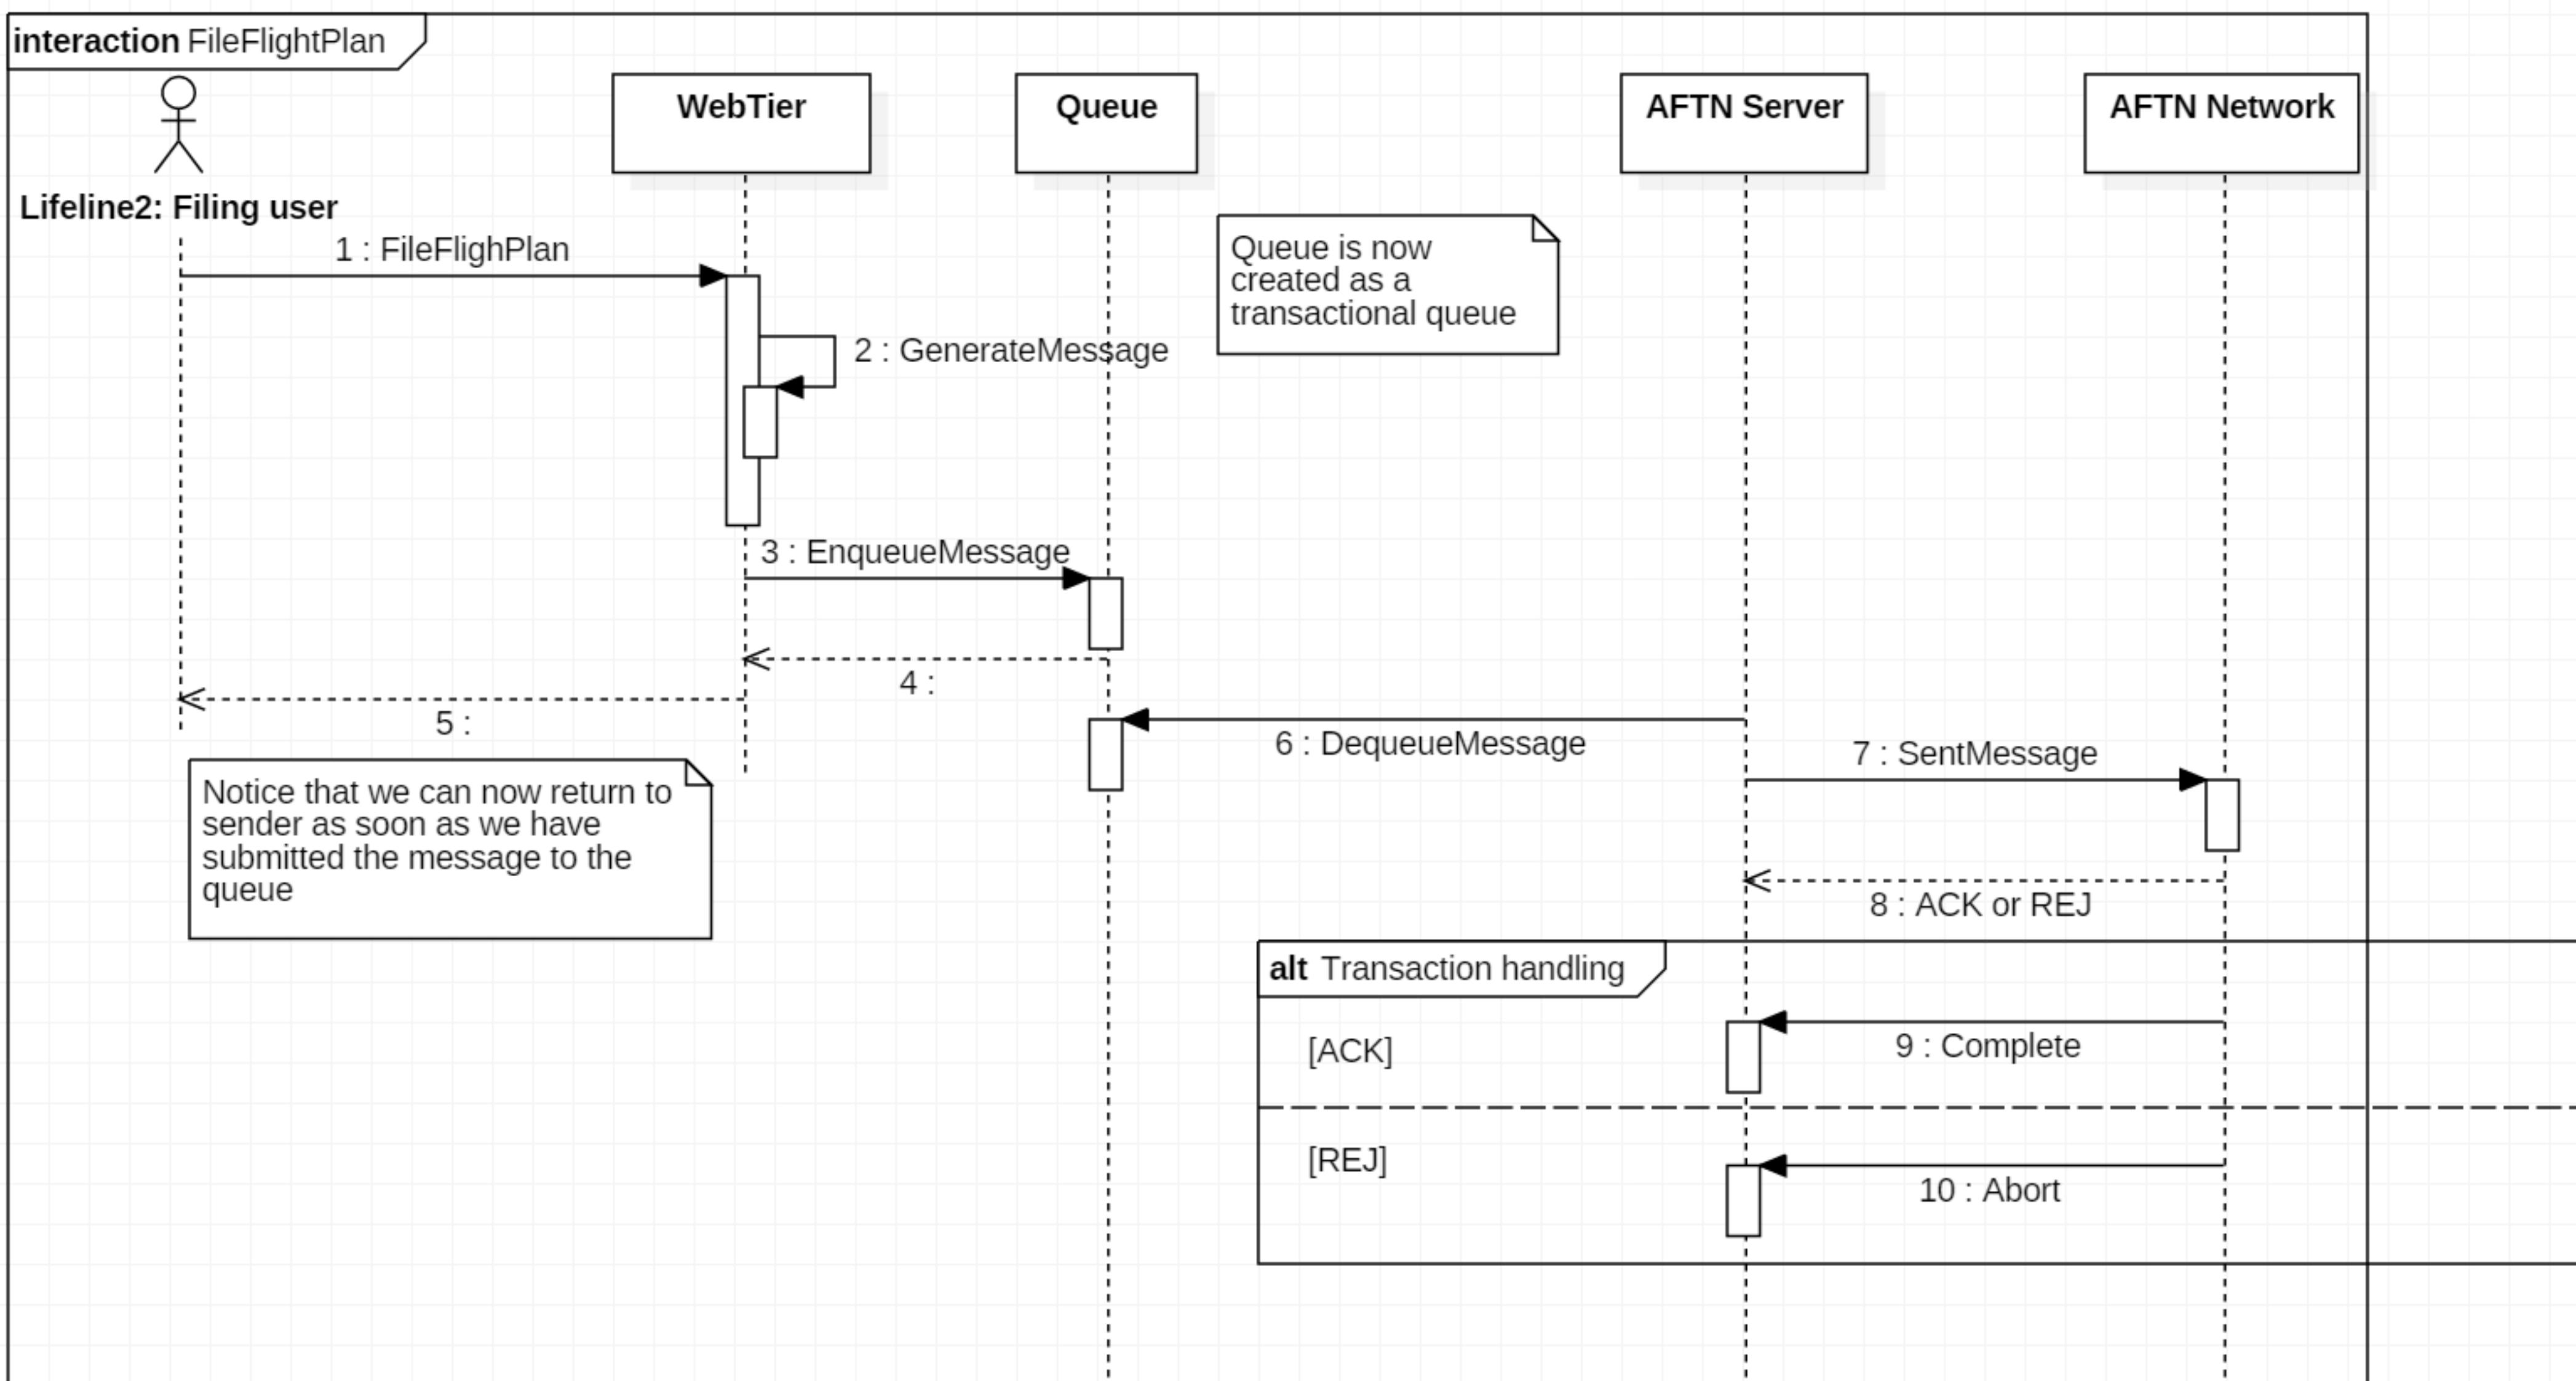


AVIATION CLOUD

Transactional systems

- A *transactional system* is a system where an operation can be rolled back if it does not succeed
- In practice
 - You begin an operation by performing an operation such as "Dequeue".
 - You then perform the action.
 - If it succeeds, you end the transaction by calling a "Complete" operation
 - If it fails, you abort the transaction by calling an "Abort" operation
 - If the server does not receive either a Complete or Abort operation, it assumes the operation has failed

Design #5 – Transactional queues



So what did we do?

- We introduced **load balancers** to help with availability
- We introduced **multi-tier** architecture to help with security and availability
- We introduced **hot-standby** for availability given the security requirement
- We introduced **queue** to help with availability and increase decoupling to the AFTN network itself
- We introduced the concept of **transactions** to help with integrity
- All of this is to cope with requirements that are mainly non-functional!

A few final thoughts

- This just showed some of the design choices made to comply with non-functional requirements in the AviationCloud filing system
- Architecture is a balance between foreseeing what is required and not ending up making overkill solutions
- "It's future Henriks problem" aka agile development :-)
- On the other hand, architectural changes are some of the most challenging ones for a production-system.